

Notes on Geometric Learning Papers

Matthew Mo

2021 年 4 月 13 日

目录

1 NeVAE	9
1.1 Encoder	9
1.2 Decoder	10
1.3 Training	11
1.4 Property Oriented Mol. Gen.	11
2 Seminar on Self/Un-Supervised Learning @ 2020/9/16	12
2.1 Self-Learning @ Video Learning	12
2.2 Transformation Equivariance vs. Invariance @ Visial Repr. Learning	14
3 AET, AVT: Autoencoding Transformations	16
4 Flow-Based Generative Models	18
4.1 Outline & Basics	18
4.2 MoFlow	19
4.2.1 GCF/Graph Conditional Flow	21
4.2.2 Validity Correction & Misc	22
4.3 GraphNVP	23
5 WGAN	23
6 GMMN+AE	24
6.1 Structure & Idea	24
6.2 Training	25
7 FoldingNet - An AntoEncoder	25

8 PointFlow: Flow-based Generative Model on Point Clouds	26
8.1 Continuous Normalizing Flow(CNF)	26
8.2 Variational Auto-Encoder	26
8.3 Model	27
9 FFJORD	29
9.1 CNF	29
9.2 Backpropagation through ODE Solutions with Adjoint Method	29
9.3 Unbiased Linear-Time Log-Density Estimation	29
10 Dequantization to Learn Discrete Distribution	30
10.1 Dequantization as Latent Variable Model	30
10.2 Variational Dequantization	31
10.3 Importance-Weighted Dequantization	31
10.4 Renyi Dequantization	32
10.5 Dequantization Distribution	32
10.6 (Choice of) Continuous Distribution	33
11 DGI: Deep Graph Infomax	33
11.1 Backgrounds, Approach, Math	33
11.2 Algorithm	34
12 GraphSAGE: Inductive Representation Learning on Graph	35
12.1 Embedding Generation/FP	35
12.2 Aggragator Selection	35
13 SGC: Simplified Graph Convolution	36
14 FastGCN	36
14.1 Method	36
14.2 Variance Reduction	37
15 GWNN: Wavelet Transform on Graph	38
15.1 Supplementary Math: Real and Complex Wavelets	38
15.2 Graph Wavelets	39
15.3 GWNN	40
15.3.1 Details	40
16 Graph Wavelets	40
16.1 经典小波变换/CWT	40
16.2 谱小波变换/SGWT	41

16.2.1	Scaling Functions	42
16.3	SGBT 的性质	42
16.3.1	Inverse SGBT	42
16.3.2	局域性	42
16.3.3	Spectral Wavelet Frames	43
16.4	Fast SGBT Approximation by Polynomials	43
16.4.1	Fast Approximation of Adjoint	44
16.4.2	Inverse Calculation	45
16.5	Implementations and Details	45
17	GMNN: Graph Markov Neural Network	46
17.1	Pseudolikelihood Variational EM	46
17.2	Inference	47
17.3	Learning	48
17.4	Optimization	48
18	ClusterGCN: Fast Deep & Large GCNs	49
18.1	Vanilla ClusterGCN: Cluster For Batch	49
18.2	Stochastic Multiple Partitions	50
18.3	Analysis of Deeper Networks	50
19	GAT: Graph Attention Network	51
20	Note on Probabilistic Graphical Models	51
20.1	Bayesian Networks	51
20.2	Undirected Networks	52
20.3	Local Probabilistic Models i.e. Specific Models Corresponds to Last 2 Sections	54
20.4	Temporal Models	55
21	RSCNN(CVPR 19')	55
21.1	Architecture	55
21.2	Details & Implementation	56
22	SimpleView(ICLR 21' Candidate)	56
22.1	Simple Review of Existing Protocols	56
22.2	Model: SimpleView	56
23	OT-Flow	57
23.1	Idea & Formulations	57
23.2	Parametrization of Model	57
23.3	Exact Hessian of Multilayer NN	58

24 Node2vec: Unsupervised Feature Learning	58
24.1 Basics	58
24.2 Biased Random Walk	59
24.3 Edge Feature	60
25 DeepWalk: Online Representation Learning	60
25.1 DeepWalk	60
25.2 SkipGram	61
25.3 Hierachichal Softmax	61
25.4 Parallelization	62
25.5 Variants	62
26 DAGNN: Towards Deeper GNN	62
26.1 Smoothness Metrics	63
26.2 Convergence of Propagation	63
26.3 DAGNN: Deep Adaptive GNN	63
27 t-SNE(t-Distributed Stochastic Neighbor Embedding)	64
27.1 SNE	64
27.2 UNI-SNE	64
27.3 t-SNE	64
27.4 Barnes-Hut-SNE	65
27.4.1 Approximating Input Simularities by Vantage-point Tree	65
27.4.2 Approximating t-SNE Gradients	65
28 Autoregressive Flows	66
28.1 Autoregressive Transformation	66
28.2 MAF: Masked Autoregressive Flow	67
28.3 IAF: Inverse Autoregressive Flow	67
28.4 Sylvester NF(UAI 18')	67
28.4.1 Idea	67
28.4.2 Parametrization of A & B	68
28.4.3 Preserving Orthogonality of Q	68
29 Contrastive Multi-view Representation Learning/MVRLG(ICML 20?)	69
29.1 Augmentations	69
29.2 Encoders	69
29.3 Training	69

30 GCC: Graph Contrastive Coding for GNN Pre-Training(KDD 20’)	71
30.1 GCC Pre-Training	71
30.2 Finetuning GCC	72
31 GIN: Graph Isomorphism Network(ICLR 19’)	72
31.1 Weisfeiler-Lehman Test	72
31.2 Math Intuitions	73
31.3 GIN	73
31.4 Graph Readout of GIN	73
32 GraphLoG: Self-Supervised Represtation Learning with Local & Global Structure	74
32.1 Preliminaries	74
32.2 Local-Inst. Stru. Learning	74
32.3 Global-Semantic Repr. Learning	75
32.3.1 Init. of HP(Hierarchical Prototypes)	75
32.3.2 Maintainance of HP	75
32.4 Sup-GraphLoG: A Supervised Baseline	76
33 Orthogonal Weights in DNNs	76
33.1 Formulation & Good Properties	76
33.2 OWN: Orthogonal Weight Normalization	76
33.2.1 Backpropagation	77
33.2.2 As Convolution	77
33.2.3 Group Based Orthogonalization: Divided Filters	78
34 OrthDNNs: Orthogonal DNNs(TPAMI 19’)	78
34.1 GE Analysis in a Robustness and Isomeric Mapping Perspective	78
34.2 GE Analysis of DNN	79
34.3 OrthDNN by SVB(Singular Value Bound)	79
34.3.1 BN Compatibility	80
34.3.2 On CNN: OrthDNN as Convolution	80
35 SimCLR: A Simple Framework for Contrastive Learning of Visual Representations	80
35.1 Ideas & Basics	80
36 BYOL: Build Your Own Latent	81
36.1 Ideas & Method	81

37 MoCo: Momentum Contrast for Unsupervised Visual Representation Learning	82
37.1 Ideas & Method	82
38 SimSiam: Exploring Simple Siamese Representation Learning	83
38.1 Intuitions	83
38.2 Method	83
39 Graph Layouts by t-SNE	85
39.1 Backgrounds	85
39.2 Method	85
39.3 tsNET	85
40 细粒度图像数据分类 by Xiangteng He	85
40.1 细粒度图像数据分类	85
40.2 RL-based 图像部件/对象识别	85
40.3 多层注意力区域辨识	85
40.4 多模态	85
41 Dirac Operator for Extrinsic Shape Analysis	86
41.1 Math	86
41.2 离散化	87
41.3 实值表示	87
41.4 有界区域: 边界条件	87
42 Mesh-Based Simulation with GNNs	87
42.1 结构	87
42.2 Adaptive Remeshing	88
43 SENet	89
44 Continuous-Time Spiking Neural Network	90
44.1 Neurons	90
44.2 Network Topology	90
44.3 突触塑性规则	90
45 Towards Deep Learning Models Resistant to Adversarial Attacks	91
45.1 Inner Maximize Prob. : 如何提出好的对抗样本	91
45.2 最小化问题	91
45.3 网络能力 & 对抗健壮性	92
46 CAS: Channel-wise Activation Suppressing Module for Adversarial Robustness	92

47 Resisting Adversarial Attacks by k-Winners-Takes-All	93
47.1 Related Work: Obfuscated Gradients	93
47.2 k -Winners-Takes-All	93
47.3 Training	94
47.4 Theory Understand of the Discontinuity	94
47.5 Related Works	94

世界はこんなに美しいべきが、僕が同じくらい醜い。
だがこの状況を変わチカラはない。

Dedicated to Amatsukaze and Elaina, Fairies From The Ideal World.

1 NeVAE

Idea VAE on graph, node-wise repr, permutation invariant.

1.1 Encoder

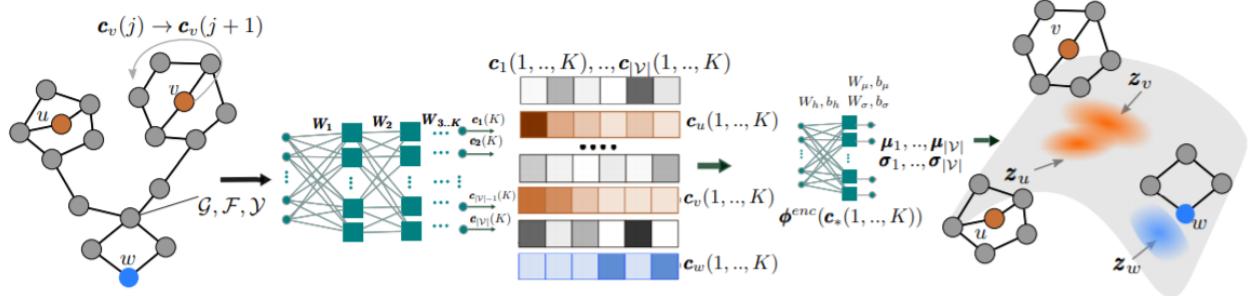


Figure 1: The encoder of our variational autoencoder for molecular graphs. From left to right, given a molecular graph \mathcal{G} with a set of node features \mathcal{F} and edge weights \mathcal{Y} , the encoder aggregates information from a different number of hops $j \leq K$ away for each node $v \in \mathcal{G}$ into an embedding vector $\mathbf{c}_v(j)$. These embeddings are fed into a differentiable function ϕ^{enc} which parameterizes the posterior distribution q_ϕ , from where the latent representation of each node in the input graph are sampled from.

GNN-like message-passing on k-hops:

$$q_\phi(\mathbf{z}_u | \mathcal{V}, \mathcal{E}, \mathcal{F}, \mathcal{Y}) \sim \mathcal{N}(\boldsymbol{\mu}_u, Diag(\boldsymbol{\sigma}_u)) \quad (1)$$

$$[\boldsymbol{\mu}_u, \boldsymbol{\sigma}_u] = \phi^{enc}(\mathbf{c}_u(k)_{k=1..K}) \quad (2)$$

$$\mathbf{c}_u(k) = \begin{cases} \mathbf{r}(\mathbf{W}_k^T \mathbf{t}_u + \mathbf{W}_k^X \mathbf{x}_u), & k = 1 \\ \mathbf{r}(\mathbf{W}_k^T \mathbf{t}_u + \mathbf{W}_k^X \mathbf{x}_u \odot \Lambda(\{y_{uv} \mathbf{c}_v(k-1)\}_v \in N(u))), & k > 1 \end{cases} \quad (3)$$

1.2 Decoder

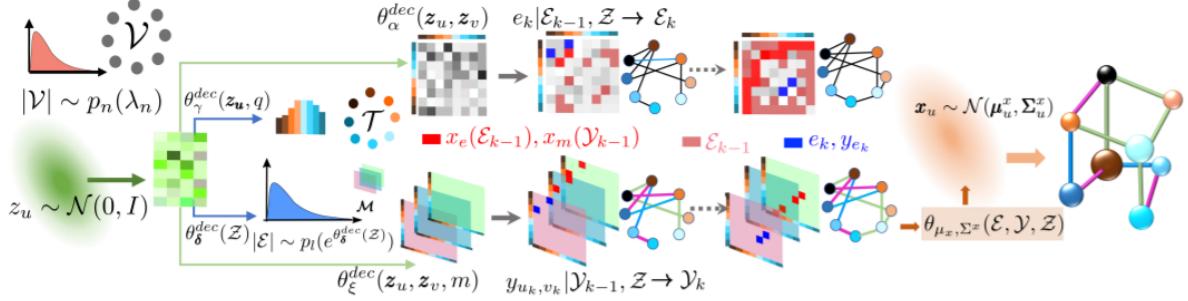


Figure 2: The decoder of our variational autoencoder for molecular graphs. From left to right, the decoder first samples the number of nodes $n = |\mathcal{V}|$ from a Poisson distribution $p_n(\lambda_n)$ and it samples a latent vector \mathbf{z}_u per node $u \in \mathcal{V}$ from $\mathcal{N}(\mathbf{0}, \mathbf{I})$. Then, for each node u , it represents all potential node feature values as an unnormalized log probability vector (or ‘logits’), where each entry is given by a nonlinearity θ_γ^{dec} of the corresponding latent representation \mathbf{z}_u , feeds this logit into a softmax distribution and samples the node features. Next, it feeds all latent vectors \mathcal{Z} into a nonlinear log intensity function $\theta_\delta^{dec}(\mathcal{Z})$ which is used to sample the number of edges. Thereafter, on the top row, it constructs a logit for all potential edges (u, v) , where each entry is given by a nonlinearity θ_α^{dec} of the corresponding latent representations $(\mathbf{z}_u, \mathbf{z}_v)$. Then, it samples the edges one by one from a soft max distribution depending on the logit and a mask $\beta_e(\mathcal{E}_{k-1})$, which gets updated every time it samples a new edge e_k . On the bottom row, it constructs a logit per edge (u, v) for all potential edge weight values m , where each entry is given by a nonlinearity θ_ξ^{dec} of the latent representations of the edge and edge weight value $(\mathbf{z}_u, \mathbf{z}_v, m)$. Then, every time it samples an edge, it samples the edge weight value from a soft max distribution depending on the corresponding logit and mask $x_m(u, v)$, which gets updated every time it samples a new y_{u_k, v_k} . Finally, for each atom u , it samples its coordinates \mathbf{x}_u from a multidimensional Gaussian distribution whose mean μ_x and variance Σ_x depends on the latent vectors of the corresponding atom and its neighbors and the underlying chemical bonds.

Decoder: gen. logits, softmax edges one-by-one, possible binary mask(for expert exp.).

$$\text{Nodes Count: } |\mathcal{V}| \sim p_l(\lambda_n) \quad (4)$$

$$\text{Latent Repr.: } \mathbf{z}_u \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \quad (5)$$

$$\text{Node Feat.: } \mathbf{f}_u = \text{softmax}_u(\theta_\gamma^{dec}(\mathbf{z}_u, q)), q \text{ is atom type} \quad (6)$$

$$\text{Edges Count: } |\mathcal{E}| \sim p_l(e^{\theta_\delta^{dec}(\mathcal{Z})}) \quad (7)$$

$$\text{Edges Gen.: } p(e = (u, v) | \mathcal{E}_{k-1}, \mathcal{V}) = \frac{\beta_e e^{\theta_\alpha^{dec}(\mathbf{z}_u, \mathbf{z}_v)}}{\sum_{e' = (u', v') \notin \mathcal{E}_{k-1}} \beta_{e'} e^{\theta_\alpha^{dec}(\mathbf{z}'_u, \mathbf{z}'_v)}} \quad (8)$$

$$\text{E. Feat. Gen.: } p(y_{uv} = m | \mathcal{Y}_{k-1}, \mathcal{V}) = \frac{\beta_m(u, v) e^{\theta_\xi^{dec}(\mathbf{z}_u, \mathbf{z}_v, m)}}{\sum_{m' \neq m} \beta'_{m'}(u, v) e^{\theta_\xi^{dec}(\mathbf{z}_u, \mathbf{z}_v, m')}}, \text{ note: not normal softmax?} \quad (9)$$

$$\text{Pos. Gen.: } p(\mathbf{x}_u | \mathcal{E}, \mathcal{Y}, \mathcal{Z}) = \mathcal{N}(\boldsymbol{\mu}_x, \boldsymbol{\Sigma}_x) \quad (10)$$

$$[\boldsymbol{\mu}_x, \boldsymbol{\Sigma}_x] = [\theta_{\mu^x}(\mathbf{r}(u)), \theta_{\Sigma^x}(\mathbf{r}(u)) \theta_{\Sigma^x}^T(\mathbf{r}(u))] \quad (11)$$

$$\mathbf{r}(u) = \mathbf{z}_u + \sum_{v \in N(u)} y_{uv} \mathbf{z}_v \quad (12)$$

1.3 Training

- **Prior:** $\mathcal{Z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- maximize evidence lower bound(ELBO)+Poisson max-likelihood:

$$\max_{\phi, \theta, \lambda_n} \frac{1}{N} \sum_{i \in [N]} \mathbb{E}_{q_\phi(\mathcal{Z}_i | \mathcal{Y}_i, \mathcal{E}_i, \mathcal{F}_i, \mathcal{Y}_i)} [\log p_\theta(\mathcal{Y}_i, \mathcal{E}_i, \mathcal{F}_i, |\mathcal{Z}_i)] - KL(q_\phi || p_z) + \log p_{\lambda_n}(n_i) \quad (13)$$

- note term $E_{q_\phi}[\log p_\theta(\mathcal{Y}_i, \mathcal{E}_i, \mathcal{F}_i, |\mathcal{Z}_i)]$ need the edges sequence specified, use BFS with random tie breaking in child-sel. step, with random selected source node $s \sim \zeta_s$! Thus

$$E_{q_\phi}[\log p_\theta(\mathcal{Y}_i, \mathcal{E}_i, \mathcal{F}_i, |\mathcal{Z}_i)] \approx E_{q_\phi}[\log \mathbb{E}_{s \sim \zeta_s} p_\theta(\mathcal{Y}_i, \mathcal{E}_i, \mathcal{F}_i, |\mathcal{Z}_i)] \quad (14)$$

$$\geq E_{q_\phi, s \sim \zeta_s} [\log p_\theta(\mathcal{Y}_i, \mathcal{E}_i, \mathcal{F}_i, |\mathcal{Z}_i)] \quad (15)$$

- **Theorem** If dist. ζ_s is independent to labels of nodes, then the learned model is permutation-invariant.
- **Proposition** Decoder defined is permutation-invariant.

1.4 Property Oriented Mol. Gen.

Train variational probabilistic decoder, to maxmize some property of mol., \Rightarrow train a supervised decoder p^* on trained decoder p_θ :

$$\min_{p(\cdot|\mathcal{Z})} \mathbb{E}_{\mathcal{Z} \sim p_z(\cdot)} \mathbb{E}_{\mathcal{E}, \mathcal{Y}, \mathcal{F} \sim p(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})} [l(\mathcal{E}, \mathcal{Y}, \mathcal{F}) + \rho \log \frac{p(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})}{p_\theta(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})}] \quad (16)$$

$$\Rightarrow \min \mathbb{E}_{\mathcal{Z} \in p_z} [KL(p(\cdot|\mathcal{Z}) || g_\theta(\cdot|\mathcal{Z}))] \quad (17)$$

$$\text{where } g_\theta(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z}) = \frac{p_\theta(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z}) \exp\left(-\frac{l(\mathcal{E}, \mathcal{Y}, \mathcal{F})}{\rho}\right)}{\mathbb{E}_{\mathcal{E}, \mathcal{Y}, \mathcal{F} \sim p_\theta(\cdot|\mathcal{Z})} [p_\theta(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z}) \exp\left(-\frac{l(\mathcal{E}, \mathcal{Y}, \mathcal{F})}{\rho}\right)]} \quad (18)$$

The above equations has a obvious solution $p^* \equiv g_\theta$, however sampling might be too slow for practical use \Rightarrow A Stochastic Gradient Approach.

Algorithm 1: PROPERTYORIENTEDDECODER: it trains a parameterized property-oriented decoder.

-
- 1: **Given:** The loss function $\ell(\cdot)$, parameter ρ , original decoder p_θ , # of iterations M , mini batch size B , and learning rate γ
 - 2: $\theta'_0 \leftarrow \theta$
 - 3: **for** $j = 1, \dots, M$ **do**
 - 4: $\mathcal{Z}_j \sim p_z(\cdot)$
 - 5: $\mathcal{D} \leftarrow \text{MINIBATCH}(p_{\theta'_j}(\cdot|\mathcal{Z}_j), B)$
 - 6: $\nabla \leftarrow 0$
 - 7: **for** $(\mathcal{E}_i, \mathcal{Y}_i, \mathcal{F}_i) \in \mathcal{D}$ **do**
 - 8: $S \leftarrow \ell(\mathcal{E}_i, \mathcal{Y}_i, \mathcal{F}_i) + \rho \log \left(p_{\theta'_j}(\mathcal{E}_i, \mathcal{Y}_i, \mathcal{F}_i | \mathcal{Z}_j) / p_\theta(\mathcal{E}_i, \mathcal{Y}_i, \mathcal{F}_i | \mathcal{Z}_j) \right)$
 - 9: $\nabla \leftarrow \nabla + (S + \rho) \nabla_{\theta'} \log p_{\theta'_j}(\mathcal{E}_i, \mathcal{Y}_i, \mathcal{F}_i | \mathcal{Z}_j)$
 - 10: $\theta'_{j+1} \leftarrow \theta'_j + \gamma \frac{\nabla}{B}$
 - 11: **Return** θ'_M
-

Use SGD to update param. θ' of $p_{\theta'}$:

$$\Delta\theta' = \alpha \nabla_{\theta'} \mathbb{E}_{\mathcal{Z} \sim p_z(\cdot)} \mathbb{E}_{\mathcal{E}, \mathcal{Y}, \mathcal{F} \sim p(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})} [l(\mathcal{E}, \mathcal{Y}, \mathcal{F}) + \rho \log \frac{p(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})}{p_{\theta}(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})}] \quad (19)$$

$$= \alpha \mathbb{E}_{\mathcal{Z} \sim p_z(\cdot)} \nabla_{\theta'} \mathbb{E}_{\mathcal{E}, \mathcal{Y}, \mathcal{F} \sim p(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})} [l(\mathcal{E}, \mathcal{Y}, \mathcal{F}) + \rho \log \frac{p(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})}{p_{\theta}(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})}] \quad (20)$$

$$\text{(by log-deriv. trick)} = \alpha \mathbb{E}_{\mathcal{Z} \sim p_z(\cdot)} \mathbb{E}_{\mathcal{E}, \mathcal{Y}, \mathcal{F} \sim p(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})} \left[(l(\mathcal{E}, \mathcal{Y}, \mathcal{F}) + \rho \log \frac{p(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})}{p_{\theta}(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})} + \rho) \nabla_{\theta'} \log p_{\theta'} \right] \quad (21)$$

by a unbiased MC estim.

$$\approx \frac{1}{M} \sum_{i \in [M]} \left[(l(\mathcal{E}, \mathcal{Y}, \mathcal{F}) + \rho \log \frac{p(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})}{p_{\theta}(\mathcal{E}, \mathcal{Y}, \mathcal{F} | \mathcal{Z})} + \rho) \nabla_{\theta'} \log p_{\theta'} \right] \quad (22)$$

2 Seminar on Self/Un-Supervised Learning @ 2020/9/16

2.1 Self-Learning @ Video Learning

Supervised success: good & sufficient data, a way different from human! \Rightarrow Linda Smith, *The Dev. of Embodied Cognition*

Paragidims:

- Use proxy task(e.g. semantics repr.) for a repr., use linear probing for downstream task.
- Use proxy task(e.g. semantics repr.) for a repr., generalizable with *zero annotation*

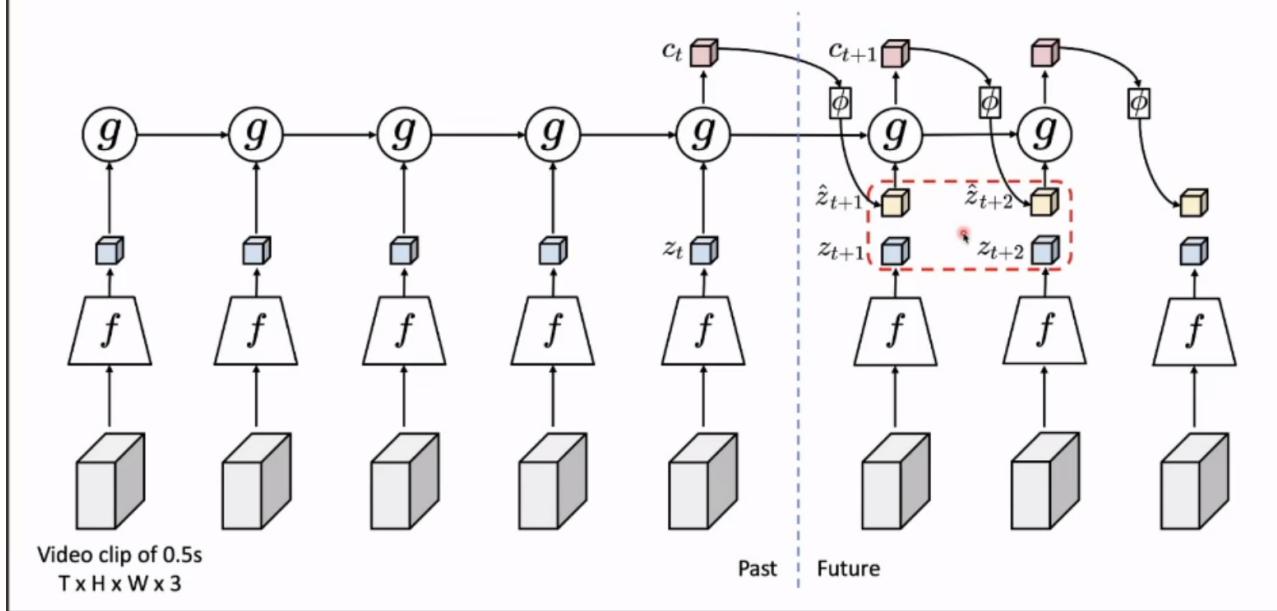
Why video-based self-supervised: like what human percepts, rich info; might with audio.

Proxy loss design: temporal info, spatial cohenrence, motions of obj., multimodal

Temporal:

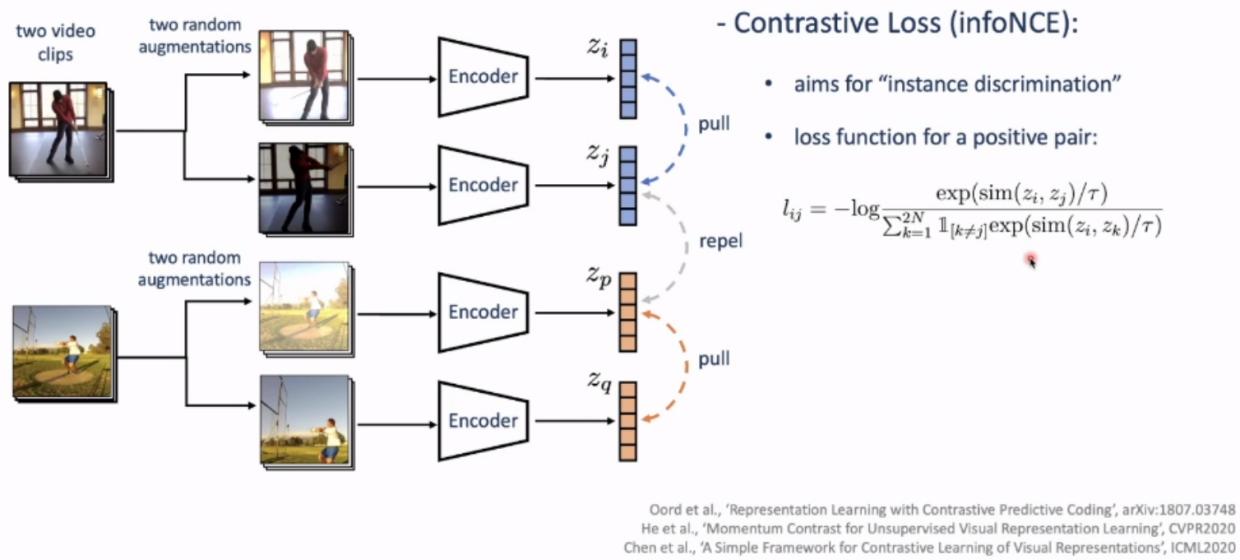
- shuffle & learn
- forward or backward?(arrow of time)
- SpeedNet: which speed(frame-rate) is normal/speed-up
- ===Weak, irrelative with downstream tasks==
- DPC: *learn repr. in predicting future in video*

Approach - DPC

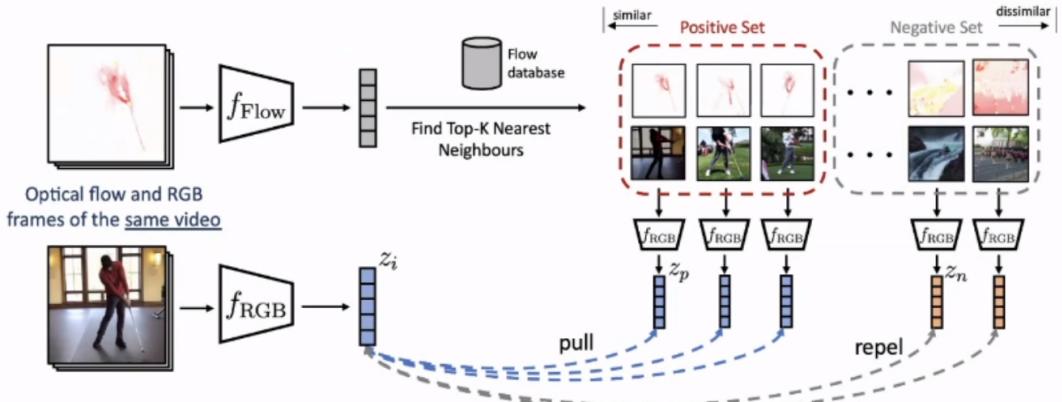


DPC Arch.:encoder-decoder like, contrast learning(infoNCE)

Self-supervised learning with videos (CoCLR)



Self-supervised learning with videos (CoCLR)



Multi-Instance Contrastive Loss (MIL-NCE):

- Features from the positive set are pulled together
- Features **NOT** from the positive set are pushed apart
- Optical flow helps RGB frames to go beyond instance discrimination

$$\mathcal{L}_{\text{CoCLR}} = -\mathbb{E} \left[\log \frac{\sum_{p \in \mathcal{P}_i} \exp(z_i \cdot z_p)}{\sum_{p \in \mathcal{P}_i} \exp(z_i \cdot z_p) + \sum_{n \in \mathcal{N}_i} \exp(z_i \cdot z_n)} \right]$$

CoCLR: SimCLR like(infoNCE), colearning multimodally with motion flow(MIL-NCE, with noise added)

Audio-Video Co-learning: train a net to check if image/audio clip are same-sourced! Get both video/audio repr.

MAST: self-supervised tracking, give 1st frame mask(seg.), predict sequential segmentations

Next:

- More efficient learning
- Scale up model to uncurated data, like GPT-1/2/3
- Design proxy task for obj-centric learning
- Design and understand **effective memory!!!**(for video task especially)
- Hand-crafted proxy task \Rightarrow Auto proxy task design?
- Theoretic: small or negative improvement, in upstream task to downstream task.
- Are there difficult task for supervised learning but easy for SSL(e.g. unable to label)?

2.2 Transformation Equivariance vs. Invariance @ Visial Repr. Learning

Contents:

- TER(Transformation Equivariance Repr.)

- AET(AutoEncoding Transformation)
 - AVT(Autoencoding Variational Transformation)
 - SAT(Semi-supervised Autoencoding Transformation)

CNN = Translation Equivariant Repr. + FC Classifier. Go beyond: Transformation Equivariant Repr. + Tranformation Invariant Classifier

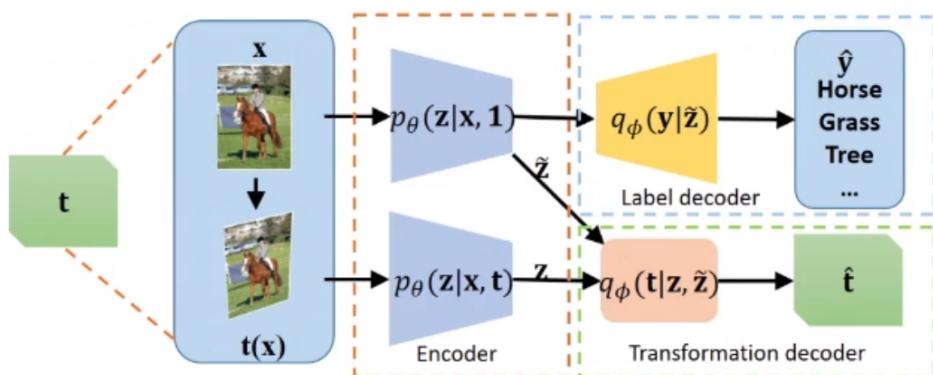
Trans. Equiv.: $E(\mathbf{t}(x)) = \rho(t)[E(x)]$. Trans. Inv. is when $\rho \equiv \mathbf{1}_E$.

Steerability: ρ is independent of sample x .

Targets: Non-linear ρ , General Transformation(e.g. recoloring)

SAT: Semi-Supervised Autoencoding Transformation

- Adding a label decoder $q_\phi(\mathbf{y}|\tilde{\mathbf{z}})$ to approximate the posterior $p_\theta(\mathbf{y}|\mathbf{x})$



Variational Bound

- By introducing label decoder and transformation decoder, we have

$$I_\theta(\mathbf{y}, \mathbf{z}; \tilde{\mathbf{z}}, \mathbf{t}) \geq \mathbb{E}_{p_\theta(\mathbf{y}, \mathbf{z}, \tilde{\mathbf{z}})} \log q_\phi(\mathbf{y} | \mathbf{z}, \tilde{\mathbf{z}}) + \mathbb{E}_{p_\theta(\mathbf{t}, \mathbf{z}, \tilde{\mathbf{z}})} \log q_\phi(\mathbf{t} | \mathbf{z}, \tilde{\mathbf{z}})$$

Label decoder
Transformation decoder

- Jointly maximizing over encoder θ and decoders ϕ

$$\max_{\theta, \phi} \mathbb{E}_{p_\theta(\mathbf{y}, \mathbf{z}, \tilde{\mathbf{z}})} \log q_\phi(\mathbf{y} | \mathbf{z}, \tilde{\mathbf{z}}) + \mathbb{E}_{p_\theta(\mathbf{t}, \mathbf{z}, \tilde{\mathbf{z}})} \log q_\phi(\mathbf{t} | \mathbf{z}, \tilde{\mathbf{z}})$$


SAT:

- add a label decoder compared to AVT.
- variational surrogate \Rightarrow cross-entropy loss on supervised data + AVT loss

Contrastive Learning: more utilized trans-invariant repr. Future: unifying trans-inv/equiv repr.

3 AET, AVT: Autoencoding Transformations

Idea autoencoders used in modeling transformations rather than images in order to learn general repr.



AET:

- use autoencoders to learn **transformations**
- trans. generated randomly for self-supervised learning
- use Siamese net as encoder backbone
- AET loss: parameterized, non-parametric, GAN-induced

Losses in AET:

- parameterized transformations: if trans. are parameterized $\mathcal{T} \in \{t_\theta | \theta \in \Theta\}$, loss can be defined as norm of param. diff.

$$l(t_\theta, \hat{t}_\theta) = \|\theta - \hat{\theta}\|.$$

- for non-parametric trans., use expected distance on source domain

$$l(t, \hat{t}) = \mathbb{E}_{x \sim X} \{ \text{dist}(t(x), \hat{t}(x)) \}$$

- GAN-induced trans.: image transformed in form $G(x, z)$, we have loss

$$l(t_z, t_{\tilde{z}} = \|z - \tilde{z}\|.)$$

Idea of AVT use prob. dist. to model trans., VAE like modeling!

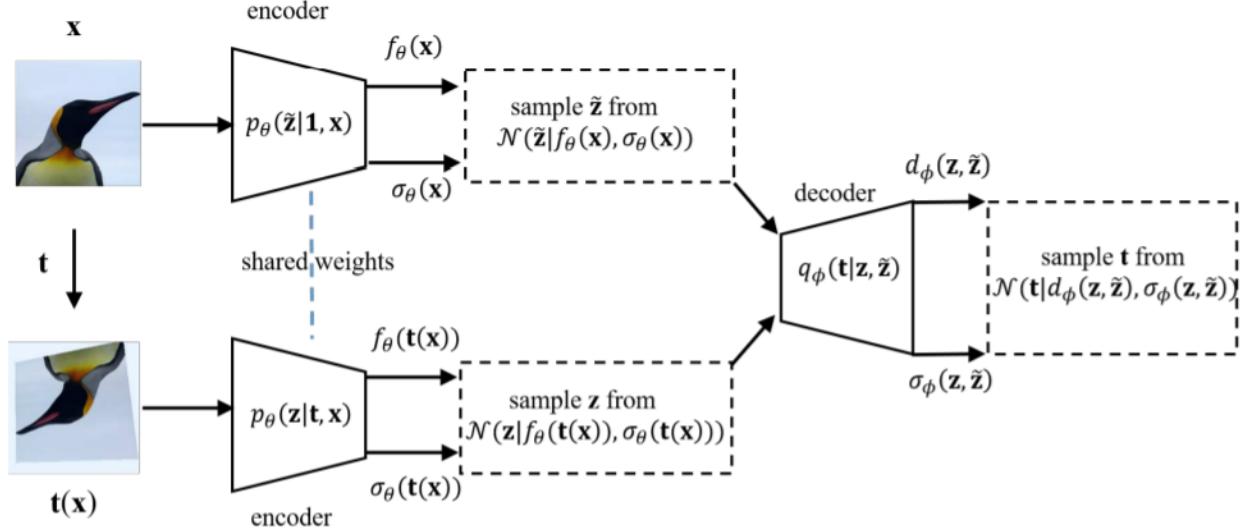


Figure 1: The architecture of the proposed AVT. The original and transformed images are fed through the encoder p_θ where $\mathbf{1}$ denotes an identity transformation to generate the representation of the original image. The resultant representations $\tilde{\mathbf{z}}$ and \mathbf{z} of original and transformed images are sampled and fed into the transformation decoder q_ϕ from which the transformation \mathbf{t} is sampled.

AVT:

- maximize mutual info $I(t; z|\tilde{z})$
- variational bound, introducing a decoder $q_\phi(t|z, \tilde{z})$:

$$I(t; z|\tilde{z}) = H(t|\tilde{z}) - H(t|z, \tilde{z}) \quad (23)$$

$$= H(t|\tilde{z}) + \mathbb{E}_{p_\theta(t, z, \tilde{z})}[p_\theta(t|z, \tilde{z})] \quad (24)$$

$$= H(t|\tilde{z}) + \mathbb{E}_{p(t, z, \tilde{z})}[q_\theta(t|z, \tilde{z})] + \mathbb{E}_{p(z, \tilde{z})}[D(p_\theta(t, z, \tilde{z})||q_\phi(t|z, \tilde{z}))] \quad (25)$$

$$\geq H(t|\tilde{z}) + \mathbb{E}_{p(t, z, \tilde{z})}[q_\phi(t|z, \tilde{z})] \equiv \tilde{I}(t; z|\tilde{z}) \quad (26)$$

$$\Rightarrow \max_{\theta, \phi} \mathbb{E}_{p(t, z, \tilde{z})}[q_\phi(t|z, \tilde{z})] \quad (27)$$

-
- specifically in batch-wise formulation:

$$\mathbb{E}_{p(t,z,\tilde{z})}[q_\phi(t|z,\tilde{z})] \approx \frac{1}{n} \sum_{i=1}^n \log \mathcal{N}(t^i | d_\phi(z^i, \tilde{z}^i), \sigma_\phi(z^i, \tilde{z}^i)) \quad (28)$$

$$\text{where } z^i = f_\theta(t^i(x^i)) + \sigma_\theta(t^i(x^i)) \odot \epsilon^i \quad (29)$$

$$\text{and } \tilde{z}^i = f_\theta(x^i) + \sigma_\theta(x^i) \odot \tilde{\epsilon}^i \quad (30)$$

$$\text{where } \epsilon^i, \tilde{\epsilon}^i \sim (\epsilon|0, I), t^i \sim p(t) \text{(predifined or so?)} \quad (31)$$

- trick: take 5 samples to full explore the distribution

4 Flow-Based Generative Models

4.1 Outline & Basics

Two random vector of same dim.:

$$X \sim P_X(x), z \sim \Pi_Z(z), \text{find mapping } f : Z \rightarrow X = x(z), \quad (32)$$

we have

$$\begin{cases} p_X(x) = \pi_Z(f^{-1}(x)) |\det J(f)|^{-1} \\ \pi_Z(z) = p_X(f(z)) |\det J(f)| \end{cases} \quad (33)$$

use a simple dist. on Z and invertibly generate $X \sim p_G(x)$: $x = G(z)$. train G^{-1} as a discriminator.
keys: invertible, easy-to-compute G^{-1} , easy-to-compute Jacobian determinant.

“Coupling Layer”:

$$\begin{cases} (\text{copy}) x_i = z_i, i \leq d \\ (\text{affine}) x_i = \beta_i z_i + r_i, d < i \leq D \end{cases} \quad (34)$$

$$\beta_{d+1, \dots, D} = F(z_{d+1, \dots, D}), \gamma_{d+1, \dots, D} = H(z_{d+1, \dots, D}) \quad (35)$$

$$J_G = \left[\begin{array}{c|c} \mathbf{I}_d & \mathbf{O} \\ \hline M(\text{non-matter}) & D(\text{diagonal}) \end{array} \right] \quad (36)$$

$$\det J_G = \prod_{k=d+1..D} \frac{\partial x_k}{\partial z_k} \quad (37)$$

Use many coupling layer to enhance expressive capability. Parts of image does not change
⇒exchange copy/affine split:

- exchange within channel
- exchange channel ⇒channel rotation use matrix/ 1×1 convolution in MoFlow

4.2 MoFlow

Summary Flow-based on molecular graphs, channel rotation as 1×1 convolution, relational GCN layer, graph conditional flow(GCF), use sigmoid rather than exp, split dimensions.

“Coupling Layer”:

$$Z_{1:d} = X_{1:d} \quad (38)$$

$$Z_{d+1:n} = X_{d+1:n} \odot \text{sigmoid}(S_\Theta(X_{1:d})) + T_\Theta(X_{1:d}) \quad (39)$$

here $S \sim$ scaling, $T \sim$ translation, both by DNNs. 每个耦合层交换上一层 copy 的 dims, 通过一个 channel 上的旋转 $W \in \mathbb{R}^{c \times c}$, 等价于一个 $1*1$ 卷积, 变换后的 Y 分为 $(Y_{1:c/2}, Y_{c/2+1,n})$ 送入下一层. 采用 split-dims 的 trick, 增加交换 channel 的模型自由度增加: $X \in \mathbb{R}^{c \times n \times n} \Rightarrow \mathbb{R}^{ch^2 \times n/h \times n/h}$

4.2.1 GCF/Graph Conditional Flow

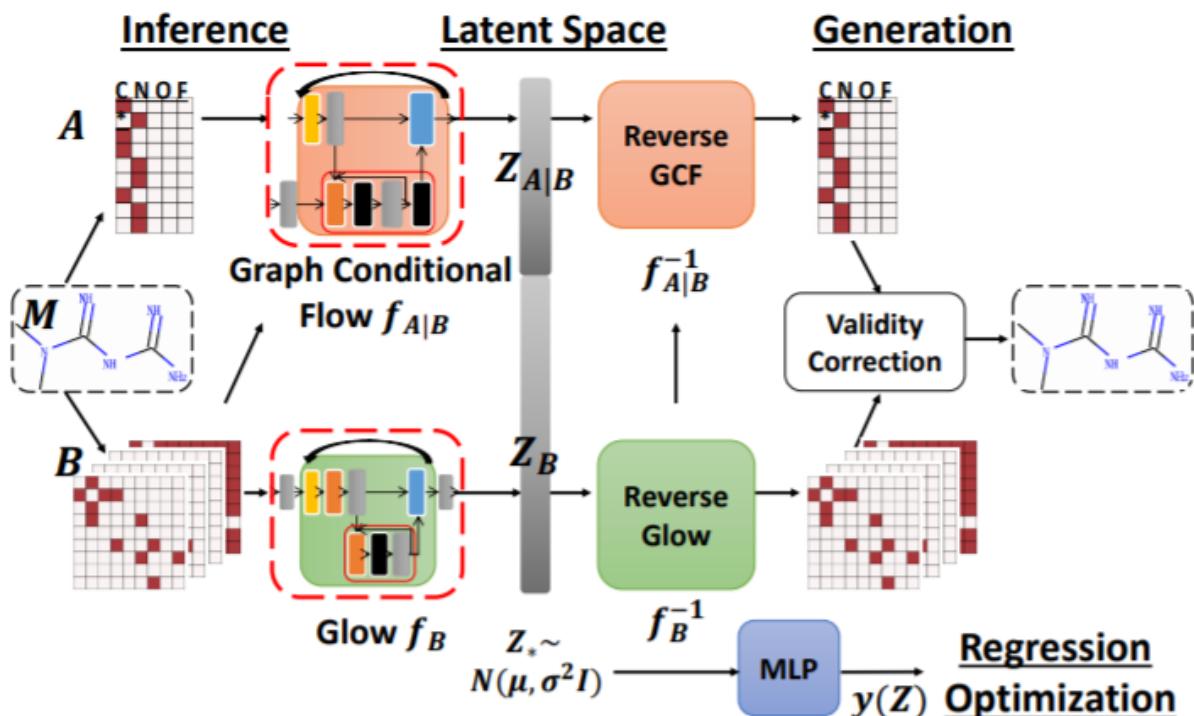


Figure 1: The outline of our MoFlow. A molecular graph M (e.g. Metformin) is represented by a feature matrix A for atoms and adjacency tensors B for bonds. **Inference:** the graph conditional flow (GCF) $f_{\mathcal{A}|\mathcal{B}}$ for atoms (Sec. 4.2) transforms A given B into conditional latent vector $Z_{A|B}$, and the Glow f_B for bonds (Sec. 4.3) transform B into latent vector Z_B . The latent space follows a spherical Gaussian distribution. **Generation:** the generation process is the reverse transformations of previous operations, followed by a validity correction (Sec. 4.4) procedure which ensures the chemical validity. We summarize MoFlow in Sec. 4.5. **Regression and optimization:** the mapping $y(Z)$ between latent space and molecular properties are used for molecular graph optimization and property prediction (Sec. 5.3, Sec. 5.4).

Def.(B conditioned flow)

We have

$$J_{A|B} = \frac{\partial f_{A|B}}{\partial(A, B)} = \left[\begin{array}{c|c} \frac{\partial f_{A|B}}{\partial A} & \frac{\partial f_{A|B}}{\partial B} \\ \hline \mathbf{O} & \mathbf{I} \end{array} \right] \quad (40)$$

$$\det J_{A|B} = \det \frac{\partial f_{A|B}}{\partial A} \quad (41)$$

GCF layer:

$$Z_{A|B} = (Z_{A_1|B}, Z_{A_2|B}), A = (A_1|A_2) \quad (42)$$

$$\begin{cases} \text{copy } Z_{A_1|B} = A_1 \\ \text{affine } Z_{A_2|B} = A_2 \odot \text{sigmoid}(S_\Theta(A_1|B)) + T_\Theta(A_1|B) \end{cases} \quad (43)$$

Special designed S, T using R-GCN:

$$\text{graphconv}(A_1) = \sum_{i=[C]} \tilde{B}_i (M \odot A) W_i + (M \odot A) W_0, \text{ where } M \text{ is the mask of split}, \quad (44)$$

$$\tilde{B} \text{ is normalized } \mathbf{B}_i : \tilde{B}_i = \mathbf{D}^{-1} \mathbf{B}_i, \quad (45)$$

$$\text{where } \mathbf{D} \text{ is the full deg-mat. } \mathbf{D} = \sum_{c,i} \mathbf{B}_{c,i,j} = \sum_c \mathbf{D}_i, \text{ computed only once!} \quad (46)$$

4.2.2 Validity Correction & Misc

使用价约束

$$\sum_{c,j} c \times B_{c,i,j} \leq \text{Valency}_i + \text{Ch}_i, \text{ where } \text{Ch}_i \text{ is formal charge} \quad (47)$$

$$(48)$$

具体的有效性校验方法:

1. 检查价约束, 满足去 2, 否则去 3
2. 返回最大连通子图
3. 第 i 个原子不满足, 对于第 i 个原子, 删去最高阶键, 去 1

这种方法试图再分子上做最小修改来满足价约束.

Note 为了防止学到的 prob. dist. 退化, 在数据集上增加 dequantization, 每个 dim 加噪声 $\sim U[0, 0.6]$

4.3 GraphNVP

5 WGAN

Idea Wasserstain 距离代替 KL/JSD 距离.

Method

- 判别器不用 sigmoid, loss 不取 log
- 判别器参数截断 \Rightarrow 为了让判别器 Lipschitz 连续.
- trick: 不用基于 momentum 的优化器 (Adam etc.), 用 RMSProp/SGD.

6 GMMN+AE

6.1 Structure & Idea

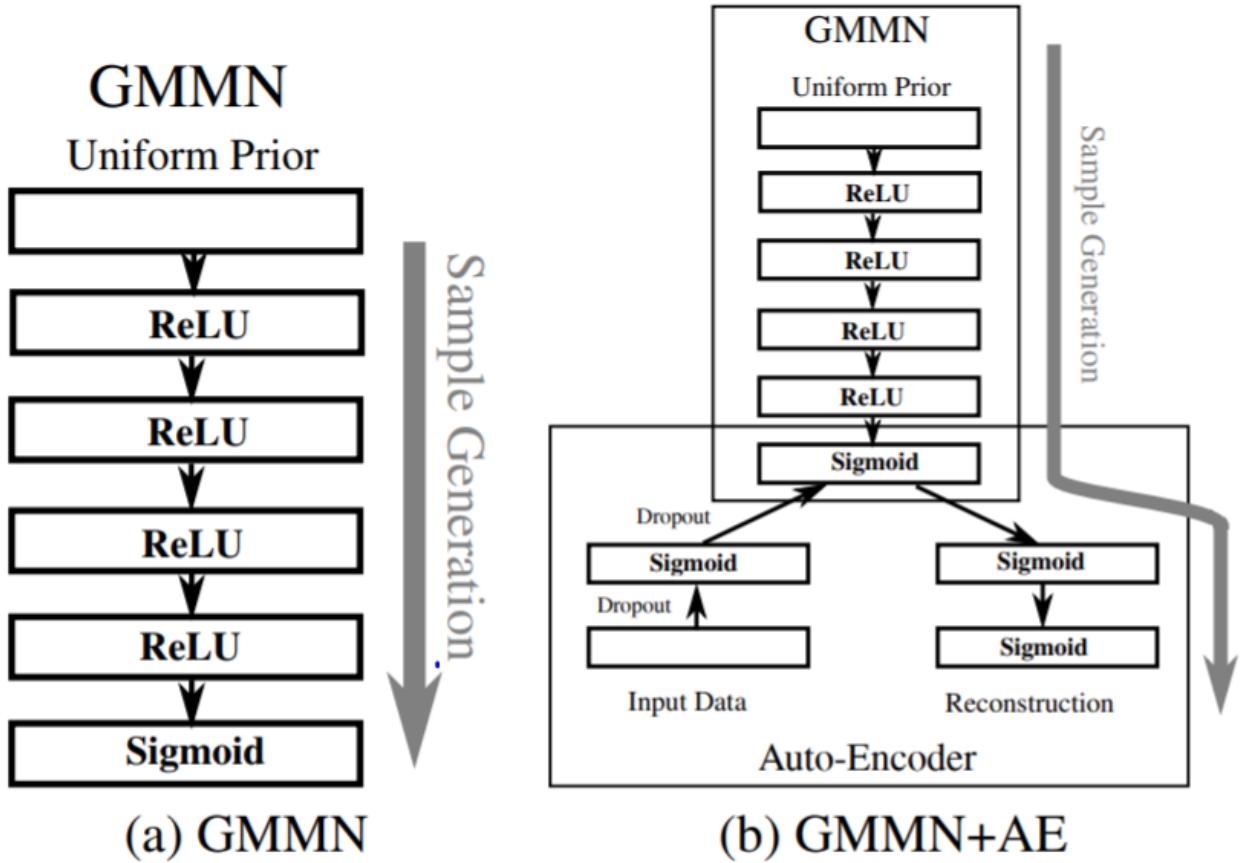


Figure 1. Example architectures of our generative moment matching networks. (a) GMMN used in the input data space. (b) GMMN used in the code space of an auto-encoder.

Use MMD(Maximum Mean Discrepancy) loss:

$$L_{MMD^2} = \left\| \frac{1}{N} \sum_i \phi(x_i) - \frac{1}{M} \sum_j \phi(y_j) \right\|^2 \quad (49)$$

$$= \frac{1}{N^2} \sum_{i,i'} K(x_i, x_{i'}) + \frac{1}{M^2} \sum_{i,i'} K(y_i, y_{i'}) - \frac{1}{NM} \sum_{i,j} K(x_i, y_j) \quad (50)$$

使用 k 阶多项式作为核, 则等价于匹配 k 阶矩! \Rightarrow 使用高斯核, 以匹配所有阶矩 (看作幂级数), 这也是 GMMN 的名字由来 (Moment-Matching):

$$K(x, y) = \exp\left(-\frac{1}{2\sigma} \|x - y\|^2\right) \quad (51)$$

设生成的数据为 $(x_i^s)_{gt}$. 为 (x_i^d) , 则偏导

$$\frac{\partial L_{MMD^2}}{\partial x_{ip}^s} = \frac{1}{\sigma} \left(\frac{2}{M^2} \sum_{j=[M]} K(x_i^s, x_j^s)(x_{jp}^s - x_{ip}^s) - \frac{2}{NM} \sum_{j=[N]} K(x_i^s, x_j^d)(x_{jp}^d - x_{ip}^s) \right) \quad (52)$$

6.2 Training

1. 逐层训练 AE
2. Finetune AE
3. 训练 GMMN

7 FoldingNet - An AutoEncoder




Figure 1. **FoldingNet Architecture**. The graph-layers are the graph-based max-pooling layers mentioned in (2) in Section 2.1. The 1st and the 2nd folding are both implemented by concatenating the codeword to the feature vectors followed by a 3-layer perceptron. Each perceptron independently applies to the feature vector of a single point as in [41], i.e., applies to the rows of the m -by- k matrix.

使用 (扩展的)Chamfer 距离

$$d_{CH}(S, \widehat{S}) = \max\left\{ \frac{1}{|S|} \sum_{x \in S} \min_{x \in \widehat{S}} \|x - \widehat{x}\|, \frac{1}{|\widehat{S}|} \sum_{\widehat{x} \in \widehat{S}} \min_{x \in S} \|x - \widehat{x}\| \right\} \quad (53)$$

这个距离让两个点云的点互相配准.

使用基于图的 Encoder: 使用的特征为局部 (KNN 上的) 协方差¹+ 位置 ($n \times 12$), 简要结构:
MLP+GNN-Aggregation+MLP⇒Codeword 其中 Graph Layers

$$\mathbf{Y} = \mathbf{A}_{\max}(\mathbf{X})\mathbf{K} \quad (54)$$

$$\mathbf{A}_{\max}(\mathbf{X})_{ij} = \text{ReLU}\left(\max_{k \in \mathcal{N}(i)} x_{kj}\right) \quad (55)$$

基于折叠的 Decoder: 重复 m 次 codeword, 和 $2d$ 格点 concat 送到 MLP(1st-folding) 得到中间折叠点云, 和 codeword concat 之后再送到第二个 folding-mlp 中得到结果.

Prop. Encoder proposed is permutation-invariant.

Prop. Decoder proposed can shape arbitrary point cloud.

8 PointFlow: Flow-based Generative Model on Point Clouds

Idea As Title

8.1 Continuous Normalizing Flow(CNF)

正则化流, 通过一系列可逆变换 f_i :

$$x = f_1 \circ \dots \circ f_n(y) \quad (56)$$

$$\log P(x) = \log P(y) - \sum_i |\log \det \mathcal{J}_{f_i}| \quad (57)$$

离散的正则化流被推广到连续的正则化流—CNFs

$$\frac{\partial y(t)}{t} = f(y(t), t) \quad (58)$$

$$\text{Thus } = y(t_0) + \int_{t_0}^{t_1} f(y(t), t) dt, y(t_0) \sim P(y) \quad (59)$$

$$\log P(x) = \log P(y(t_0)) - \int_{t_0}^{t_1} \mathcal{T}r\left(\frac{\partial f}{\partial y(t)}\right) dt \quad (60)$$

一个黑盒 ODE 求解器可以用于估计流的输出和输入的梯度!

8.2 Variational Auto-Encoder

Optimize ELBO

$$\log P_\theta(X) \geq \log P_\theta(X) - D_{KL}(Q_\phi(z|X)||P_\theta(z|X)) \quad (61)$$

$$= \mathbb{E}_{Q_\phi(z|X)}[\log P_\theta(X|z)] - D_{KL}(Q_\phi(z|X)||P_\psi\theta(z)) \quad (62)$$

¹回忆协方差公式

$$\text{Cov}(\mathbf{X}) = \mathbb{E}[(\mathbf{X} - \mathbb{E}\mathbf{X})(\mathbf{X} - \mathbb{E}\mathbf{X})^T]$$

8.3 Model

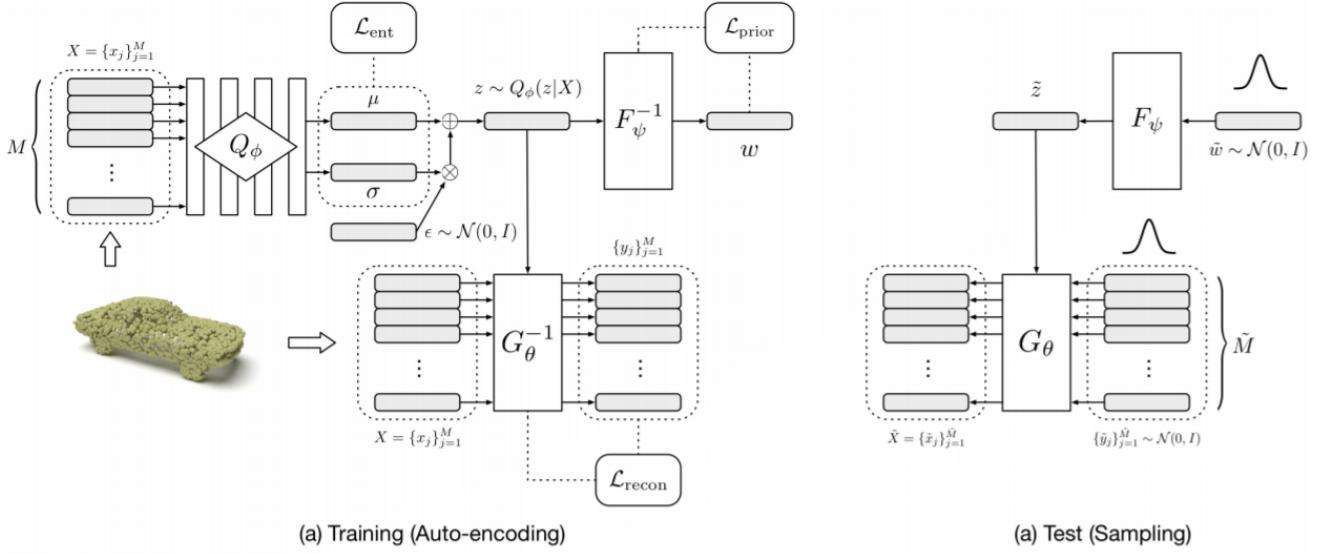


Figure 2: Model architecture. (a) At training time, the encoder Q_ϕ infers a posterior over shape representations given an input point cloud X , and samples a shape representation z from it. We then compute the probability of z in the prior distribution ($\mathcal{L}_{\text{prior}}$) through a inverse CNF F_ψ^{-1} , and compute the reconstruction likelihood of X ($\mathcal{L}_{\text{recon}}$) through another inverse CNF G_θ^{-1} conditioned on z . The model is trained end-to-end to maximize the evidence lower bound (ELBO), which is the sum of $\mathcal{L}_{\text{prior}}$, $\mathcal{L}_{\text{recon}}$, and \mathcal{L}_{ent} (the entropy of the posterior $Q_\phi(z|X)$). (b) At test time, we sample a shape representation \tilde{z} by sampling \tilde{w} from a Gaussian prior and transforming it with F_ψ . To sample points from the shape represented by \tilde{z} , we first sample points from the 3-D Gaussian prior and then move them according to the CNF parameterized by \tilde{z} .

Summary VAE-like. Decoder: Flow-based, i.e. CNF; Prior: CNF-based; Encoder: some simple permutation-invariant encoder.

Notations

$$z \sim \text{Latent Repr. for Shape} \quad (63)$$

$$y \sim \text{Simple Distribution/Source Dist. to be Transformed} \quad (64)$$

$$x \sim \text{Point Cloud} \quad (65)$$

$$(66)$$

Point cloud lld

$$\log P_\theta(X|z) = \sum_{x \in X} \log P_\theta(x|z) \quad (67)$$

model $P(x|z)$ by 条件 CNF

$$x = G_\theta(y(t_0); z) \quad (68)$$

$$= y(t_0) + \int_{t_0}^{t_1} g_\theta(y(t), t; z) dt, y(t_0) \sim P(y) = \mathcal{N}(0, I) \quad (69)$$

reconstruction lld:

$$\log P(x) = \log P(y(t_0)) - \int_{t_0}^{t_1} \mathcal{J}_{g_\theta(t)} dt \quad (70)$$

虽然用高斯分布的先验在 shape repr. 上可行, 但是有证据证明这受限的分布先验在 VAE 中会限制性能. 使用另一个 CNF 来参数化可学习的先验来减少影响

$$D_{KL}(Q_\phi(z|X)||P_\psi\theta(z)) = \mathbb{E}_{Q_\phi(z|X)}[\log P_\psi\theta(z)] - H(P_\psi\theta(z)) \quad (71)$$

obtain P_ψ by $P(w) \sim \mathcal{N}(0, I)$ and CNF

$$z = F_\psi(w(t_0)) \quad (72)$$

$$\triangleq w(t_0) + \int_{t_0}^{t_1} f_\psi(w(t), t) dt, w(t_0) \sim P(w) = \mathcal{N}(0, I) \quad (73)$$

log-probability

$$\log P(x) = \log P(F_\psi^{-1}(z)) - \int_{t_0}^{t_1} \mathcal{J}_{f_\psi(t)} dt \quad (74)$$

最终的 loss term(ELBO)

$$\begin{aligned} \mathcal{L}(X; \phi, \psi, \theta) &= \mathbb{E}_{Q_\phi(z|x)} [\log P_\psi(z) + \log P_\theta(X | z)] + H[Q_\phi(z | X)] \\ &= \mathbb{E}_{Q_\phi(z|x)} \left[\log P(F_\psi^{-1}(z)) - \int_{t_0}^{t_1} \text{Tr} \left(\frac{\partial f_\psi}{\partial w(t)} \right) dt \right. \\ &\quad \left. + \sum_{x \in X} \left(\log P(G_\theta^{-1}(x; z)) - \int_{t_0}^{t_1} \text{Tr} \left(\frac{\partial g_\theta}{\partial y(t)} \right) dt \right) \right] \\ &\quad + H[Q_\phi(z | X)] \end{aligned} \quad (75)$$

can be interpretes in 3 parts:

1. Prior: $\mathcal{L}_{\text{prior}}(X; \psi, \phi) \triangleq \mathbb{E}_{Q_\phi(z|x)} [\log P_\psi(z)]$, use reparametrization to MC-sample:

$$\mathbb{E}_{Q_\phi(z|x)} [\log P_\psi(z)] \approx \frac{1}{L} \sum_{l=1}^L \log P_\psi(\mu + \epsilon_l \odot \sigma) \quad (76)$$

2. Recon. ld.: $\mathcal{L}_{\text{recon}}(X; \theta, \phi) \triangleq \mathbb{E}_{Q_\phi(z|x)} [\log P_\theta(X | z)]$, 依然使用 MC 采样估计.

3. Posterior Entropy: $\mathcal{L}_{\text{ent}}(X; \phi) \triangleq H[Q_\phi(z | X)]$, has form

$$H[Q_\phi(z | X)] = \frac{d}{2}(1 + \ln(2\pi)) + \sum_{i=1}^d \ln \sigma_i \quad (77)$$

9 FFJORD

9.1 CNF

use some base dist. $\mathbf{z}_0 \sim p_{z_0}(\mathbf{z}_0)$, 通过含时 ODE 得到要建模的分布

$$\mathbf{z}(t_0) = \mathbf{z}_0 \quad (78)$$

$$\frac{\partial \mathbf{z}}{\partial t} = f(\mathbf{z}(t), t; \theta) \quad (79)$$

log-pdf 的方程 (*instantaneous change of variables form.*)

$$\frac{\partial \log p(\mathbf{z}(t))}{\partial t} = -\text{Tr}\left(\frac{\partial f}{\partial \mathbf{z}(t)}\right) \quad (80)$$

$$\log p(\mathbf{z}(t_1)) = \log p(\mathbf{z}(t_0)) - \int_{t_0}^{t_1} \text{Tr}\left(\frac{\partial f}{\partial \mathbf{z}(t)}\right) dt \quad (81)$$

$$\underbrace{\begin{bmatrix} \mathbf{z}_0 \\ \log p(\mathbf{x}) - \log p_{z_0}(\mathbf{z}_0) \end{bmatrix}}_{\text{solutions}} = \underbrace{\int_{t_0}^{t_1} \begin{bmatrix} f(\mathbf{z}(t), t; \theta) \\ -\text{Tr}\left(\frac{\partial f}{\partial \mathbf{z}(t)}\right) \end{bmatrix} dt}_{\text{dynamics}}, \underbrace{\begin{bmatrix} \mathbf{z}(t_1) \\ \log p(\mathbf{x}) - \log p(\mathbf{z}(t_1)) \end{bmatrix}}_{\text{initial values}} = \begin{bmatrix} \mathbf{x} \\ 0 \end{bmatrix} \quad (82)$$

9.2 Backpropagation through ODE Solutions with Adjoint Method

Problem: calc. deriv. based on loss func.

$$L(\mathbf{z}(t_1)) = L\left(\int_{t_0}^{t_1} f(\mathbf{z}(t), t; \theta) dt\right) \quad (83)$$

Pontryagin(1962) 证明

$$\frac{dL}{d\theta} = - \int_{t_1}^{t_0} \left(\frac{\partial L}{\partial \mathbf{z}(t)}\right)^T \frac{\partial f(\mathbf{z}(t), t; \theta)}{\partial \theta} dt \quad (84)$$

值 $-\partial L / \partial \mathbf{z}(t)$ 称为 ODE 的伴随状态 (adjoint state). 使用一个 black-box ODE solver 来计算 $\mathbf{z}(t_1)$, 再用初值 $\partial L / \partial \mathbf{z}(t_1)$ 送进这个 ODE solver 来计算 (84)

9.3 Unbiased Linear-Time Log-Density Estimation

Hutchinson Estimator:

$$\text{Tr}(A) = E_{p(\epsilon)} [\epsilon^T A \epsilon] \quad (85)$$

holds if $\mathbb{E}[\epsilon] = 0$, $\text{Cov}\epsilon = I$ to avoid randomness, fix noise at each round of solving ODE

$$\begin{aligned} \log p(\mathbf{z}(t_1)) &= \log p(\mathbf{z}(t_0)) - \int_{t_0}^{t_1} \text{Tr}\left(\frac{\partial f}{\partial \mathbf{z}(t)}\right) dt \\ &= \log p(\mathbf{z}(t_0)) - \int_{t_0}^{t_1} \mathbb{E}_{p(\epsilon)} \left[\epsilon^T \frac{\partial f}{\partial \mathbf{z}(t)} \epsilon \right] dt \\ &= \log p(\mathbf{z}(t_0)) - \mathbb{E}_{p(\epsilon)} \left[\int_{t_0}^{t_1} \epsilon^T \frac{\partial f}{\partial \mathbf{z}(t)} \epsilon dt \right] \end{aligned} \quad (86)$$

噪声分布可以选为高斯分布/Rademacher 分布² 并且向量和 Jacobian 的乘积, i.e. $\epsilon \frac{\partial f}{\partial \mathbf{z}(t)}$, 可以快速算出 (通过 auto-diff)

Trick: Bottleneck width H to reduce variance of estimator.

Algorithm 1 Unbiased stochastic log-density estimation using the FFJORD model

Require: dynamics f_θ , start time t_0 , stop time t_1 , minibatch of samples \mathbf{x} .

```

 $\epsilon \leftarrow \text{sample\_unit\_variance}(\mathbf{x}.\text{shape})$                                 ▷ Sample  $\epsilon$  outside of the integral
function  $f_{aug}([\mathbf{z}_t, \log p_t], t)$ :                                         ▷ Augment  $f$  with log-density dynamics.
     $f_t \leftarrow f_\theta(\mathbf{z}(t), t)$                                               ▷ Evaluate dynamics
     $g \leftarrow \epsilon^T \frac{\partial f}{\partial \mathbf{z}}|_{\mathbf{z}(t)}$                          ▷ Compute vector-Jacobian product with automatic differentiation
     $\tilde{\text{Tr}} = \text{matrix\_multiply}(g, \epsilon)$                                ▷ Unbiased estimate of  $\text{Tr}(\frac{\partial f}{\partial \mathbf{z}})$  with  $\epsilon^T \frac{\partial f}{\partial \mathbf{z}} \epsilon$ 
    return  $[f_t, -\tilde{\text{Tr}}]$                                                  ▷ Concatenate dynamics of state and log-density
end function
 $[\mathbf{z}, \Delta_{logp}] \leftarrow \text{odeint}(f_{aug}, [\mathbf{x}, \vec{0}], t_0, t_1)$    ▷ Solve the ODE, ie.  $\int_{t_0}^{t_1} f_{aug}([\mathbf{z}(t), \log p(\mathbf{z}(t))], t) dt$ 
 $\log \hat{p}(\mathbf{x}) \leftarrow \log p_{\mathbf{z}_0}(\mathbf{z}) - \Delta_{logp}$                            ▷ Add change in log-density
return  $\log \hat{p}(\mathbf{x})$ 
```

10 Dequantization to Learn Discrete Distribution

为了近似一个离散空间上的 pd., 需要通过在数据点上加入噪声, 使用“去量化”技巧 (dequantization). 可变性更好的 noise \Rightarrow 更紧的下界 \Rightarrow learned noise?.

Theorem 加入合适的噪声后的连续随机变量的 ld(likelihood) 是对应离散随机变量 ld 的下界.

10.1 Dequantization as Latent Variable Model

$$P_{model}(x) = \int P_\theta(x|v)p(v)dv, \quad (88)$$

$$\text{where } P_\theta(x|v) = \mathbb{1}[v \in B_\theta(x)] \quad (89)$$

称 $P_\theta(x|v)$ 是量化子 (quantizer). 不同的量化子导致了不同的去量化方法. Half-infinite dequant. for bin. var.: $B(x) = \{x \cdot u | u \in \mathbb{R}_+^D\}, x \in -1, 1$; Hypercube dequant. for grid var.(images etc.): $B(x) = \{x + u | u \in [0, 1]^D\}$

上述积分难以计算, 引入去量化子 $q_\phi(v|x)$, 注意它具有不重叠的紧支撑集, 为此标记 $u = v + x$

² 在 $\{-1, 1\}$ 上均匀分布的离散分布

$$f(k) = \begin{cases} 1/2 & \text{if } k = -1 \\ 1/2 & \text{if } k = +1 \\ 0 & \text{otherwise} \end{cases} \quad (87)$$

$$P_{model}(x) = \int \frac{q_\phi(u|x) P_\theta(x|v)p(v)}{q_\phi(u|x)} dv \quad (90)$$

$$= \mathbb{E}_{u \sim q_\phi(u|x)} \left[\frac{P_\theta(x|v)p(v)}{q_\phi(u|x)} \right] \quad (91)$$

现有的方法常常使用格点积分作为离散和连续模型的区分

$$P(x) = \int_{[0,1)^D} p(x+u) du \quad (92)$$

下面提出三种 dequant. : i) variational inference, ii) weighted importance sampling, iii) variational Renyi approx.

10.2 Variational Dequantization

根据 Jensen 不等式, 得到 lld 的变分代理函数

$$\log P_{model}(x) \geq \mathbb{E}_{u \sim q_\phi(u|x)} \left[\log \frac{P_\theta(x|v)p_\theta(v)}{q_\phi(u|x)} \right] \quad (93)$$

注意去量化子要有紧支撑集, 所以对其输出进行 sigmoid; 并且进而有

$$\log P_{model}(x) \geq \mathbb{E}_{u \sim q_\phi(u|x)} [\log p_\theta(v)] + \mathbb{H}[q_\phi] \quad (94)$$

熵一项防止了概率分布退化到离散点上的 delta-peak, 从而推出了变分去量化 (vi dequant.)

10.3 Importance-Weighted Dequantization

除此之外, 还可以把去量化分布看作 proposal dist., 用采样多次替代 Jensen 不等式:

$$\log P_{model}(x) \geq \log \left[\frac{1}{K} \sum_{k \in [K]} \frac{P_\theta(x|v_k)p_\theta(v_k)}{q_\phi(u_k|x)} \right] \quad (95)$$

若提案分布限制于紧支撑集上, 则有

$$\log P_{model}(x) \geq \log \left[\frac{1}{K} \sum_{k \in [K]} \frac{p_\theta(v_k)}{q_\phi(u_k|x)} \right] \quad (96)$$

$$= \log [w_k(x)] \quad (97)$$

$$\text{where } w_k(x) \triangleq \frac{p_\theta(v_k)}{q_\phi(u_k|x)} \quad (98)$$

若 $K \rightarrow \infty$, 则取等号, 否则给出了 lld 的一个下界 (*iw-bound*), 故给出了 vi 界的更好估计 \Rightarrow iw-dequant.

10.4 Renyi Dequantization

vi/iw-去量化都可以看作变分 Renyi 去量化的特例. lld 可以用 Renyi Divergence 提供下界

$$\log P_{model}(x) \geq \frac{1}{1-\alpha} \log \left[\frac{1}{K} \sum_{k \in [K]} \left(\frac{P_\theta(x|v_k)p_\theta(v_k)}{q_\phi(u_k|x)} \right)^{1-\alpha} \right] \quad (99)$$

$$\text{where } \alpha \in [0, 1) \quad (100)$$

vi-bound $\alpha \rightarrow 1$, iw-bound $\alpha = 0$. [Li & Turner, 2016] 考虑小于 0 的 α , 这可能在低采样数时提供更紧的下界 (当 $K \rightarrow \infty$, 实际上提供了一个上界). 令 $\alpha = -\infty$, 得到 VR-max(variational Renyi max-approximation), 一个 iw-bound 的快速估计

$$\log P_{model}(x) \approx \log \max_k w_k(x) \quad (101)$$

Detail 用 Cholesky 分解计算协方差矩阵 $\Lambda = \Gamma \Gamma^T$, Γ^T 可学习.

10.5 Dequantization Distribution

Uniform Dequant. $q_\phi(u|x)$ 是 uniform in $\mathcal{B}(x)$

Gaussian Dequant. 更具表达力的是条件 logit-正态分布 (cond. logit-normal dist.)

$$q_\phi(u|x) = \text{sigmoid}(\mathcal{N}(\mu_\phi(x), \Sigma_\phi(x))) \quad (102)$$

Flow-based Dequant.

$$q_\phi(u|x) = q_\phi(\varepsilon = f_\phi(\text{sigmoid}^{-1}(u); x)|x) \det \mathbf{J} \quad (103)$$

由基分布 $q_\phi(\varepsilon|x)$ 和流双射 $f \in \mathbb{R}^D \rightarrow \mathbb{R}^D$ 组成. 这里的基分布采用对角高斯分布, 以及两种双射: coupling layer/flow/bipartite 和 autoregressive.

Bipartite Dequant. (Dinh et al., 2017) 使用流模型的耦合层:

$$\begin{cases} (\text{copy}) u_1 = \varepsilon + 1 \\ (\text{affine}) u_2 = \varepsilon_2 \odot s_\phi(\varepsilon_1; x) + t_\phi(\varepsilon_1) \end{cases} \quad (104)$$

$$(105)$$

为了保证所有分量都被变换, 应用另一个更改了 copy 层位置的耦合层.

Autoregressive Dequant./ARD (Kingma et al., 2016) 使用一个自回归模型

$$[m, s] = ARM_\phi(\varepsilon, h) \quad (106)$$

$$u = s \odot \varepsilon + m \quad (107)$$

其中 h 是上下文变量, 基于条件变量 x , 通过网络 s 计算出来.

10.6 (Choice of) Continuous Distribution

可以按前一节那样任意地选择量化子 $p_\theta(v)$. 但是在训练中需要采样 $v \sim p_\theta(v)$, 故使用自回归模型是很慢的, 故只考虑对角协方差/正常协方差的高斯分布和二分 flow-based 模型.

11 DGI: Deep Graph Infomax

11.1 Backgrounds, Approach, Math

Target Learn a encoder $\mathcal{E} = \mathcal{E}(\mathbf{X}, \mathbf{A}) : \mathbb{R}^{N \times F} \times \mathbb{R}^{N \times N} \rightarrow \mathbb{R}^{N \times F}$.

Approach 最大化局部互信息. 使用 *Readout* 函数来获得全局图特征 $\vec{s} = \mathcal{R}(\mathcal{E}(\mathbf{X}, \mathbf{A}))$. 为了能够计算 MI, 引入判别器 $\mathcal{D} : \mathbb{R}^F \times \mathbb{R}^F \rightarrow \mathbb{R}$, $\mathcal{D}(\vec{h}_i, \vec{s})$ 代表了两个图的 (repr.) 相似度. 判别器的负样本通过把一个图和一个不同的图联系在一起组成. 对于多图场景 (ModelNet/Molecule Graphs) 这可以通过采样其他图得到; 对于单图情景 (Cora etc.), 必须定义一个 (随机) 损坏函数 $\mathcal{C} : \mathbb{R}^{N \times F} \times \mathbb{R}^{N \times N} \rightarrow \mathbb{R}^{M \times F} \times \mathbb{R}^{M \times M}$

为此, 使用 contrastive loss

$$\mathcal{L} = \frac{1}{N+M} \left(\sum_{i=1}^N \mathbb{E}_{(\mathbf{X}, \mathbf{A})} \left[\log \mathcal{D}(\vec{h}_i, \vec{s}) \right] + \sum_{j=1}^M \mathbb{E}_{(\tilde{\mathbf{X}}, \tilde{\mathbf{A}})} \left[\log \left(1 - \mathcal{D}(\vec{h}_j, \vec{s}) \right) \right] \right) \quad (108)$$

这个 Jensen-Shannon divergence 本质上是互信息的 estimator!

Lemma 1. $\{\mathbf{X}^{(k)}\}_{k \in [|X|]}$ 从 $p(\mathbf{X})$ 中取出的一系列节点表示, 且 $p(\mathbf{X}^{(k)}) = p(\mathbf{X}^{(k')}) \forall k, k'$, 并且 $\mathcal{R}(\odot)$ 是确定性 Readout 函数, $\vec{s}^{(k)} = \mathcal{R}(\mathbf{X}^{(k)})$, 具有边缘分布 $p(\vec{s})$. 则基于联合分布的最优分类器 $p(\mathbf{X}, \vec{s})$ 和边缘分布的乘积 $p(\mathbf{X})p(\vec{s})$ 的误差有上界 $\text{Err}^* = \frac{1}{2} \sum_{k=1}^{|X|} p(\vec{s}^{(k)})^2$. 当 \mathcal{R} 是单射时达到上界.

Corollary 1. 此后都假设 \mathcal{R} 是单射, 假设 \vec{s} 的状态不少于 $|\mathbf{X}|$, 则最优全局表示满足 $|\vec{s}^*| = |\mathbf{X}|$.

Theorem 1. $\vec{s}^* = \operatorname{argmax}_{\vec{s}} I(\mathbf{X}; \vec{s})$

Theorem 2. 令 $\mathbf{X}_i^{(k)} = \{\vec{x}_j\}_{j \in n(\mathbf{X}^{(k)}, i)}$, 是第 k 层图卷积的特征, $\vec{h}_i = \mathcal{E}(\mathbf{X}_i^{(k)})$, 假设 $|\mathbf{X}_i| = |\mathbf{X}| = |\vec{s}| \geq |\vec{h}_i|$, 则最小化 $p(\vec{h}_i, \vec{s})$ and $p(\vec{h}_i)p(\vec{s})$ 的 \vec{h}_i 也让 MI 最大化.

11.2 Algorithm

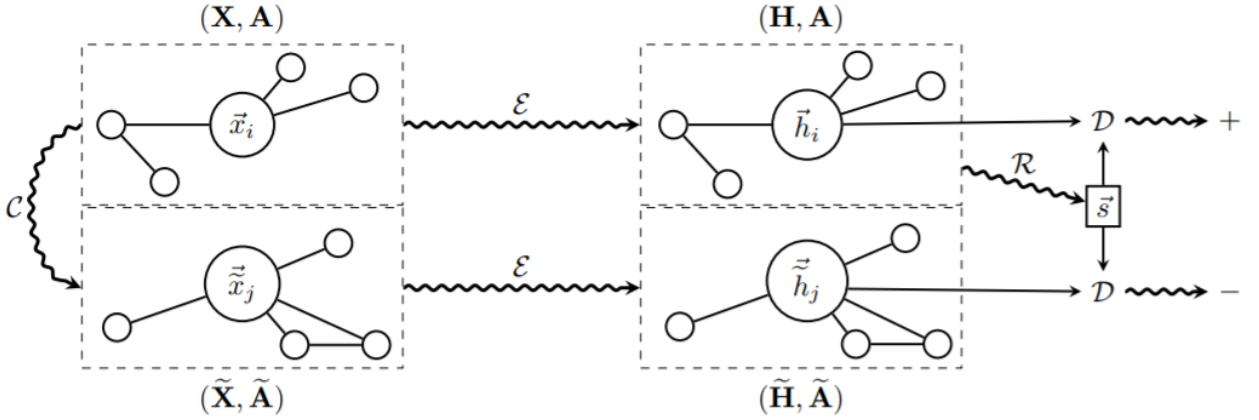


Figure 1: A high-level overview of Deep Graph Infomax. Refer to Section 3.4 for more details.

1. 从损坏函数中采样 $(\tilde{\mathbf{X}}, \tilde{\mathbf{A}}) \sim \mathcal{C}(\mathbf{X}, \mathbf{A})$
2. 获得正负样本的 patch node-repr., $\mathbf{H} = \mathcal{E}(\mathbf{X}, \mathbf{A}) = \{\vec{h}_1, \vec{h}_2, \dots, \vec{h}_N\}$, $\tilde{\mathbf{H}} = \mathcal{E}(\tilde{\mathbf{X}}, \tilde{\mathbf{A}}) = \{\tilde{h}_1, \tilde{h}_2, \dots, \tilde{h}_M\}$
3. 获得全局特征表示 $\vec{s} = \mathcal{R}(\mathbf{H})$.
4. 根据方程 (108) 更新 $\mathcal{R}, \mathcal{D}, \mathcal{E}$ 参数.

Details 使用 PReLU, 迁移学习任务上 (transductive, Cora, Citeseer, PubMed) 使用 GCN

$$\mathcal{E}(\mathbf{X}, \mathbf{A}) = \sigma \left(\hat{\mathbf{D}}^{-\frac{1}{2}} \hat{\mathbf{A}} \hat{\mathbf{D}}^{-\frac{1}{2}} \mathbf{X} \Theta \right) \quad (109)$$

在推断任务上 (inductive, Reddit) 使用 mean-aggr 和 GraphSAGE-GCN

$$\text{MP}(\mathbf{X}, \mathbf{A}) = \hat{\mathbf{D}}^{-1} \hat{\mathbf{A}} \mathbf{X} \Theta \quad (110)$$

在多图任务上 (PPI) 使用三层带有 dense skip conn. 的 mean-pooling 层

$$\begin{aligned} \mathbf{H}_1 &= \sigma(\text{MP}_1(\mathbf{X}, \mathbf{A})) \\ \mathbf{H}_2 &= \sigma(\text{MP}_2(\mathbf{H}_1 + \mathbf{X} \mathbf{W}_{\text{skip}}, \mathbf{A})) \\ \mathcal{E}(\mathbf{X}, \mathbf{A}) &= \sigma(\text{MP}_3(\mathbf{H}_2 + \mathbf{H}_1 + \mathbf{X} \mathbf{W}_{\text{skip}}, \mathbf{A})) \end{aligned} \quad (111)$$

在 Readout 函数上使用简单的 graph-mean-aggr

$$\mathcal{R}(\mathbf{H}) = \sigma \left(\frac{1}{N} \sum_{i=1}^N \vec{h}_i \right) \quad (112)$$

在判别器上使用简单的双线性打分函数

$$\mathcal{D}(\vec{h}_i, \vec{s}) = \sigma(\vec{h}_i^T \mathbf{W} \vec{s}) \quad (113)$$

12 GraphSAGE: Inductive Representation Learning on Graph

12.1 Embedding Generation/FP

Idea 在 k-hops 上逐层做 aggr.! Weisfeiler-Lehman 图同构检验的连续推广.

Algorithm 1: GraphSAGE embedding generation (i.e., forward propagation) algorithm

Input : Graph $\mathcal{G}(\mathcal{V}, \mathcal{E})$; input features $\{\mathbf{x}_v, \forall v \in \mathcal{V}\}$; depth K ; weight matrices $\mathbf{W}^k, \forall k \in \{1, \dots, K\}$; non-linearity σ ; differentiable aggregator functions $\text{AGGREGATE}_k, \forall k \in \{1, \dots, K\}$; neighborhood function $\mathcal{N} : v \rightarrow 2^{\mathcal{V}}$

Output: Vector representations \mathbf{z}_v for all $v \in \mathcal{V}$

```

1  $\mathbf{h}_v^0 \leftarrow \mathbf{x}_v, \forall v \in \mathcal{V}$  ;
2 for  $k = 1 \dots K$  do
3   for  $v \in \mathcal{V}$  do
4      $\mathbf{h}_{\mathcal{N}(v)}^k \leftarrow \text{AGGREGATE}_k(\{\mathbf{h}_u^{k-1}, \forall u \in \mathcal{N}(v)\})$ ;
5      $\mathbf{h}_v^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_v^{k-1}, \mathbf{h}_{\mathcal{N}(v)}^k))$ 
6   end
7    $\mathbf{h}_v^k \leftarrow \mathbf{h}_v^k / \|\mathbf{h}_v^k\|_2, \forall v \in \mathcal{V}$ 
8 end
9  $\mathbf{z}_v \leftarrow \mathbf{h}_v^K, \forall v \in \mathcal{V}$ 

```

使用固定大小的邻域函数 $\mathcal{N}(v)$ 以使用固定大小的权重 \mathbf{W} , 本工作使用邻域上的均匀采样. (?) 那么非均匀或者随时间变化的采样呢?) 为了进行图上的无监督学习, 引入 graph-loss(鼓励相邻节点具有相似的学到的表示)

$$J_{\mathcal{G}}(\mathbf{z}_u) = -\log(\sigma(\mathbf{z}_u^\top \mathbf{z}_v)) - Q \cdot \mathbb{E}_{v_n \sim P_n(v)} \log(\sigma(-\mathbf{z}_u^\top \mathbf{z}_{v_n})) \quad (114)$$

这里 σ 是 sigmoid 函数, v 是从 u 开始的固定长的随机游走序列上的节点, P_n 是负样本分布.

12.2 Aggregator Selection

Mean Aggregator 和 GCN 不同, mean-aggr. 的 repr. 和上一层的表示 concat, 可以看作 skip-conn., 大幅改善了性能.

$$\mathbf{h}_v^k \leftarrow \sigma(\mathbf{W} \cdot \text{MEAN}(\{\mathbf{h}_v^{k-1}\} \cup \{\mathbf{h}_u^{k-1}, \forall u \in \mathcal{N}(v)\})) \quad (115)$$

LSTM Aggregator 由于 LSTM 并不是内蕴轮换不变的, 所以使用结点的随机打乱作为输入.

Pooling Aggregator

$$\text{AGGREGATE}_k^{\text{pool}} = \max(\{\sigma(\mathbf{W}_{\text{pool}} \mathbf{h}_{u_i}^k + \mathbf{b}), \forall u_i \in \mathcal{N}(v)\}) \quad (116)$$

注意是 MLP+max-pooling.

13 SGC: Simplified Graph Convolution

回顾 GCN 中的图卷积, node-wise

$$\mathbf{h}_i^{(k)} \leftarrow \frac{1}{d_i + 1} \mathbf{h}_i^{(k-1)} + \sum_{j=1}^n \frac{a_{ij}}{\sqrt{(d_i + 1)(d_j + 1)}} \mathbf{h}_j^{(k-1)} \quad (117)$$

matrix-repr.

$$\begin{aligned} \mathbf{S} &= \tilde{\mathbf{D}}^{-\frac{1}{2}} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-\frac{1}{2}} \\ \overline{\mathbf{H}}^{(k)} &\leftarrow \mathbf{S} \mathbf{H}^{(k-1)} \end{aligned} \quad (118)$$

每一层的 feat.-trans. 和最后的分类器

$$\begin{aligned} \mathbf{H}^{(k)} &\leftarrow \text{ReLU} \left(\overline{\mathbf{H}}^{(k)} \Theta^{(k)} \right) \\ \hat{\mathbf{Y}}_{\text{GCN}} &= \text{softmax} \left(\mathbf{S} \mathbf{H}^{(K-1)} \Theta^{(K)} \right) \end{aligned} \quad (119)$$

SGC 直接在 k-hops 上聚合 (可以看作在 k-hop 连接图上聚集)

$$\hat{\mathbf{Y}}_{\text{SGC}} = \text{softmax} (\mathbf{S}^K \mathbf{X} \Theta) \quad (120)$$

这是一个凸优化问题, 可以通过二阶方法或者 SGD 来求解.

回顾 Ch

Net,

$$\mathbf{U} \hat{\mathbf{G}} \mathbf{U}^\top \mathbf{x} \approx \sum_{i=0}^k \theta_i \Delta^i \mathbf{x} = \mathbf{U} \left(\sum_{i=0}^k \theta_i^i \right) \mathbf{U}^\top \mathbf{x} \quad (121)$$

$$(122)$$

14 FastGCN

14.1 Method

回忆 GCN

$$\tilde{H}^{(l+1)} = \hat{A} H^{(l)} W^{(l)}, \quad H^{(l+1)} = \sigma \left(\tilde{H}^{(l+1)} \right), \quad l = 0, \dots, M-1, \quad L = \frac{1}{n} \sum_{i=1}^n g \left(H^{(M)}(i, :) \right) \quad (123)$$

写成泛函/积分变换的形式

$$\begin{aligned} \tilde{h}^{(l+1)}(v) &= \int \hat{A}(v, u) h^{(l)}(u) W^{(l)} dP(u), \quad h^{(l+1)}(v) = \sigma \left(\tilde{h}^{(l+1)}(v) \right), \quad l = 0, \dots, M-1 \\ L &= \mathbb{E}_{v \sim P} [g(h^{(M)}(v))] = \int g(h^{(M)}(v)) dP(v) \end{aligned} \quad (124)$$

把每个节点看作是 (连续 iid) 随机变量! 写成这种形式可以便利地使用 Monte-Carlo estimator 来估计, 每层使用 t_l 个采样来计算

$$\tilde{h}_{t_{l+1}}^{(l+1)}(v) := \frac{1}{t_l} \sum_{j=1}^{t_l} \hat{A} \left(v, u_j^{(l)} \right) h_{t_l}^{(l)} \left(u_j^{(l)} \right) W^{(l)}, \quad h_{t_{l+1}}^{(l+1)}(v) := \sigma \left(\tilde{h}_{t_{l+1}}^{(l+1)}(v) \right), \quad l = 0, \dots, M-1 \quad (125)$$

损失的估计 (这个估计是相容的 (以 1 概率收敛至真实值))

$$L_{t_0, t_1, \dots, t_M} := \frac{1}{t_M} \sum_{i=1}^{t_M} g \left(h_{t_M}^{(M)} \left(u_i^{(M)} \right) \right) \quad (126)$$

对于 mini-batch

$$L_{\text{batch}} = \frac{1}{t_M} \sum_{i=1}^{t_M} g \left(H^{(M)} \left(u_i^{(M)}, : \right) \right) \quad (127)$$

以及每一层的 FP

$$H^{(l+1)}(v, :) = \sigma \left(\frac{n}{t_l} \sum_{j=1}^{t_l} \hat{A} \left(v, u_j^{(l)} \right) H^{(l)} \left(u_j^{(l)}, : \right) W^{(l)} \right), \quad l = 0, \dots, M-1 \quad (128)$$

其中 n 是图节点数量, 作为正则化系数 (从矩阵形式到积分形式).

14.2 Variance Reduction

Summary Utilize Importance Sampling, Degree Weighted.

Use notations

	Function	Samples	Num. samples	
Layer $l+1$; random variable v	$\tilde{h}_{t_{l+1}}^{(l+1)}(v) \rightarrow y(v)$	$u_i^{(l+1)} \rightarrow v_i$	$t_{l+1} \rightarrow s$	(129)
Layer l ; random variable u	$h_{t_l}^{(l)}(u)W^{(l)} \rightarrow x(u)$	$u_j^{(l)} \rightarrow u_j$	$t_l \rightarrow t$	

consider layer repr.

$$G := \frac{1}{s} \sum_{i=1}^s y(v_i) = \frac{1}{s} \sum_{i=1}^s \left(\frac{1}{t} \sum_{j=1}^t \hat{A}(v_i, u_j) x(u_j) \right) \quad (130)$$

compute it's variance

$$\text{Var}\{G\} = R + \frac{1}{st} \iint \hat{A}(v, u)^2 x(u)^2 dP(u) dP(v) \quad (131)$$

$$\text{where } R = \frac{1}{s} \left(1 - \frac{1}{t} \right) \int e(v)^2 dP(v) - \frac{1}{s} \left(\int e(v) dP(v) \right)^2, e(v) = \int \hat{A}(v, u) x(u) dP(u) \quad (132)$$

第一项很难再优化, 由于取决于下一层的采样. 优化第二项, 引入新的采样分布, 则为了保持 G 期望不变

$$y_Q(v) := \frac{1}{t} \sum_{j=1}^t \hat{A}(v, u_j) x(u_j) \left(\frac{dP(u)}{dQ(u)} \Big|_{u_j} \right), \quad u_1, \dots, u_t \sim Q \quad (133)$$

此时

$$G_Q := \frac{1}{s} \sum_{i=1}^s y_Q(v_i) = \frac{1}{s} \sum_{i=1}^s \left(\frac{1}{t} \sum_{j=1}^t \hat{A}(v_i, u_j) x(u_j) \left(\frac{dP(u)}{dQ(u)} \Big|_{u_j} \right) \right) \quad (134)$$

Theorem 当

$$dQ(u) = \frac{b(u)|x(u)|dP(u)}{\int b(u)|x(u)|dP(u)} \quad \text{where} \quad b(u) = \left[\int \hat{A}(v, u)^2 dP(v) \right]^{\frac{1}{2}} \quad (135)$$

时, 方差最小, 为

$$\text{Var}\{G_Q\} = R + \frac{1}{st} \left[\int b(u)|x(u)|dP(u) \right]^2 \quad (136)$$

然而实际上 $|x(u)|$ 会变化且难以计算, 直接用 (... 那你论证半天为个啥...)

$$dQ(u) = \frac{b(u)^2 dP(u)}{\int b(u)^2 dP(u)} \quad (137)$$

MC 形式

$$q(u) = \|\hat{A}(:, u)\|^2 / \sum_{u' \in V} \left\| \hat{A}(:, u') \right\|^2, \quad u \in V \quad (138)$$

即和节点度数正比, 此时的每一层 FP 公式

$$H^{(l+1)}(v, :) = \sigma \left(\frac{1}{t_l} \sum_{j=1}^{t_l} \frac{\hat{A}(v, u_j^{(l)}) H^{(l)}(u_j^{(l)}, :) W^{(l)}}{q(u_j^{(l)})} \right), \quad u_j^{(l)} \sim q, \quad l = 0, \dots, M-1 \quad (139)$$

15 GWNN: Wavelet Transform on Graph

15.1 Supplementary Math: Real and Complex Wavelets

Function $\psi \in L^2(\mathbb{R})$ called **orthogonal wavelet**, if it could be used to define a orthogonal complete basis of Hilbert space $L^2(\mathbb{R})$. Given ψ , the basis are

$$\psi_{jk}(x) = 2^{\frac{j}{2}} \psi(2^j x - k) \quad (140)$$

under normal inner-product on $L^2(\mathbb{R})$, it's orthogonal

$$\begin{aligned} \langle \psi_{jk}, \psi_{lm} \rangle &= \int_{-\infty}^{\infty} \psi_{jk}(x) \overline{\psi_{lm}(x)} dx \\ &= \delta_{jl} \delta_{km} \end{aligned} \quad (141)$$

Integral Wavelet Transform

$$[W_\psi f](a, b) = \frac{1}{\sqrt{|a|}} \int_{-\infty}^{\infty} \overline{\psi\left(\frac{x-b}{a}\right)} f(x) dx \quad (142)$$

wavelet coefficient given by

$$c_{jk} = [W_\psi f](2^{-j}, k2^{-j}) \quad (143)$$

Meyer Wavelet in frequency-domain defined

$$\Psi(\omega) := \begin{cases} \frac{1}{\sqrt{2\pi}} \sin\left(\frac{\pi}{2}\nu\left(\frac{3|\omega|}{2\pi} - 1\right)\right) e^{j\omega/2} & \text{if } 2\pi/3 < |\omega| < 4\pi/3 \\ \frac{1}{\sqrt{2\pi}} \cos\left(\frac{\pi}{2}\nu\left(\frac{3|\omega|}{4\pi} - 1\right)\right) e^{j\omega/2} & \text{if } 4\pi/3 < |\omega| < 8\pi/3 \\ 0 & \text{otherwise} \end{cases} \quad (144)$$

where (standard impl.)

$$\nu(x) := \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } 0 < x < 1 \\ 1 & \text{if } x > 1 \end{cases} \quad (145)$$

it can also be

$$\nu(x) := \begin{cases} x^4(35 - 84x + 70x^2 - 20x^3) & \text{if } 0 < x < 1 \\ 0 & \text{otherwise.} \end{cases} \quad (146)$$

in time-domain a close form is obtained

$$\phi(t) = \begin{cases} \frac{2}{3} + \frac{4}{3\pi} & t = 0 \\ \frac{\sin(\frac{2\pi}{3}t) + \frac{4}{3}t \cos(\frac{4\pi}{3}t)}{\pi t - \frac{16\pi}{9}t^3} & \text{otherwise} \end{cases} \quad (147)$$

and $\psi(t) = \psi_1(t) + \psi_2(t)$,

$$\begin{aligned} \psi_1(t) &= \frac{\frac{4}{3\pi}(t-\frac{1}{2})\cos[\frac{2\pi}{3}(t-\frac{1}{2})] - \frac{1}{\pi}\sin[\frac{4\pi}{3}(t-\frac{1}{2})]}{(t-\frac{1}{2}) - \frac{16}{9}(t-\frac{1}{2})^3} \\ \psi_2(t) &= \frac{\frac{8}{3\pi}(t-\frac{1}{2})\cos[\frac{8\pi}{3}(t-\frac{1}{2})] + \frac{1}{\pi}\sin[\frac{4\pi}{3}(t-\frac{1}{2})]}{(t-\frac{1}{2}) - \frac{64}{9}(t-\frac{1}{2})^3} \end{aligned} \quad (148)$$

Mexican Hat Wavelet 1d-form Ricker Wavelet, 2nd deriv. of Gaussian dist.

$$\psi(t) = \frac{2}{\sqrt{3\sigma\pi^{1/4}}} \left(1 - \left(\frac{t}{\sigma}\right)^2\right) e^{-\frac{t^2}{2\sigma^2}} \quad (149)$$

2d-form Marr Wavelet

$$\psi(x, y) = \frac{1}{\pi\sigma^4} \left(1 - \frac{1}{2} \left(\frac{x^2 + y^2}{\sigma^2}\right)\right) e^{-\frac{x^2 + y^2}{2\sigma^2}} \quad (150)$$

Morlet Wavelet

$$\Psi_\sigma(t) = c_\sigma \pi^{-\frac{1}{4}} e^{-\frac{1}{2}t^2} (e^{i\sigma t} - \kappa_\sigma) \quad (151)$$

scale factor

$$c_\sigma = \left(1 + e^{-\sigma^2} - 2e^{-\frac{3}{4}\sigma^2}\right)^{-\frac{1}{2}} \quad (152)$$

15.2 Graph Wavelets

定义一系列图上的小波 $\psi_s = \{\psi_{si}\}$, ψ_{si} 代表以结点 i 为中心, 尺度为 s 的小波, 数学上可以写成

$$\psi_s = \mathbf{U} \mathbf{G}_s \mathbf{U}^\top \quad (153)$$

其中 \mathbf{U} 是拉普拉斯矩阵的特征向量, $\mathbf{G}_s = \text{diag}(g(s\lambda_1), \dots, g(s\lambda_n))$, $g(s\lambda_i) = e^{s\lambda_i}$ (... 就这? 这不是说 $\mathbf{G}_s = \exp(s\Lambda)$?) 图上的小波变换

$$\hat{\mathbf{x}} = \psi_s^{-1} \mathbf{x} \quad (154)$$

小波基的卷积

$$\mathbf{x} *_{\mathcal{G}} \mathbf{y} = \psi_s ((\psi_s^{-1} \mathbf{y}) \odot (\psi_s^{-1} \mathbf{x})) \quad (155)$$

15.3 GWNN

GWNN Layer

$$\mathbf{X}_{[:,j]}^{m+1} = h \left(\psi_s \sum_{i=1}^p \mathbf{F}_{i,j}^m \psi_s^{-1} \mathbf{X}_{[:,i]}^m \right) \quad j = 1, \dots, q \quad (156)$$

in node-wise favor

$$\mathbf{x}_j^{m+1} = h \left(\psi_s \sum_{i=1}^p \mathbf{F}_{i,j}^m \psi_s^{-1} \mathbf{x}_i^m \right) \quad j = 1, \dots, q \quad (157)$$

where \mathbf{F} is diagonal. On inductive missons(Cora etc.), 使用两层 (ReLU,softmax)GWNN. Parameters $O(npq)$, bad! Detach feat. trans. and graph conv.(as if GCN)

$$\mathbf{X}^{m+1} = h \left(\psi_s \mathbf{F}^m \psi_s^{-1} \mathbf{X}^m \mathbf{W} \right) \quad (158)$$

Advantages

1. 高效性: 小波基可以通过快速方法得到 (Chebyshev 估计,m 阶对应复杂度 $O(m|E|)$), 无需昂贵的 EVD).
2. 高稀疏性.
3. 局部化卷积.
4. 可变的邻域.

15.3.1 Details

1. \mathbf{F} 是一个对角阵 (特征向量的滤波器)
2. 只用了一个尺度 (严格地说是两个 $s, -s$), 核函数是 heat kernel: e^{-t}
3. 可以用pygsp包的内建函数来计算 Chebyshev 系数
 - pygsp.filters.approximations.compute_cheby_coeff(filter, order)
 - pygsp.filters.approximations.cheby_op(G, c, signal)
4. 源代码里用了一个 trick, 即在 $N \times N$ 单位阵上应用cheby_op(G, c, I) 来得到 ψ_s 的稀疏表示. 最后还用 L1 范数归一化.
5. Shapes: $\mathbf{X}^m, \mathbf{X}^{m'} \in \mathbb{R}^{N \times F}, \psi_s \in \mathbb{R}^{F \times N}, \psi_s^{-1} \in \mathbb{R}^{N \times F}, \mathbf{F} = \text{Diag}(\mathbf{f}) \in \mathbb{R}^{N \times N}$,

16 Graph Wavelets

16.1 经典小波变换/CWT

小波

$$\psi_{s,a}(x) = \frac{1}{s} \psi \left(\frac{x-a}{s} \right) \quad (159)$$

(经典) 小波变换/CWT

$$W_f(s, a) = \int_{-\infty}^{\infty} \frac{1}{s} \psi^* \left(\frac{x-a}{s} \right) f(x) dx \quad (160)$$

可逆, 若满足 admissibility cond.

$$\int_0^{\infty} \frac{|\hat{\psi}(\omega)|^2}{\omega} d\omega = C_{\psi} < \infty \quad (161)$$

逆变换/IWT

$$f(x) = \frac{1}{C_{\psi}} \int_0^{\infty} \int_{-\infty}^{\infty} W_f(s, a) \psi_{s,a}(x) \frac{das}{s} \quad (162)$$

定义算子

$$T^s f(a) = W_f(s, a) \quad (163)$$

有

$$\bar{\psi}_s(x) = \frac{1}{s} \psi^* \left(\frac{-x}{s} \right) \quad (164)$$

则有

$$\begin{aligned} (T^s f)(a) &= \int_{-\infty}^{\infty} \frac{1}{s} \psi^* \left(\frac{x-a}{s} \right) f(x) dx = \int_{-\infty}^{\infty} \bar{\psi}_s(a-x) f(x) dx \\ &= (\bar{\psi}_s \star f)(a) \end{aligned} \quad (165)$$

频域上有

$$\widehat{T^s f}(\omega) = \hat{\psi}_s(\omega) \hat{f}(\omega) \quad (166)$$

以及

$$\hat{\psi}_s(\omega) = \hat{\psi}^*(s\omega) \quad (167)$$

那么

$$(T^s f)(x) = \frac{1}{2\pi} \int_{-\infty}^{\infty} e^{i\omega x} \hat{\psi}^*(s\omega) \hat{f}(\omega) d\omega \quad (168)$$

16.2 谱小波变换/SGWT

SGWT 核 $g \Rightarrow T_g = g(\mathcal{L})$, 有频谱

$$\widehat{T_g f}(\ell) = g(\lambda_\ell) \hat{f}(\ell) \quad (169)$$

使用 IFT

$$(T_g f)(m) = \sum_{\ell=0}^{N-1} g(\lambda_\ell) \hat{f}(\ell) \chi_\ell(m) \quad (170)$$

局域化的图小波 $\psi_{t,n} = T_g^t \delta_n$, 展开得

$$\psi_{t,n}(m) = \sum_{\ell=0}^{N-1} g(t\lambda_\ell) \chi_\ell^*(n) \chi_\ell(m) \quad (171)$$

小波系数

$$W_f(t, n) = (T_g^t f)(n) = \sum_{\ell=0}^{N-1} g(t\lambda_\ell) \hat{f}(\ell) \chi_\ell(n) \quad (172)$$

16.2.1 Scaling Functions

小波都和第一特征向量 χ_0 正交, 并和特征值接近 0 的 eig-vec 几乎正交. 于是引入尺度函数, 类似地通过一个函数 $h : \mathbb{R}^+ \rightarrow \mathbb{R}$ 定义, 满足 $h(0) = 0, h(\infty) = 0, \phi_n = T_h \delta_n = h(\mathcal{L}) \delta_n$, 系数 $S_f(n) = \langle \phi_n, f \rangle$.

将会在之后看到, 当 $G(\lambda) = h(\lambda)^2 + \sum_{j=1}^J g(t_j \lambda)^2$ 有界且离开 0 时, 可以达到稳定近似.

16.3 SGWT 的性质

16.3.1 Inverse SGWT

Lemma 若 SGWT 核满足 admissibility cond.

$$\int_0^\infty \frac{g^2(x)}{x} dx = C_g < \infty \quad (173)$$

且 $g(0) = 0$, 则

$$\frac{1}{C_g} \sum_{n=1}^N \int_0^\infty W_f(t, n) \psi_{t,n}(m) \frac{dt}{t} = f^\#(m) \quad (174)$$

且

$$f = f^\# + \hat{f}(0)\chi_0 \quad (175)$$

16.3.2 局域性

Lemma 定义 $d_G(m, n)$ 为结点最短路径长度 (不考虑边权). 若 $d_G(m, n) > s$, $(\mathcal{L}^s)_{m,n} = 0$

Lemma Let $\psi_{t,n} = T_g^t \delta_n$ and $\tilde{\psi}_{t,n} = T_{\tilde{g}}^t \delta_n$ be the wavelets at scale t generated by the kernels g and \tilde{g} . If $|g(t\lambda) - \tilde{g}(t\lambda)| \leq M(t)$ for all $\lambda \in [0, \lambda_{N-1}]$, then $|\psi_{t,n}(m) - \tilde{\psi}_{t,n}(m)| \leq M(t)$ for each vertex m . Additionally, $\|\psi_{t,n} - \tilde{\psi}_{t,n}\|_2 \leq \sqrt{N}M(t)$

Lemma Let g be $K+1$ times continuously differentiable, satisfying $g(0) = 0, g^{(r)}(0) = 0$ for all $r < K$, and $g^{(K)}(0) = C \neq 0$. Assume that there is some $t' > 0$ such that $|g^{(K+1)}(\lambda)| \leq B$ for all $\lambda \in [0, t'\lambda_{N-1}]$. Then, for $\tilde{g}(t\lambda) = (C/K!)(t\lambda)^K$ we have

$$M(t) = \sup_{\lambda \in [0, \lambda_{N-1}]} |g(t\lambda) - \tilde{g}(t\lambda)| \leq t^{K+1} \frac{\lambda_{N-1}^{K+1}}{(K+1)!} B$$

for all $t < t'$

Theorem Let G be a weighted graph with Laplacian \mathcal{L} . Let g be a kernel satisfying the hypothesis of Lemma 5.4, with constants t' and B . Let m and n be vertices of G such that $d_G(m, n) > K$. Then there exist constants D and t'' , such that

$$\frac{\psi_{t,n}(m)}{\|\psi_{t,n}\|} \leq Dt$$

for all $t < \min(t', t'')$

16.3.3 Spectral Wavelet Frames

使用中必然使用 J 个 t 的离散采样, 导致 NJ 个小波和 N 个伸缩函数(尺度函数). 我们称一个在离散化的尺度上的小波为一个帧. 一些 Hilbert 空间上的向量组成的帧 $\Gamma_k \in \mathcal{H}$, 不等式

$$A\|f\|^2 \leq \sum_k |\langle f, \Gamma_k \rangle|^2 \leq B\|f\|^2$$

控制了数值稳定性.

Theorem Given a set of scales $\{t_j\}_{j=1}^J$, the set $F = \{\phi_n\}_{n=1}^N \cup \{\psi_{t_j, n}\}_{j=1}^J N_{n=1}$ forms a frame with bounds A, B given by

$$\begin{aligned} A &= \min_{\lambda \in [0, \lambda_{N-1}]} G(\lambda) \\ B &= \max_{\lambda \in [0, \lambda_{N-1}]} G(\lambda) \end{aligned}$$

where $G(\lambda) = h^2(\lambda) + \sum_j g(t_j \lambda)^2$

16.4 Fast SGWT Approximation by Polynomials

Lemma (多项式逼近的有效性) Let $\lambda_{\max} \geq \lambda_{N-1}$ be any upper bound on the spectrum of \mathcal{L} . For fixed $t > 0$, let $p(x)$ be a polynomial approximant of $g(tx)$ with L_∞ error $B = \sup_{x \in [0, \lambda_{\max}]} |g(tx) - p(x)|$. Then the approximate wavelet coefficients $\tilde{W}_f(t, n) = (p(\mathcal{L})f)_n$ satisfy

$$|W_f(t, n) - \tilde{W}_f(t, n)| \leq B\|f\|$$

获得这么一个估计在使用时往往需要知道一个特征值上界的估计 λ_{\max} , 但这是个很容易的问题, 只需要做一些矩阵-向量乘法即可, 比如 Arnoldi 迭代或者 Jacobi-Davidson 算法.

使用 Chebyshev 多项式逼近: 由数值分析得知 Chebyshev 时同阶多项式逼近性能最好的.

$$T_0(\lambda) = 1, T_1(\lambda) = \lambda, \quad (176)$$

$$T_j(\lambda) = 2\lambda T_{j-1}(\lambda) - T_{j-2}(\lambda) \quad (177)$$

$$T_n(x) = \cos(n \arccos(x)) \quad (178)$$

使用变换 $x = a(y + 1), a = \lambda_{\max}/2$ 来把 x 变换到 $[-1, 1]$ 上. 假设使用一系列离散化的尺度 t_n , 记偏移的 CP $\bar{T}(x) = T_k\left(\frac{x-a}{a}\right)$, 可写

$$g(t_n x) = \frac{1}{2}c_{n,0} + \sum_{k=1}^{\infty} c_{n,k} \bar{T}_k(x) \quad (179)$$

系数

$$c_{n,k} = \frac{2}{\pi} \int_0^\pi \cos(k\theta) g(t_n(a(\cos(\theta) + 1))) d\theta \quad (180)$$

(简单的数值积分即可, 计算快速). 对于任何尺度系 t_j , 截断级数到 M_j 项来逼近核函数 g , 我们

有估计

$$\tilde{W}_f(t_j, n) = \left(\frac{1}{2} c_{j,0} f + \sum_{k=1}^{M_j} c_{j,k} \bar{T}_k(\mathcal{L}) f \right)_n \quad (181)$$

$$\Rightarrow \tilde{W}_f(t_j, :) = \frac{1}{2} c_{j,0} f + \sum_{k=1}^{M_j} c_{j,k} \bar{T}_k(\mathcal{L}) f \quad (182)$$

$$\tilde{S}_f(n) = \left(\frac{1}{2} c_{0,0} f + \sum_{k=1}^{M_0} c_{0,k} \bar{T}_k(\mathcal{L}) f \right)_n \quad (183)$$

$$\Rightarrow \tilde{S}_f = \frac{1}{2} c_{0,0} f + \sum_{k=1}^{M_0} c_{0,k} \bar{T}_k(\mathcal{L}) f \quad (184)$$

$$\text{with efficient comp. of } \bar{T}_k(\mathcal{L}) f = \frac{2}{a} (\mathcal{L} - I) (\bar{T}_{k-1}(\mathcal{L}) f) - \bar{T}_{k-2}(\mathcal{L}) f \quad (185)$$

16.4.1 Fast Approximation of Adjoint

可以认为 $W : \mathbb{R}^N \rightarrow \mathbb{R}^{N(J+1)}$ 是一个线性变换, 且 $Wf = \left((T_h f)^T, (T_g^{t_1} f)^T, \dots, (T_g^{t_J} f)^T \right)^T$, 考虑其多项式估计 $\tilde{W} = \left((p_0(\mathcal{L}) f)^T, (p_1(\mathcal{L}) f)^T, \dots, (p_J(\mathcal{L}) f)^T \right)^T$, 这里展示其伴随算子的快速近似算法. 有

$$\begin{aligned} \langle \eta, Wf \rangle_{N(J+1)} &= \langle \eta_0, T_h f \rangle + \sum_{j=1}^J \langle \eta_j, T_g^{t_j} f \rangle_N \\ &= \langle T_h^* \eta_0, f \rangle + \left\langle \sum_{j=1}^J (T_g^{t_j})^* \eta_j, f \right\rangle_N = \left\langle T_h \eta_0 + \sum_{j=1}^J T_g^{t_j} \eta_j, f \right\rangle_N \end{aligned} \quad (186)$$

这表明

$$W^* \eta = T_h \eta_0 + \sum_{j=1}^J T_g^{t_j} \eta_j \quad (187)$$

为了计算逆变换 (伪逆 $L = (W^* W)^{-1} W^*$ 是差异 norm 最小的逆变换), 计算变换和其伴随算子的乘积

$$\tilde{W}^* \tilde{W} f = \sum_{j=0}^J p_j(\mathcal{L}) (p_j(\mathcal{L}) f) = \left(\sum_{j=0}^J (p_j(\mathcal{L}))^2 \right) f \quad (188)$$

记 $P(x) = \sum_{j=0}^J (p_j(x))^2 = \frac{1}{2} d_0 + \sum_{k=1}^{M^*} d_k \bar{T}_k(x)$, $M^* \max(M_j)$, 有公式

$$T_k(x) T_l(x) = \frac{1}{2} (T_{k+l}(x) + T_{|k-l|}(x)) \quad (189)$$

设 $c'_{j,k} = c_{j,k}$ for $k \geq 1$ and $c'_{j,0} = \frac{1}{2} c_{j,0}$ 有

$$d'_{j,k} = \begin{cases} \frac{1}{2} \left(c'_{j,0} 2 + \sum_{i=0}^{M_n} c'_{j,i} 2 \right) & \text{if } k = 0 \\ \frac{1}{2} \left(\sum_{i=0}^k c'_{j,i} c'_{j,k-i} + \sum_{i=0}^{M_j-k} c'_{j,i} c'_{j,k+i} + \sum_{i=k}^{M_j} c'_{j,i} c'_{j,i-k} \right) & \text{if } 0 < k \leq M_j \\ \frac{1}{2} \left(\sum_{i=k-M_j}^{M_j} c'_{j,i} c'_{j,k-i} \right) & \text{if } M_j < k \leq 2M_j \end{cases} \quad (190)$$

设

$$d_{n,0} = 2d'_{j,0} \text{ and } d_{j,k} = d'_{j,k} \text{ for } k \geq 1, \text{ and setting } d_k = \sum_{j=0}^J d_{j,k} \quad (191)$$

则有

$$\tilde{W}^* \tilde{W} f = P(\mathcal{L})f = \frac{1}{2}d_0 f + \sum_{k=1}^{M^*} d_k \bar{T}_k(\mathcal{L})f \quad (192)$$

16.4.2 Inverse Calculation

使用伪逆 $L = (W^*W)^{-1} W^*$, 给定小波系数 c , 可以通过方程

$$(W^*W) f = W^*c \quad (193)$$

计算原信号. 直接解是困难的, 使用快速共轭梯度法, 或者经典的帧算法 (frame algorithm).

16.5 Implementations and Details

A good example:

$$g(x; \alpha, \beta, x_1, x_2) = \begin{cases} x_1^\alpha x^{-\alpha} & \text{for } x < x_1 \\ s(x) & \text{for } x_1 \leq x \leq x_2 \\ x_2^\beta x^{-\beta} & \text{for } x > x_2 \end{cases} \quad (194)$$

其中 $s(x)$ 是三次样条. 伸缩函数 $h(x) = h(0) \exp\left(-\left(\frac{x}{0.6\lambda_{\min}}\right)^4\right)$, 尺度按从小到大对数线性间隔选取, $t_1 = x_2/\lambda_{\min}$, $t_I = x_2/\lambda_{\max}$, $\lambda_{\min} = \lambda_{\max}/K$

17 GMNN: Graph Markov Neural Network




Figure 1. Framework overview. Yellow and grey squares are labeled and unlabeled objects. Grey/white grids are attributes. Histograms are label distributions of objects. Orange triple circles are object representations. GMNN is trained by alternating between an E-step and an M-step. See Sec. 4.4 for the detailed explanation.

Idea 使用统计关系学习 (SRL) 建模

$$p(\mathbf{y}_V | \mathbf{x}_V) = \frac{1}{Z(\mathbf{x}_V)} \prod_{(n_i, n_j) \in E} \psi_{i,j}(\mathbf{y}_{n_i}, \mathbf{y}_{n_j}, \mathbf{x}_V) \quad (195)$$

而 GNN 模型则忽略 labels 之间的关系

$$p(\mathbf{y}_V | \mathbf{x}_V) = \prod_{n \in V} p(\mathbf{y}_n | \mathbf{x}_V) \quad (196)$$

具体上讲, GMNN 使用一个条件随机场 (CRF)+ 平均场近似 (mean-field approx.) 来建模, 并用 EM 算法来优化.

17.1 Pseudolikelihood Variational EM

优化 ELBO

$$\begin{aligned} \log p_\phi(\mathbf{y}_L | \mathbf{x}_V) &\geq \\ \mathbb{E}_{q_\theta(\mathbf{y}_U | \mathbf{x}_V)} [\log p_\phi(\mathbf{y}_L, \mathbf{y}_U | \mathbf{x}_V) - \log q_\theta(\mathbf{y}_U | \mathbf{x}_V)] \end{aligned} \quad (197)$$

这里 q_θ 是任意分布, 当

$$q_\theta(\mathbf{y}_U \mid \mathbf{x}_V) = p_\phi(\mathbf{y}_U \mid \mathbf{y}_L, \mathbf{x}_V) \quad (198)$$

取等号. 使用经典的 EM 算法来学习! 然而 p_ϕ 中的配分函数难以计算, 使用以下 psedo-ld

$$\begin{aligned} \ell_{PL}(\phi) &\triangleq \mathbb{E}_{q_\theta(\mathbf{y}_U \mid \mathbf{x}_V)} \left[\sum_{n \in V} \log p_\phi(\mathbf{y}_n \mid \mathbf{y}_{V \setminus n}, \mathbf{x}_V) \right] \\ &= \mathbb{E}_{q_\theta(\mathbf{y}_U \mid \mathbf{x}_V)} \left[\sum_{n \in V} \log p_\phi(\mathbf{y}_n \mid \mathbf{y}_{\text{NB}(n)}, \mathbf{x}_V) \right] \end{aligned} \quad (199)$$

以上伪似然函数广泛应用于 Markov 学习中.

17.2 Inference

这一步设计计算后验分布 $p_\phi(\mathbf{y}_U \mid \mathbf{y}_L, \mathbf{x}_V)$, 但这是困难的, 使用另一个变分分布来计算, 并使用平均场近似

$$q_\theta(\mathbf{y}_U \mid \mathbf{x}_V) = \prod_{n \in U} q_\theta(\mathbf{y}_n \mid \mathbf{x}_V) \quad (200)$$

使用一个 GNN 来参数化上述公式的每一项

$$q_\theta(\mathbf{y}_n \mid \mathbf{x}_V) = \text{MLP}[\text{Cat}(\mathbf{y}_n \mid \text{softmax}(W_\theta \mathbf{h}_{\theta,n}))] \quad (201)$$

根据平均场近似, 最优值为

$$\begin{aligned} \log q^*(\mathbf{y}_n \mid \mathbf{x}_V) &= \\ \mathbb{E}_{q_\theta(\mathbf{y}_{\text{NB}(n) \cap U} \mid \mathbf{x}_V)} [\log p_\phi(\mathbf{y}_n \mid \mathbf{y}_{\text{NB}(n)}, \mathbf{x}_V)] + \text{const.} \end{aligned} \quad (202)$$

使用 Monte-Carlo 估计

$$\begin{aligned} \mathbb{E}_{q_\theta(\mathbf{y}_{\text{NB}(n) \cap U} \mid \mathbf{x}_V)} [\log p_\phi(\mathbf{y}_n \mid \mathbf{y}_{\text{NB}(n)}, \mathbf{x}_V)] \\ \simeq \log p_\phi(\mathbf{y}_n \mid \hat{\mathbf{y}}_{\text{NB}(n)}, \mathbf{x}_V) \end{aligned} \quad (203)$$

其中 $\hat{\mathbf{y}}_{\text{NB}(n)} = \{\hat{\mathbf{y}}_{n'}\}_{n' \in \text{NB}(n)}$, 且对于任何 unlabeled neighbors, 使用采样的标签 $\hat{\mathbf{y}}_{n'} \sim q_\theta(\mathbf{y}_{n'} \mid \mathbf{x}_V)$, 实践中发现只取一个 (unlabeled) 样本几乎和取很多样本效果相当 (!), 效率考虑只取一个, 综上,

$$q^*(\mathbf{y}_n \mid \mathbf{x}_V) \approx p_\phi(\mathbf{y}_n \mid \hat{\mathbf{y}}_{\text{NB}(n)}, \mathbf{x}_V) \quad (204)$$

那么我们可以把后者作为 (最大化) 目标, 然后最小化 KL 散度

$$\text{KL}(p_\phi(\mathbf{y}_n \mid \hat{\mathbf{y}}_{\text{NB}(n)}, \mathbf{x}_V) \parallel q_\theta(\mathbf{y}_n \mid \mathbf{x}_V)) \quad (205)$$

进一步还是用并行更新策略, 独立的为每个 unlabeled node 优化

$$O_{\theta,U} = \sum_{n \in U} \mathbb{E}_{p_\phi(\mathbf{y}_n \mid \hat{\mathbf{y}}_{\text{NB}(n)}, \mathbf{x}_V)} [\log q_\theta(\mathbf{y}_n \mid \mathbf{x}_V)] \quad (206)$$

以及在 labeled node 上优化

$$O_{\theta,L} = \sum_{n \in L} \log q_\theta(\mathbf{y}_n \mid \mathbf{x}_V) \quad (207)$$

最终的 loss

$$O_\theta = O_{\theta,U} + O_{\theta,L} \quad (208)$$

Algorithm 1 Optimization Algorithm

Input: A graph G , some labeled objects (L, \mathbf{y}_L) .

Output: Object labels \mathbf{y}_U for unlabeled objects U .

Pre-train q_θ with \mathbf{y}_L according to Eq. (11).

while not converge **do**

□ M-Step: Learning Procedure

 Annotate unlabeled objects with q_θ .

 Denote the sampled labels as $\hat{\mathbf{y}}_U$.

 Set $\hat{\mathbf{y}}_V = (\mathbf{y}_L, \hat{\mathbf{y}}_U)$ and update p_ϕ with Eq. (14).

□ E-Step: Inference Procedure

 Annotate unlabeled objects with p_ϕ and $\hat{\mathbf{y}}_V$.

 Denote the predicted label distribution as $p_\phi(\mathbf{y}_U)$.

 Update q_θ with Eq. (10), (11) based on $p_\phi(\mathbf{y}_U), \mathbf{y}_L$.

end while

Classify each unlabeled object n based on $q_\theta(\mathbf{y}_n | \mathbf{x}_V)$.

17.3 Learning

直接使用 GNN 来建模, 而非使用势函数

$$p_\phi(\mathbf{y}_n | \mathbf{y}_{NB(n)}, \mathbf{x}_V) = \text{Cat}(\mathbf{y}_n | \text{softmax}(W_\phi \mathbf{h}_{\phi,n})) \quad (209)$$

还可以使用 SRL 中的 techniques, 同时把 $\mathbf{y}_{NB(n)}, \mathbf{x}_{NB(n)}$ 送到 GNN 中作为 in-feature. 最终的优化目标

$$O_\phi = \sum_{n \in V} \log p_\phi(\hat{\mathbf{y}}_n | \hat{\mathbf{y}}_{NB(n)}, \mathbf{x}_V) \quad (210)$$

17.4 Optimization

在有标签数据上预训练 q_θ , 再 EM. 最后用 q_θ 来预测 (往往比用 p_ϕ 准确率高)

	GCN [9]	Vanilla SGD	GraphSAGE [5]	FastGCN [1]	VR-GCN [2]
Time complexity	$O(L\ A\ _0 F + LNF^2)$	$O(d^L NF^2)$	$O(r^L NF^2)$	$O(rLN F^2)$	$O(L\ A\ _0 F + LNF^2 + r^L NF^2)$
Memory complexity	$O(LNF + LF^2)$	$O(bd^L F + LF^2)$	$O(br^L F + LF^2)$	$O(brLF + LF^2)$	$O(LNF + LF^2)$

18 ClusterGCN: Fast Deep & Large GCNs

18.1 Vanilla ClusterGCN: Cluster For Batch

GCN 需要整个 epoch 来更新一次梯度, 使用 mini-batch SGD 可能可以增加性能, 为此使用 batch-estimator

$$\frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \nabla \text{loss}(y_i, z_i^{(L)}) \quad (211)$$

估计 loss. 但这会增加一整个 epoch 的计算时间, SGD 导致了 node-repr. 聚合了 $O(d^L)$ 个邻居的信息, 导致 BP 复杂度很高. 为此定义 embedding utilization(嵌入效用), 为一个节点的表示在 BP 中被重复利用的次数. 在 GCN 中很高, 每层都为 d , 但是在 GraphSAGE/FastGCN 中是一个很低的常数, 由于 k-hops 很难重叠.

为此, 考虑到一个 batch 的 emb. util. 是其中的边数 $\|A_{\mathcal{B}, \mathcal{B}}\|_0$, 故提出想法: 每次取出边数最大的(导出)子图. 对于一个图 G , 有分割

$$\bar{G} = [G_1, \dots, G_c] = [\{\mathcal{V}_1, \mathcal{E}_1\}, \dots, \{\mathcal{V}_c, \mathcal{E}_c\}] \quad (212)$$

据此, 有

$$A = \bar{A} + \Delta = \begin{bmatrix} A_{11} & \cdots & A_{1c} \\ \vdots & \ddots & \vdots \\ A_{c1} & \cdots & A_{cc} \end{bmatrix} \quad (213)$$

以及

$$\bar{A} = \begin{bmatrix} A_{11} & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & A_{cc} \end{bmatrix}, \Delta = \begin{bmatrix} 0 & \cdots & A_{1c} \\ \vdots & \ddots & \vdots \\ A_{c1} & \cdots & 0 \end{bmatrix} \quad (214)$$

令 \bar{A}' 是正则化后的邻接矩阵, FP 公式变得对于 cluster 分离

$$\begin{aligned} Z^{(L)} &= \bar{A}' \sigma(\bar{A}' \sigma(\dots \sigma(\bar{A}' X W^{(0)}) W^{(1)}) \dots) W^{(L-1)} \\ &= \begin{bmatrix} \bar{A}'_{11} \sigma(\bar{A}'_{11} \sigma(\dots \sigma(\bar{A}'_{11} X_1 W^{(0)}) W^{(1)}) \dots) W^{(L-1)} \\ \vdots \\ \bar{A}'_{cc} \sigma(\bar{A}'_{cc} \sigma(\dots \sigma(\bar{A}'_{cc} X_c W^{(0)}) W^{(1)}) \dots) W^{(L-1)} \end{bmatrix} \end{aligned} \quad (215)$$

loss 同理

$$\mathcal{L}_{\bar{A}'} = \sum_t \frac{|\mathcal{V}_t|}{N} \mathcal{L}_{\bar{A}'_{tt}}, \mathcal{L}_{\bar{A}'_{tt}} = \frac{1}{|\mathcal{V}_t|} \sum_{i \in \mathcal{V}_t} \text{loss}(y_i, z_i^{(L)}) \quad (216)$$

使用图节点聚类方法来产生分割 (Metis or Graclus), 本作中使用的是 METIS 算法³

³

18.2 Stochastic Multiple Partitions

以上分割算法的问题：固定地排除了一些边；并且倾向于把相似的结点放在一起，可能引入 bias。解决方案：先分割出相对大量的聚类，再随机选取一些聚类并在一起作为 batch。加快收敛。

Algorithm 1: Cluster GCN

Input: Graph A , feature X , label Y ;

Output: Node representation \bar{X}

- 1 Partition graph nodes into c clusters $\mathcal{V}_1, \mathcal{V}_2, \dots, \mathcal{V}_c$ by METIS;
 - 2 **for** $iter = 1, \dots, max_iter$ **do**
 - 3 Randomly choose q clusters, t_1, \dots, t_q from \mathcal{V} without replacement;
 - 4 Form the subgraph \bar{G} with nodes $\bar{\mathcal{V}} = [\mathcal{V}_{t_1}, \mathcal{V}_{t_2}, \dots, \mathcal{V}_{t_q}]$ and links $A_{\bar{\mathcal{V}}, \bar{\mathcal{V}}}$;
 - 5 Compute $g \leftarrow \nabla \mathcal{L}_{A_{\bar{\mathcal{V}}, \bar{\mathcal{V}}}}$ (loss on the subgraph $A_{\bar{\mathcal{V}}, \bar{\mathcal{V}}}$) ;
 - 6 Conduct Adam update using gradient estimator g
 - 7 Output: $\{W_l\}_{l=1}^L$
-

18.3 Analysis of Deeper Networks

一个方法是增加 residual-links

$$X^{(l+1)} = \sigma(A' X^{(l)} W^{(l)}) + X^{(l)} \quad (217)$$

另一个想法是

$$X^{(l+1)} = \sigma((A' + I) X^{(l)} W^{(l)}) \quad (218)$$

(from Wikipedia) METIS is a software package for graph partitioning that implements various multilevel algorithms. METIS' multilevel approach has three phases and comes with several algorithms for each phase:

1. Coarsen the graph by generating a sequence of graphs G_0, G_1, \dots, G_N , where G_0 is the original graph and for each $0 \leq i \leq j \leq N$, the number of vertices in G_i is greater than the number of vertices in G_j .
2. Compute a partition of G_N
3. Project the partition back through the sequence in the order of G_N, \dots, G_0 , refining it with respect to each graph.

The final partition computed during the third phase (the refined partition projected onto G_0) is a partition of the original graph.

用于强调上一层的 embedding, 为了提供数值稳定性, 使用度正则化 (区别于 GCN 的对称正则化)

$$\tilde{A} = (D + I)^{-1}(A + I) \quad (219)$$

以及 FP 公式

$$X^{(l+1)} = \sigma \left((\tilde{A} + \lambda \operatorname{diag}(\tilde{A})) X^{(l)} W^{(l)} \right) \quad (220)$$

实验证明这提高了深层网络的性能.

19 GAT: Graph Attention Network

GAT 层:

1. feat. trans. $\mathbf{h}' = \mathbf{W}\mathbf{h}$
2. atten. coeff. $e_{ij} = a(\mathbf{h}'_i, \mathbf{h}'_j)$
3. atten. on neighbors $\boldsymbol{\alpha}_i = \operatorname{softmax}(\mathbf{e}_i), \alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k \in \mathcal{N}_i} \exp(e_{ik})}$, 本文中使用单层 MLP+concat
feat. 作为注意力层, 则有

$$\alpha_{ij} = \frac{\exp \left(\operatorname{LeakyReLU} \left(\vec{\mathbf{a}}^T [\mathbf{W}\vec{h}_i \| \mathbf{W}\vec{h}_j] \right) \right)}{\sum_{k \in \mathcal{N}_i} \exp \left(\operatorname{LeakyReLU} \left(\vec{\mathbf{a}}^T [\mathbf{W}\vec{h}_i \| \mathbf{W}\vec{h}_k] \right) \right)} \quad (221)$$

4. 进一步, 使用 multi-head atten.

$$\vec{h}'_i = \|\sum_{k=1}^K \alpha_{ij}^k \mathbf{W}^k \vec{h}_j\| \quad (222)$$

最终层则使用 mean-aggr 而非 concat

$$\vec{h}'_i = \sigma \left(\frac{1}{K} \sum_{k=1}^K \sum_{j \in \mathcal{N}_i} \alpha_{ij}^k \mathbf{W}^k \vec{h}_j \right) \quad (223)$$

5. trick: transductive task 上使用随机采样的固定大小的邻域结点

20 Note on Probabilistic Graphical Models

20.1 Bayesian Networks

Theorem (d-分离的完备性) 几乎所有 (在测度意义上) 的能被 BN 表征的概率分布 $P(\text{CSDs})$ 都满足: 若两节点 d-分离, 则它们条件独立.

Theorem (I-等价判定) 若两个 BN 有相同的骨架 (无向图基底) 和相同的 v-结构 ($X \rightarrow Z \leftarrow Y$)
朴素贝叶斯, 贝叶斯网络 (一个 DAG)

20.2 Undirected Networks

Definition 一个(或一些)r.v. D 的因子是一个函数 $\phi : \text{dom}(D) \rightarrow \mathbb{R}$. 并且定义因子的乘积, $\phi_1 : \text{dom}((X_i) \cup (Y_j)) \rightarrow \mathbb{R}, \phi_2 : \text{dom}((Y_j) \cup (Z_k)) \rightarrow \mathbb{R}, \psi(\mathbf{X}, \mathbf{Y}, \mathbf{Z}) = \phi_1(\mathbf{X}, \mathbf{Y}) \times \phi_2(\mathbf{Y}, \mathbf{Z})$.

Definition 一个被

$$\Phi = \{\phi_1(\mathbf{D}_1), \dots, \phi_K(\mathbf{D}_K)\}$$

参数化的 Gibbs 分布 P_Φ 满足

$$P_\Phi(X_1, \dots, X_n) = \frac{1}{Z} \tilde{P}_\Phi(X_1, \dots, X_n) \quad (224)$$

且

$$\tilde{P}_\Phi(X_1, \dots, X_n) = \phi_1(\mathbf{D}_1) \times \phi_2(\mathbf{D}_2) \times \dots \times \phi_m(\mathbf{D}_m) \quad (225)$$

其中配分函数

$$Z = \sum_{X_1, \dots, X_n} \tilde{P}_\Phi(X_1, \dots, X_n) \quad (226)$$

Definition MN 的约化(reduction) $\mathcal{H}[\mathbf{u}]$ 和 $P_\Phi[\mathbf{u}]$ 定义为在变量集合 \mathbf{U} 上取值后的分布/图. 且他们是一一对应的.

Definition \mathbf{X}, \mathbf{Y} 关于 \mathbf{Z} 分离, 若前二者之间没有不通过 \mathbf{Z} 的路径.

Theorem P 是 Gibbs 分布, factorize $MN\mathcal{H}$, 则后者是前者的 I-map(即独立关系包含前者).

Theorem (Hammersley-Clifford) \mathcal{H} 是 MN, 是前者结点上的正分布 P 的 I-map, 则 P 是 Gibbs 分布, 且 factorize $MN\mathcal{H}$.

Theorem 若 \mathbf{X}, \mathbf{Y} 关于 \mathbf{Z} 不分离, 那么 \mathbf{X}, \mathbf{Y} 关于 \mathbf{Z} 不独立.

类似的, 我们可以说在几乎所有分布上独立可以推出在图上分离.

Definition

$$\mathcal{I}_p(\mathcal{H}) = \{(X \perp Y \mid \mathcal{X} - \{X, Y\}) : X - Y \notin \mathcal{H}\} \quad (227)$$

是 pairwise-separation of \mathcal{H} ,

$$\mathcal{I}_\ell(\mathcal{H}) = \{(X \perp \mathcal{X} - \{X\} - \text{MB}_{\mathcal{H}}(X) \mid \text{MB}_{\mathcal{H}}(X)) : X \in \mathcal{X}\} \quad (228)$$

是 markov-blanket of \mathcal{H} .

Theorem 以下结论等价 1. $P \models \mathcal{I}_\ell(\mathcal{H})$. 2. $P \models \mathcal{I}_p(\mathcal{H})$. 3. $P \models \mathcal{I}(\mathcal{H})$.

Definition $\phi(\mathbf{D}) = \exp(-\epsilon(\mathbf{D}))$, $\epsilon(\mathbf{D})$ 是能量函数.

Definition 20.1 一个分布 P 是一个 log-linear 模型, 在 \mathcal{H} 上, 若它和以下参数关联:

1. a set of features $\mathcal{F} = \{f_1(\mathbf{D}_1), \dots, f_k(\mathbf{D}_k)\}$, where each \mathbf{D}_i is a complete subgraph in \mathcal{H} ,
2. a set of weights w_1, \dots, w_k such that

$$P(X_1, \dots, X_n) = \frac{1}{Z} \exp \left[- \sum_{i=1}^k w_i f_i(\mathbf{D}_i) \right]$$

Example • Ising Model: 二元 r.v. $X_i \in \{-1, +1\}$, $\epsilon_{i,j}(x_i, x_j) = w_{i,j} x_i x_j$, 有能量函数

$$P(\xi) = \frac{1}{Z} \exp \left(- \sum_{i < j} w_{i,j} x_i x_j - \sum_i u_i x_i \right) \quad (229)$$

-
- Boltzmann Dist.: 二元 r.v. $X_i \in \{0, 1\}$, 边上的能量如同 Ising 模型, 但每个随机变量都分配了 pdf sigmoid(z), $z = -\left(\sum_j w_{i,j} x_j\right) - w_i$
 - Metric CRF: 使用 CRF 来标注图节点, 有能量函数

$$E(x_1, \dots, x_n) = \sum_i \epsilon_i(x_i) + \sum_{(i,j) \in \mathcal{E}} \epsilon_{i,j}(x_i, x_j) \quad (230)$$

其中能量函数的取法导致了不同的模型, Ising Model:

$$\epsilon_{i,j}(x_i, x_j) = \begin{cases} 0 & x_i = x_j \\ \lambda_{i,j} & x_i \neq x_j \end{cases} = \delta_{x_i, x_j} \lambda_{i,j} \quad (231)$$

Potts Model 定义了结点的度规函数 μ (需要满足非负性, 自反性, 三角不等式) 用于能量函数, 并适用于多种标签的情况. 若一个度规满足前二者(非负性, 自反性), 则称其为 semi-metric/半度规的. CV 中常用的能量函数截断范数

$$\epsilon(x_i, x_j) = \min\left(c \|x_i - x_j\|_p, \text{dist}_{\max}\right) \quad (232)$$

Definition 20.2 令 $\ell(\xi) = \log P(\xi)$ Canonical energy on clique, 关于一个特定的赋值

$$\xi^* = (x_1^*, \dots, x_n^*)$$

$$\epsilon_D^*(d) = \sum_{Z \subset D} (-1)^{|D-Z|} \ell(d_Z, \xi_{-Z}^*) \quad (233)$$

Proposition 20.3 Let \mathcal{B} be a Bayesian network over \mathcal{X} and $\mathbf{E} = e$ an obseruation. Let $\mathbf{W} = \mathcal{X} - \mathbf{E}$. Then $P_{\mathcal{B}}(\mathbf{W} | e)$ is a Gibbs distribution defined by the factors $\Phi = \{\phi_{X_i}\}_{X_i \in \mathcal{X}}$, where

$$\phi_{X_i} = P_{\mathcal{B}}(X_i | \text{Pa}_{X_i}) [\mathbf{E} = e]$$

The partition function for this Gibbs distribution is $P(e)$

Definition 20.4 Moralized map for BNG 定义为一个同样节点的无向图 $M[\mathcal{G}]$, 其中一条边 (X, Y) 存在若在 \mathcal{G} 中有一条有向边连接, 或者他们是 moral 的(具有相同的子结点).

Proposition 20.5 For BN \mathcal{G} , $M[\mathcal{G}]$ 是极小 I-map.

Proposition 20.6 Let $\mathbf{X}, \mathbf{Y}, \mathbf{Z}$ be three disjoint sets of nodes in a Bayesian network \mathcal{G} . Let $\mathbf{U} = \mathbf{X} \cup \mathbf{Y} \cup \mathbf{Z}$, and let $\mathcal{G}' = \mathcal{G}^+[\mathbf{U}]$ be the induced Bayesian network over $\mathbf{U} \cup \text{Ancestors}_{\mathcal{G}}$. Let \mathcal{H} be the moralized graph $M[\mathcal{G}']$. Then $d - \text{sep}_{\mathcal{G}}(\mathbf{X}; \mathbf{Y} | \mathbf{Z})$ if and only if $\text{sep}_{\mathcal{H}}(\mathbf{X}; \mathbf{Y} | \mathbf{Z})$

Theorem 20.7 若 $MN\mathcal{H}$ 是弦图, 则存在 BN \mathcal{G} such that $\mathcal{I}(\mathcal{H}) = \mathcal{I}(\mathcal{G})$

Definition 20.8 CRF 是一个无向图 \mathcal{H} , 节点为 $\mathbf{X} \cup \mathbf{Y}$, 带有因子 $\phi_1(\mathbf{D}_1), \dots, \phi_m(\mathbf{D}_m)$ such that each $\mathbf{D}_i \not\subseteq \mathbf{X}$, models dist. such as

$$\begin{aligned} P(\mathbf{Y} | \mathbf{X}) &= \frac{1}{Z(\mathbf{X})} \tilde{P}(\mathbf{Y}, \mathbf{X}) \\ \tilde{P}(\mathbf{Y}, \mathbf{X}) &= \prod_{i=1}^m \phi_i(\mathbf{D}_i) \\ Z(\mathbf{X}) &= \sum_Y \tilde{P}(\mathbf{Y}, \mathbf{X}) \end{aligned} \quad (234)$$

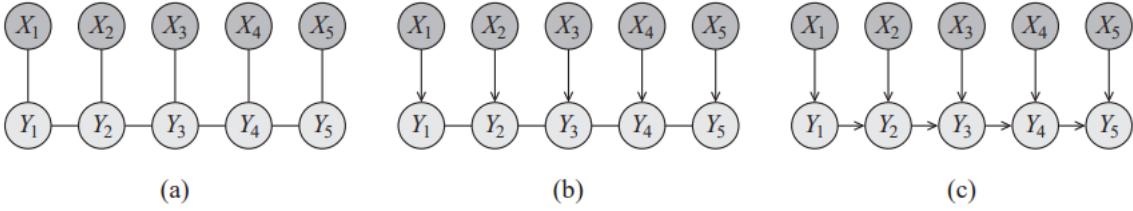


Figure 4.14 Different linear-chain graphical models: (a) a linear-chain-structured conditional random field, where the feature variables are denoted using grayed-out ovals; (b) a partially directed variant; (c) a fully directed, non-equivalent model. The X_i 's are assumed to be always observed when the network is used, and hence they are shown as darker gray.

Example 考虑只有一个 Y 的 CRF(朴素 Markov 模型), 能量函数

$$\phi_i(X_i, Y) = \exp\{w_i I\{X_i = 1, Y = 1\}\} \quad (235)$$

我们可以得到

$$P(Y = 1 | x_1, \dots, x_k) = \text{sigmoid}\left(w_0 + \sum_{i=1}^k w_i x_i\right) \quad (236)$$

a sigmoid-regression model!

20.3 Local Probabilistic Models | i.e. Specific Models Corresponds to Last 2 Sections

表式 \Rightarrow 复杂度极高!

确定性 CPD 由父节点的函数决定

$$P(x | \text{pa}_X) = \begin{cases} 1 & x = f(\text{pa}_X) \\ 0 & \text{otherwise} \end{cases} \quad (237)$$

树形 CPD: 类似于决策树, 但每个节点都 annotate 一个子结点上的分布

基于规则的 CPD

noisy-or CPD

$$\begin{aligned} P(y^0 | X_1, \dots, X_k) &= (1 - \lambda_0) \prod_{i:X_i=x_i^1} (1 - \lambda_i) \\ P(y^1 | X_1, \dots, X_k) &= 1 - \left[(1 - \lambda_0) \prod_{i:X_i=x_i^1} (1 - \lambda_i) \right] \end{aligned} \quad (238)$$

sigmoid CPD

$$P(y^1 | X_1, \dots, X_k) = \text{sigmoid}\left(w_0 + \sum_{i=1}^k w_i X_i\right) \quad (239)$$

Gaussian CPD

$$p(Y | x_1, \dots, x_k) = \mathcal{N}(\beta_0 + \beta_1 x_1 + \dots + \beta_k x_k; \sigma^2) \quad (240)$$

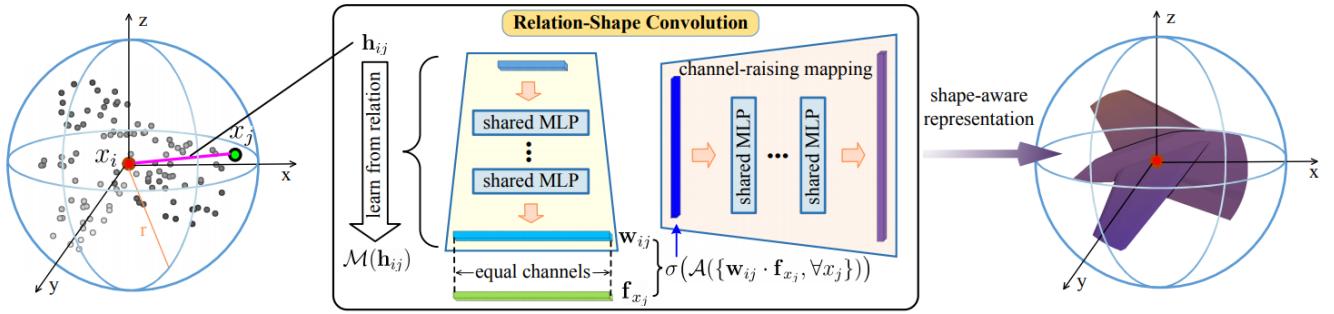


Figure 2. Overview of relation-shape convolution (RS-Conv). The key is to learn from relation. Specifically, the convolutional weight for x_j is converted to \mathbf{w}_{ij} , which learns a mapping \mathcal{M} (Eq. (2)) on predefined geometric relation vector \mathbf{h}_{ij} . In this way, the inductive convolutional representation $\sigma(\mathcal{A}(\{\mathbf{w}_{ij} \cdot \mathbf{f}_{x_j}, \forall x_j\}))$ (Eq. (3)) can expressively reason the spatial layout of points, resulting in discriminative shape awareness. As in image CNN [34], further channel-raising mapping is conducted for a more powerful shape-aware representation.

写成向量，则为

$$p(Y | \mathbf{x}) = \mathcal{N}(\beta_0 + \boldsymbol{\beta}^T \mathbf{x}; \sigma^2) \quad (241)$$

条件线性高斯模型 (CLG)

$$p(X | \mathbf{u}, \mathbf{y}) = \mathcal{N}\left(a_{\mathbf{u}, 0} + \sum_{i=1}^k a_{\mathbf{u}, i} y_i; \sigma_{\mathbf{u}}^2\right) \quad (242)$$

Definition 20.9 (*conditional Bayesian networks*) 条件贝叶斯网络 \mathcal{G} 是一个 DAG, 节点是分离的三个集合的并 $\mathbf{X} \cup \mathbf{Y} \cup \mathbf{Z}$, \mathbf{X} 没有父节点, 称作输入, \mathbf{Y} 称作输出, 且条件概率分布由链式法则定义

$$P_{\mathcal{B}}(\mathbf{Y}, \mathbf{Z} | \mathbf{X}) = \prod_{X \in Y \cup Z} P(X | \text{Pa}_X^{\mathcal{G}}) \quad (243)$$

. 边缘分布由求和给出

$$P_{\mathcal{B}}(\mathbf{Y} | \mathbf{X}) = \sum_{\mathbf{Z}} P_{\mathcal{B}}(\mathbf{Y}, \mathbf{Z} | \mathbf{X}) \quad (244)$$

20.4 Temporal Models

21 RSCNN(CVPR 19')

21.1 Architecture

Idea 使用空间卷积/spatial conv., 在球形邻域上.

一个广义卷积

$$\mathbf{f}_{\text{sub}} = \sigma(\mathcal{A}(\{\mathcal{T}(\mathbf{f}_{x_j}), \forall x_j\})), d_{ij} < r \forall x_j \in \mathcal{N}(x_i) \quad (245)$$

要想是这个卷积 permut.-invar., 函数 \mathcal{A}, \mathcal{T} 必须分别是对称的和 shared.

使用 shape-aware/geometric info 函数 \mathcal{M} (shared MLP 建模) 代替传统卷积

$$\mathcal{T}(\mathbf{f}_{x_j}) = \mathbf{w}_{ij} \cdot \mathbf{f}_{x_j} = \mathcal{M}(\mathbf{h}_{ij}) \cdot \mathbf{f}_{x_j} \quad (246)$$

则卷积形式变为

$$\mathbf{f}_{P_{\text{sub}}} = \sigma(\mathcal{A}(\{\mathcal{M}(\mathbf{h}_{ij}) \cdot \mathbf{f}_{x_j}, \forall x_j\})) \quad (247)$$

为了和 CNN 相对应, 使用 channel-raising MLP 来增多 channels.

最终, 这个卷积具有以下性质: permut. invar., 对于刚性变换的健壮性, shared weights, interacted point geometric.

21.2 Details & Implementation

使用 ReLU 激活函数, 使用 BN, \mathcal{M} 使用三层 MLP, aggr. f. 为 max-pooling. Low-level 几何表示 $\mathbf{h}_{ij} = [\mathbf{x}_i - \mathbf{x}_j, \mathbf{x}_i, \mathbf{x}_j]$, channel-raising 使用单层 MLP.

点云采样方面, 使用原点云的 FPS 采样. 使用 3-scale 邻域 (不同于 PN++ 的 MSG)

22 SimpleView(ICLR 21' Candidate)

22.1 Simple Review of Existing Protocols

数据增强: 包括抖动, y-轴随机旋转, 随即平移和缩放. ModelNet40 由于已经对齐, y-轴随机旋转会降低性能.

Model	PointNet(++)	DGCNN	RSCNN
All	随机旋转/平移	随机旋转/平移	

输入点数 PN(++) 使用 1024 个固定输入. PointCNN, RSCNN 使用每个 epoch 重采样的点.

Loss 大多数方法使用交叉熵, DGCNN 使用了平滑了的交叉熵 (label 经过平滑, 这个方法在所有结构上提高了性能)

模型选择 PN(++) 使用最终收敛的模型, DGCNN/RSCNN 使用测试集上的最好模型.

模型聚合 PN(++) 在 inference-time 把最终模型在不同旋转角度的输入上做判定 (10 次), 然后投票决定. RSCNN, DensePoint 在不同尺寸和角度的输入上判定 (300 次), 然后投票决定. DGCNN 完全没有投票.

比较性能, 本文提出的方法使用随机平移/缩放强化和 smooth-loss, 并且为了不利用任何测试集的信息, 使用 final model sel.

22.2 Model: SimpleView

Idea 使用多个视角的深度图像!

具体上, 使用六个 view(水平面四个, z 轴两个, 实验上这样性能最好), 并且在每张深度图上使用 ResNet18/4 骨架 (ResNet18, 滤波器数量为 1/4), concat 连接特征.

23 OT-Flow

23.1 Idea & Formulations

Formulation(based on FFJORD)

$$\partial_t \begin{bmatrix} z(\mathbf{x}, t) \\ \ell(\mathbf{x}, t) \end{bmatrix} = \begin{bmatrix} \mathbf{v}(z(\mathbf{x}, t), t; \boldsymbol{\theta}) \\ \text{tr}(\nabla \mathbf{v}(z(\mathbf{x}, t), t; \boldsymbol{\theta})) \end{bmatrix}, \quad \begin{bmatrix} z(\mathbf{x}, 0) \\ \ell(\mathbf{x}, 0) \end{bmatrix} = \begin{bmatrix} \mathbf{x} \\ 0 \end{bmatrix} \quad (248)$$

在 FFJORD 的基础

$$\min_{\boldsymbol{\theta}} \mathbb{E}_{\rho_0(\mathbf{x})} \left\{ C(\mathbf{x}, T) := \frac{1}{2} \|z(\mathbf{x}, T)\|^2 - \ell(\mathbf{x}, T) + \frac{d}{2} \log(2\pi) \right\} \quad (249)$$

上增加最优输运代价

$$L(\mathbf{x}, T) = \int_0^T \frac{1}{2} \|\mathbf{v}(z(\mathbf{x}, t), t)\|^2 dt \quad (250)$$

满足上两个 cost 的和最小化时, 则必定存在势函数

$$\mathbf{v}(\mathbf{x}, t; \boldsymbol{\theta}) = -\nabla \Phi(\mathbf{x}, t; \boldsymbol{\theta}) \quad (251)$$

并且满足 HJB 方程 (Hamilton-Jacobi-Bellman Eq.)

$$-\partial_t \Phi(\mathbf{x}, t) + \frac{1}{2} \|\nabla \Phi(z(\mathbf{x}, t), t)\|^2 = 0, \quad \Phi(\mathbf{x}, T) = G(\mathbf{x}) \quad (252)$$

故引入惩罚项

$$R(\mathbf{x}, T) = \int_0^T \left| \partial_t \Phi(z(\mathbf{x}, t), t) - \frac{1}{2} \|\nabla \Phi(z(\mathbf{x}, t), t)\|^2 \right| dt \quad (253)$$

本工作直接不建模梯度函数 \mathbf{v} , 而是直接建模势函数 Φ .

23.2 Parametrization of Model

势函数

$$\Phi(\mathbf{s}; \boldsymbol{\theta}) = \mathbf{w}^\top N(\mathbf{s}; \boldsymbol{\theta}_N) + \frac{1}{2} \mathbf{s}^\top (\mathbf{A}^\top \mathbf{A}) \mathbf{s} + \mathbf{b}^\top \mathbf{s} + c, \quad \text{where } \boldsymbol{\theta} = (\mathbf{w}, \boldsymbol{\theta}_N, \mathbf{A}, \mathbf{b}, c) \quad (254)$$

其中 N 是一个 NN(这里用的是一个简单的两层 ResNet), $\mathbf{A} \in \mathbb{R}^{r \times (d+1)}$, 且限制 rank $r = \max(10, d)$ 这里后面三项建模了一个二次势函数, 也即一个线性动力系统, NN 则建模了非线性部分.

ResNet 结构

$$\begin{aligned} \mathbf{u}_0 &= \sigma(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0) \\ N(\mathbf{s}; \boldsymbol{\theta}_N) &= \mathbf{u}_1 = \mathbf{u}_0 + h \sigma(\mathbf{K}_1 \mathbf{u}_0 + \mathbf{b}_1) \end{aligned} \quad (255)$$

梯度计算

$$\nabla_s \Phi(\mathbf{s}; \boldsymbol{\theta}) = \nabla_s N(\mathbf{s}; \boldsymbol{\theta}_N) \mathbf{w} + (\mathbf{A}^\top \mathbf{A}) \mathbf{s} + \mathbf{b} \quad (256)$$

Hessian Trace 计算

$$\text{tr}(\nabla^2 \Phi(\mathbf{s}; \boldsymbol{\theta})) = \text{tr}(\mathbf{E}^\top \nabla_s^2(N(\mathbf{s}; \boldsymbol{\theta}_N) \mathbf{w}) \mathbf{E}) + \text{tr}(\mathbf{E}^\top (\mathbf{A}^\top \mathbf{A}) \mathbf{E}) \quad (257)$$

后一项是 trivial 的, \mathbf{E} 是 $\mathbb{R}^{(d+1)}$ 标准正交基的前 d 项, ResNet 项可以得到一个闭形式

$$\begin{aligned}\text{tr}(\mathbf{E}^\top \nabla_{\mathbf{s}}^2(N(\mathbf{s}; \boldsymbol{\theta}_N) \mathbf{w}) \mathbf{E}) &= t_0 + ht_1, \quad \text{where} \\ t_0 &= (\sigma''(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0) \odot \mathbf{z}_1)^\top ((\mathbf{K}_0 \mathbf{E}) \odot (\mathbf{K}_0 \mathbf{E})) \mathbf{1} \\ t_1 &= (\sigma''(\mathbf{K}_1 \mathbf{u}_0 + \mathbf{b}_1) \odot \mathbf{w})^\top ((\mathbf{K}_1 \nabla_{\mathbf{s}} \mathbf{u}_0^\top) \odot (\mathbf{K}_1 \nabla_{\mathbf{s}} \mathbf{u}_0^\top)) \mathbf{1}\end{aligned}\tag{258}$$

第一层的 Hessian 计算复杂度 $O(md)$, 之后每多一层为 $O(m^2d)$, 总复杂度则为 $O(d)$

23.3 Exact Hessian of Multilayer NN

Exact Trace Computation Using (13) and the same E , we compute the trace in one forward pass through the layers. The trace of the first ResNet layer is

$$\begin{aligned}t_0 &= \text{tr}(\mathbf{E}^\top \nabla_{\mathbf{s}}(\mathbf{K}_0^\top \text{diag}(\sigma''(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0)) \mathbf{z}_1) \mathbf{E}) \\ &= \text{tr}(\mathbf{E}^\top \mathbf{K}_0^\top \text{diag}(\sigma''(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0) \odot \mathbf{z}_1) \mathbf{K}_0 \mathbf{E}) \\ &= (\sigma''(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0) \odot \mathbf{z}_1)^\top ((\mathbf{K}_0 \mathbf{E}) \odot (\mathbf{K}_0 \mathbf{E})) \mathbf{1}\end{aligned}$$

using the same notation as (14). For the last step, we used the diagonality of the middle matrix. Computing t_0 requires $\mathcal{O}(m \cdot d)$ FLOPS when first squaring the elements in the first d columns of \mathbf{K}_0 , then summing those columns, and finally one inner product.

To compute the trace of the entire ResNet, we continue with the remaining rows in (27) in reverse order to obtain

$$\text{tr}(\mathbf{E}^\top \nabla_{\mathbf{s}}^2(N(\mathbf{s}; \boldsymbol{\theta}_N) \mathbf{w}) \mathbf{E}) = t_0 + h \sum_{i=1}^M t_i$$

where t_i is computed as

$$\begin{aligned}t_i &= \text{tr}(\mathbf{J}_{i-1}^\top \nabla_{\mathbf{s}}(\mathbf{K}_i^\top \text{diag}(\sigma''(\mathbf{K}_i \mathbf{u}_{i-1}(\mathbf{s}) + \mathbf{b}_i)) \mathbf{z}_{i+1}) \mathbf{J}_{i-1}) \\ &= \text{tr}(\mathbf{J}_{i-1}^\top \mathbf{K}_i^\top \text{diag}(\sigma''(\mathbf{K}_i \mathbf{u}_{i-1} + \mathbf{b}_i) \odot \mathbf{z}_{i+1}) \mathbf{K}_i \mathbf{J}_{i-1}) \\ &= (\sigma''(\mathbf{K}_i \mathbf{u}_{i-1} + \mathbf{b}_i) \odot \mathbf{z}_{i+1})^\top ((\mathbf{K}_i \mathbf{J}_{i-1}) \odot (\mathbf{K}_i \mathbf{J}_{i-1})) \mathbf{1}\end{aligned}$$

Here, $\mathbf{J}_{i-1} = \nabla_{\mathbf{s}} \mathbf{u}_{i-1}^\top \in \mathbb{R}^{m \times d}$ is a Jacobian matrix, which can be updated and over-written in the forward pass at a computational cost of $\mathcal{O}(m^2 \cdot d)$ FLOPS. The J update follows:

$$\begin{aligned}\nabla_{\mathbf{s}} \mathbf{u}_i^\top &= \nabla_{\mathbf{s}} \mathbf{u}_{i-1} + h \sigma'(\mathbf{K}_i \mathbf{u}_{i-1} + \mathbf{b}_i) \mathbf{K}_i^\top \nabla_{\mathbf{s}} \mathbf{u}_{i-1} \\ J &\leftarrow J + h \sigma'(\mathbf{K}_i \mathbf{u}_{i-1} + \mathbf{b}_i) \mathbf{K}_i^\top J\end{aligned}$$

24 Node2vec: Unsupervised Feature Learning

24.1 Basics

MF approx. + lld optimization, 在某种采样策略 S 下:

$$\max_f \sum_{u \in V} \log \Pr(N_S(u) | f(u))\tag{259}$$

Algorithm 1 The node2vec algorithm.

LearnFeatures (Graph $G = (V, E, W)$, Dimensions d , Walks per node r , Walk length l , Context size k , Return p , In-out q)
 $\pi = \text{PreprocessModifiedWeights}(G, p, q)$
 $G' = (V, E, \pi)$
Initialize $walks$ to Empty
for $iter = 1$ **to** r **do**
 for all nodes $u \in V$ **do**
 $walk = \text{node2vecWalk}(G', u, l)$
 Append $walk$ to $walks$
 $f = \text{StochasticGradientDescent}(k, d, walks)$
 return f

node2vecWalk (Graph $G' = (V, E, \pi)$, Start node u , Length l)
 Initialize $walk$ to $[u]$
 for $walk_iter = 1$ **to** l **do**
 $curr = walk[-1]$
 $V_{curr} = \text{GetNeighbors}(curr, G')$
 $s = \text{AliasSample}(V_{curr}, \pi)$
 Append s to $walk$
 return $walk$

进一步使用 MF(邻域内条件独立)

$$\Pr(N_S(u) | f(u)) = \prod_{n_i \in N_S(u)} \Pr(n_i | f(u)) \quad (260)$$

特征空间对称性(点之间)给出

$$\Pr(n_i | f(u)) = \frac{\exp(f(n_i) \cdot f(u))}{\sum_{v \in V} \exp(f(v) \cdot f(u))} \quad (261)$$

最后优化目标为

$$\max_f \sum_{u \in V} \left[-\log Z_u + \sum_{n_i \in N_S(u)} f(n_i) \cdot f(u) \right] \quad (262)$$

24.2 Biased Random Walk

不同与传统的BFS/DFS, 采用一种折衷的方法(二阶Markov随机游走), 设上一步为($t \rightarrow v$), 则下一步的转移概率为 $\pi_{vx} = \alpha_{pq}(t, x)$, 其中

$$\alpha_{pq}(t, x) = \begin{cases} \frac{1}{p} & \text{if } d_{tx} = 0 \\ 1 & \text{if } d_{tx} = 1 \\ \frac{1}{q} & \text{if } d_{tx} = 2 \end{cases} \quad (263)$$

其中 d_{uv} 是节点最短路径.

考虑其中的参数

1. Return Param. p , 控制了回到之前的点的概率. 设为较高的值 ($> \max(q, 1)$) 可以防止回到原点, 设为较低的值 ($< \max(p, 1)$) 则会鼓励在原点附近探索.
2. In-out Param. q . 大于 1 的值偏向于探索原点附近的节点, 小于 1 的值偏向于 DFS 那样的远离探索.

l 长度的游走可以为 k 个节点生成 $l - k$ 大小的邻域, 总时间复杂度为 $O(\frac{l}{k(l-k)})$

24.3 Edge Feature

简单的在节点间使用 edge feature generator 即可

Operator	Symbol	Definition	
Average	\oplus	$[f(u) \oplus f(v)]_i = \frac{f_i(u) + f_i(v)}{2}$	
Hadamard	\square	$[f(u) \square f(v)]_i = f_i(u) * f_i(v)$	(264)
Weighted-L1	$\ \cdot\ _{\bar{1}}$	$\ f(u) \cdot f(v)\ _{\bar{1}i} = f_i(u) - f_i(v) $	
Weighted-L2	$\ \cdot\ _{\bar{2}}$	$\ f(u) \cdot f(v)\ _{\bar{2}i} = f_i(u) - f_i(v) ^2$	

25 DeepWalk: Online Representation Learning

Note 自然语言中词语的出现 pdf 和社交图中节点在短随机游走中出现的概率都近似服从幂律分布.

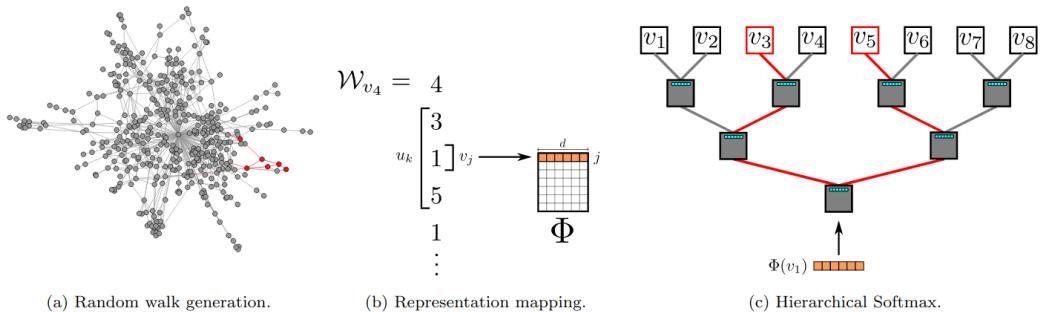


Figure 3: Overview of DEEPWALK. We slide a window of length $2w + 1$ over the random walk \mathcal{W}_{v_4} , mapping the central vertex v_1 to its representation $\Phi(v_1)$. Hierarchical Softmax factors out $\Pr(v_3 \mid \Phi(v_1))$ and $\Pr(v_5 \mid \Phi(v_1))$ over sequences of probability distributions corresponding to the paths starting at the root and ending at v_3 and v_5 . The representation Φ is updated to maximize the probability of v_1 co-occurring with its context $\{v_3, v_5\}$.

25.1 DeepWalk

从图中的每个节点开始 (使用一个随机生成的二叉生成树来指定顺序), 进行随机游走, 长度不定, 在邻域上均匀采样决定下一个节点, 接着使用 SkipGram 算法来更新节点表示. 每一个生成的随机游走为 \mathcal{W}_{v_i} , 长度为 t . 注意, 特征的形式为表式 $\Phi \in \mathbb{R}^{|V| \times d}$

Algorithm 1 DEEPWALK(G, w, d, γ, t)

Input: graph $G(V, E)$ window size w embedding size d walks per vertex γ walk length t **Output:** matrix of vertex representations $\Phi \in \mathbb{R}^{|V| \times d}$ 1: Initialization: Sample Φ from $\mathcal{U}^{|V| \times d}$ 2: Build a binary Tree T from V 3: **for** $i = 0$ to γ **do**4: $\mathcal{O} = \text{Shuffle}(V)$ 5: **for each** $v_i \in \mathcal{O}$ **do**6: $\mathcal{W}_{v_i} = \text{RandomWalk}(G, v_i, t)$ 7: SkipGram($\Phi, \mathcal{W}_{v_i}, w$)8: **end for**9: **end for**

25.2 SkipGram

Algorithm 2 SkipGram($\Phi, \mathcal{W}_{v_i}, w$)

1: **for each** $v_j \in \mathcal{W}_{v_i}$ **do**2: **for each** $u_k \in \mathcal{W}_{v_i}[j - w : j + w]$ **do**3: $J(\Phi) = -\log \Pr(u_k | \Phi(v_j))$ 4: $\Phi = \Phi - \alpha * \frac{\partial J}{\partial \Phi}$ 5: **end for**6: **end for**

SkipGram 是一个最大化一句话中词汇 co-occurrence 概率的算法。最大化每个窗口中的 lld $J(\Phi) = -\log \Pr(u_k | \Phi(v_j))$ 。为了方便计算 lld，引入 Hierachical Softmax。

25.3 Hierachical Softmax

把每一个节点放到一个二叉树的树叶上，然后每个节点的 lld 为按照一个从根到叶子节点的路径（的乘积）

$$\Pr(u_k | \Phi(v_j)) = \prod_{l=1}^{\lceil \log |V| \rceil} \Pr(b_l | \Phi(v_j)) \quad (265)$$

中间每一层都是这样,一个节点的所有条件 lld 的计算代价为 $O(|V| \log |V|)$ 还可以通过 Huffman 树来让常用的节点到根的长度更小.⁴

25.4 Parallelization

由于每个随机游走过程的 SGD 相对独立,可以并行化并使用 ASGD 来进行参数更新.

25.5 Variants

Streaming Learning: 在没有整个图的知识的情况下学习,此时应该不使用递减学习率(退火),可能也无法显式地建立树,如果能知道节点数的上限,则可以用那个最大值来建树.若具有对于节点出现频率的先验知识,则可以使用 Huffman 编码来建树.

Non-random Walks: 有些图是有特定的生成结构的,我们可以利用这些生成结构来指定随机游走的顺序.

26 DAGNN: Towards Deeper GNN




Figure 5: An illustration of the proposed Deep Adaptive Graph Neural Network (DAGNN). For clarity, we show the pipeline to generate the prediction for one node. Notation letters are consistent with Eq.(8) but bold lowercase versions are applied to denote representation vectors. s is the projection vector that computes retainment scores for representations generating from various receptive fields. s_0, s_1, s_2 , and s_k represent the retainment scores of z, h_1, h_2 , and h_k , respectively.

⁴ A brief supplement from NLP: 中间节点每一项都是一个二分类器/Logistic Reg.:

$$p(d_j^w | \mathbf{x}_w, \theta_{j-1}^w) = \begin{cases} \sigma(\mathbf{x}_w^\top \theta_{j-1}^w), & d_j^w = 0 \\ 1 - \sigma(\mathbf{x}_w^\top \theta_{j-1}^w), & d_j^w = 1 \end{cases} \quad (266)$$

其中 w 是那个 word/节点, d_j^w 是中间节点/二叉树指示编码, 写成一个式子为

$$p(d_j^w | \mathbf{x}_w, \theta_{j-1}^w) = [\sigma(\mathbf{x}_w^\top \theta_{j-1}^w)]^{1-d_j^w} \cdot [1 - \sigma(\mathbf{x}_w^\top \theta_{j-1}^w)]^{d_j^w} \quad (267)$$

lld 为

$$\begin{aligned} \mathcal{L}_w &= \log \prod_{j=2}^{l^w} \left\{ [\sigma(\mathbf{x}_w^\top \theta_{j-1}^w)]^{1-d_j^w} \cdot [1 - \sigma(\mathbf{x}_w^\top \theta_{j-1}^w)]^{d_j^w} \right\} \\ &= \sum_{j=2}^{l^w} \{(1 - d_j^w) \cdot \log [\sigma(\mathbf{x}_w^\top \theta_{j-1}^w)] + d_j^w \cdot \log [1 - \sigma(\mathbf{x}_w^\top \theta_{j-1}^w)]\} \end{aligned} \quad (268)$$

26.1 Smoothness Metrics

使用欧式距离为相似度度量

$$D(\mathbf{x}_i, \mathbf{x}_j) = \frac{1}{2} \left\| \frac{\mathbf{x}_i}{\|\mathbf{x}_i\|} - \frac{\mathbf{x}_j}{\|\mathbf{x}_j\|} \right\|, \quad (269)$$

点到图的平滑度

$$SMV_i = \frac{1}{n-1} \sum_{j \in V, j \neq i} D(x_i, x_j) \quad (270)$$

图的总平滑度度量

$$SMV_G = \frac{1}{n} \sum_{i \in V} SMV_i \quad (271)$$

GCN 随着层数增加, 特征的平滑度缓慢下降, 但准确度迅速下降(数层). 这可能是由于 propag. 和特征变换的耦合导致的. 解耦了的 SGC 则在 75-100 层以后准确度/平滑度迅速下降 (over-smoothing 问题).

26.2 Convergence of Propagation

Theorem 26.1 给定图 G , $\widehat{A}_{\oplus} = \tilde{D}^{-1} \tilde{A}$ and $\widehat{A}_{\odot} = \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$, $\Psi(x) = \frac{x}{\text{sum}(x)}$, $\Phi(x) = \frac{x}{\|x\|}$.

$$\lim_{k \rightarrow \infty} \widehat{A}_{\oplus}^k = \Pi_{\oplus}$$

, 其中 Π_{\oplus} 每行都是

$$\pi_{\oplus} = \Psi(e\bar{D})$$

Theorem 26.2 给定图 G ,

$$\lim_{k \rightarrow \infty} \widehat{A}_{\odot}^k = \Pi_{\odot}$$

, 其中

$$\Pi_{\odot} = \Phi \left(\tilde{D}^{\frac{1}{2}} e^T \right) \left(\Phi \left(\tilde{D}^{\frac{1}{2}} e^T \right) \right)^T$$

这两个定理说明了, 如果使用无限层的 propagation, 会导致传递矩阵的收敛和退化, 进而导致不可分的 feat. repr./ over-smoothing. 这不可避免的是一个问题, 所以我们应该更加关心收敛速度.

26.3 DAGNN: Deep Adaptive GNN

DAGNN 的结构如下

$$\begin{aligned} Z &= \text{MLP}(X) && \in \mathbb{R}^{n \times c} \\ H_{\ell} &= \widehat{A}^{\ell} Z, \ell = 1, 2, \dots, k && \in \mathbb{R}^{n \times c} \\ H &= \text{stack}(Z, H_1, \dots, H_k) && \in \mathbb{R}^{n \times c} \\ S &= \sigma(Hs) && \in \mathbb{R}^{n \times (k+1) \times 1} \\ \tilde{S} &= \text{reshape}(S) && \in \mathbb{R}^{n \times 1 \times (k+1)} \\ X_{\text{out}} &= \text{softmax}(\text{squeeze}(\tilde{S}H)) && \in \mathbb{R}^{n \times (k+1) \times c} \end{aligned} \quad (272)$$

这里使用对称正规化的传播矩阵 (GCN-like), $s \in \mathbb{R}^{c \times 1}$ 是 (小型嵌入式 MLP 中的) 的可训练的投影向量 (计算出的 \tilde{S} 为赋予不同大小 receptive fields 的特征向量权重). DAGNN 没有 FC 层! 输出就直接为类别预测分数.

27 t-SNE(t-Distributed Stochastic Neighbor Embedding)

27.1 SNE

Target $f : X \rightarrow Y \in R^{3or2}$, 一个降维映射

将欧式距离变为条件概率 $p_{j|i}$, 常用的概率如归一化的 Gaussian

$$p_{j|i} = \frac{\exp\left(-\|x_i - x_j\|^2 / 2\sigma_i^2\right)}{\sum_{k \neq i} \exp\left(-\|x_i - x_k\|^2 / 2\sigma_i^2\right)} \quad (273)$$

并且设置为自身相似度为 0: $p_{i|i} = 0$

在低维度下的相似度也类似定义, 并固定方差为 $1/\sqrt{2}$

$$q_{j|i} = \frac{\exp\left(-\|y_i - y_j\|^2\right)}{\sum_{k \neq i} \exp\left(-\|y_i - y_k\|^2\right)} \quad (274)$$

则优化变化前后的两个分布的 KL 散度

$$C = \sum_i KL(P_i | Q_i) = \sum_{i,j} p_{j|i} \log \left\{ \frac{p_{j|i}}{q_{j|i}} \right\} \quad (275)$$

困惑度 (perplexity)

$$\text{Perp}(P_i) = 2^{H(P_i)} H(P_i) = - \sum_j p_{j|i} \log_2 p_{j|i} \quad (276)$$

用于选择方差. 使用二分搜索困惑度指标找到最优方差.

在优化开始阶段可以加入一些 Gaussian noise, 之后如同退火逐渐减少噪声幅度, 可以避免局部最优解. 无法避免 crowding 问题.

27.2 UNI-SNE

给低维空间一个均匀分布基准. 可通过退火逐渐减小这个基准

27.3 t-SNE

使用对称的联合 $\text{pdf}p_{ij} = \frac{1}{2}(p_{i|j} + p_{j|i})$ 来解决不对称性. 同时低维分布改为 t 分布

$$f(t) = \frac{\Gamma(\frac{v+1}{2})}{\sqrt{v\pi}\Gamma(\frac{v}{2})} \left(1 + \frac{t^2}{v}\right)^{-\frac{v+1}{2}} \quad (277)$$

$v = 1$ 时为 Cauchy 分布

$$f(t) = \frac{1}{\pi(1 + t^2)}$$

Tricks

- 提前压缩: 开始初始化时点离得近些, 方便聚类中心移动. 可通过 L2 正则项的引入实现?
- 提前夸大: 开始优化阶段 p_{ij} 进行扩大, 避免太小导致优化太慢

Cons

- 主要用于可视化, 难以用于特征提取.
- 倾向于保存局部特征. 对于内蕴维度 (intrinsic dim.) 较高的数据集不可能完整映射.
- 没有唯一解. 没有预估. 训练太慢 ($O(n^2)$), 后续有基于树的改进.

27.4 Barnes-Hut-SNE

27.4.1 Approximating Input Similarities by Vantage-point Tree

使用一定数量的最近邻而非全部点, 来估计相似度.

$$p_{j|i} = \begin{cases} \frac{\exp(-d(\mathbf{x}_i, \mathbf{x}_j)^2 / 2\sigma_i^2)}{\sum_{k \in \mathcal{N}_i} \exp(-d(\mathbf{x}_i, \mathbf{x}_k)^2 / 2\sigma_i^2)}, & \text{if } j \in \mathcal{N}_i \\ 0, & \text{otherwise} \end{cases} \quad (278)$$

$$p_{ij} = \frac{p_{j|i} + p_{i|j}}{2N}$$

最近邻集合的选取可以通过 VPT 来找到, 时间代价为 $O(uN \log N)$, u 为 perplexity.

一个 VPT 中, 每个节点保存了一个数据对象和一个以其为中心的球. 所有非叶节点都有两个孩子, 左儿子保存了所有在球内部的数据对象, 右儿子则保存了所有在外的数据对象. VPT 通过一个一个遍历数据对象构建, 每次根据在外/内便利节点, 并且创建新节点, 其半径为父节点所有对象到他的距离中位数.

一次最近邻搜索可以用 VPT 上的 DFS 来实现, 计算所有节点到目标节点的距离, 维护已经找到的最近邻和到最远近邻的距离 τ . τ 决定了是否要继续探索: 若左节点里可能有比它更近的节点, 搜索左节点, 右边节点同理. 若目标节点在左节点的球中, 先搜索左节点, 右边同理.

27.4.2 Approximating t-SNE Gradients

有

$$\frac{\partial C}{\partial \mathbf{y}_i} = 4(F_{\text{attr}} - F_{\text{rep}}) = 4 \left(\sum_{j \neq i} p_{ij} q_{ij} Z(\mathbf{y}_i - \mathbf{y}_j) - \sum_{j \neq i} q_{ij}^2 Z(\mathbf{y}_i - \mathbf{y}_j) \right) \quad (279)$$

前半部分可以方便的算出, 后半部分可利用 Barnes-Hut 算法快速近似, $O(N \log N)$.

考虑两个点 y_j, y_k 很接近, 那么他们在 y_i 梯度中的贡献就很相近. BH 算法利用这一点, 在 embed. dist. 上建立一个 quad-tree 来估计总梯度.

Quad-Tree 是一个树, 每个节点代表了一个矩形, 非叶节点有四个子节点, 代表了划分为个象限的四个矩形. 叶节点包含最多一个 embedding 点. 在每个节点, 保存矩形的质心 y_{cell} , 和总包含点数. 一个 N 个点的 quad-tree 可以在 $O(N)$ 时间构建. 每个 cell 中对总梯度的贡献相似, 所以

$$\sum_{j \in cell} q_{ij}^2 Z(\mathbf{y}_i - \mathbf{y}_j) \approx N_{cell} q_{i,cell}^2 Z(\mathbf{y}_i - \mathbf{y}_{cell}) \quad (280)$$

并且定义单元配分函数

$$q_{i,cell} Z = \left(1 + \|\mathbf{y}_i - \mathbf{y}_{cell}\|^2\right)^{-1} \quad (281)$$

定义 trade-off factor, 衡量一个单元是否可以作为整体参与计算

$$\|\mathbf{y}_i - \mathbf{y}_{cell}\|^2 / r_{cell} < \theta \quad (282)$$

5

Dual-Tree Algorithms: 使用 cell-cell 距离来进一步减少计算.

28 Autoregressive Flows

Planar Flow

$$f(\mathbf{z}) = \mathbf{z} + \mathbf{u}h(\mathbf{w}^T \mathbf{z} + b) \quad (283)$$

行列式为

$$\left| \det \frac{\partial f}{\partial \mathbf{z}} \right| = \left| 1 + \mathbf{u}^T h'(\mathbf{w}^T \mathbf{z} + b) \mathbf{w} \right| \quad (284)$$

可以看做空间中的超平面, 收缩或者扩张其附近的空间.

Radial Flow

$$f(\mathbf{z}) = \mathbf{z} + \beta h(\alpha, r)(\mathbf{z} - \mathbf{z}_0) \quad (285)$$

其中 $r = \|\mathbf{z} - \mathbf{z}_0\|_2$, $h(\alpha, r) = \frac{1}{\alpha+r}$ 类似的, 可以看作超空间的一个球, 收缩或者扩张其中的空间.

由于这些都是比较 sparse 的变换, 只影响空间的一小部分, 需要很多层才能对高维空间有效.

28.1 Autoregressive Transformation

使用 per-dim 的 AF

$$\begin{aligned} y_1 &= \mu_1 + \sigma_1 z_1 \\ y_i &= \mu(\mathbf{y}_{1:i-1}) + \sigma(\mathbf{y}_{1:i-1}) z_i \end{aligned} \quad (286)$$

Jacobian 是下三角矩阵, 行列式容易计算

$$|detJ| = \left| \prod_i \sigma(\mathbf{y}_{1:i-1}) \right| \quad (287)$$

⁵In other words, more easy to comprehend:

$$r_{cell} > \|y_i - y_{cell}\|^2 / \theta$$

则不行

逆变换为

$$z_i = \frac{y_i - \mu(\mathbf{y}_{1:i-1})}{\sigma(\mathbf{y}_{1:i-1})} \quad (288)$$

由于无法并行计算所有维度, 必须顺序计算, 计算代价很高.

28.2 MAF: Masked Autoregressive Flow

MAF 用上文中的公式(286)进行变换, 这导致了他采样时极为缓慢. 在图像生成中尤为如此, 不过作为 VAE 的先验倒是可以接受 (如 1000 维).

28.3 IAF: Inverse Autoregressive Flow

IAF 使用(288)来进行重参数化 pdf. 使用之前的逆变换作为变换

$$y_i = z_i \sigma(\mathbf{z}_{1:i-1}) + \mu(\mathbf{z}_{1:i-1}) \quad (289)$$

此时所有的 σ, μ 可以并行获得! i.e.

$$\mathbf{y} = \mathbf{z} \circ \sigma(\mathbf{z}) + \mu(\mathbf{z}) \quad (290)$$

假设 $\mathbf{z}_k = \mathbf{z}, \mathbf{z}_{k+1} = \mathbf{y}$, 则有

$$\frac{\partial \mathbf{z}_k}{\partial \mathbf{z}_{k-1}} = \underbrace{\frac{\partial \mu_k}{\partial \mathbf{z}_{k-1}} + \frac{\partial \sigma_k}{\partial \mathbf{z}_{k-1}} \text{diag}(\mathbf{z}_{k-1})}_{\text{lower triangular with zeros on the diagonal}} + \text{diag}(\sigma_k) \underbrace{\frac{\partial \mathbf{z}_{k-1}}{\partial \mathbf{z}_{k-1}}}_{=\mathbf{I}} \quad (291)$$

以及行列式

$$\det\left(\frac{\partial \mathbf{z}_k}{\partial \mathbf{z}_{k-1}}\right) = \prod_{i=1}^d \sigma_{k,i} \quad (292)$$

最终 log-pdf 可以写作

$$\log q_K(\mathbf{z}_k) = \log q(\mathbf{z}) - \sum_{k=0}^K \sum_{i=1}^d \log \sigma_{k,i} \quad (293)$$

由于他的逆是 MAF, 所以从目标分布倒过来求 density 需要计算所有逆, 这就和 MAF 一样难于计算, 虽然仍是可能的.

28.4 Sylvester NF(UAI 18')

28.4.1 Idea

考虑单层 MLP 作为流

$$\mathbf{z}' = \mathbf{z} + \mathbf{A}h(\mathbf{B}\mathbf{z} + \mathbf{b}) \quad (294)$$

Jacobian 可以从 Sylvester 行列式恒等式推出. Sylvester 恒等式指出, 对于矩阵 $\mathbf{A} \in \mathbb{R}^{D \times M}, \mathbf{B} \in \mathbb{R}^{M \times D}$, 有

$$\det(\mathbf{I}_D + \mathbf{AB}) = \det(\mathbf{I}_M + \mathbf{BA}) \quad (295)$$

据此, 上述 MLP 的 Jacobian 行列式为

$$\det\left(\frac{\partial \mathbf{z}'}{\partial \mathbf{z}}\right) = \det(\mathbf{I}_M + \text{diag}(h'(\mathbf{B}\mathbf{z} + \mathbf{b})) \mathbf{BA}) \quad (296)$$

28.4.2 Parametrization of A & B

一般来说,MLP 作为流不是可逆的, 提出以下特例

$$\mathbf{z}' = \mathbf{z} + \mathbf{Q}\mathbf{R}h\left(\tilde{\mathbf{R}}\mathbf{Q}^T\mathbf{z} + \mathbf{b}\right) = \phi(\mathbf{z}) \quad (297)$$

, 其中 $\mathbf{R}, \tilde{\mathbf{R}}$ 是上三角矩阵, $\mathbf{Q} = (\mathbf{q}_1 \dots \mathbf{q}_M)$ 是一个正交基构成的矩阵/正交矩阵. 则 Jacobian 行列式变为

$$\begin{aligned} \det \mathbf{J} &= \det \left(\mathbf{I}_M + \text{diag} \left(h' \left(\tilde{\mathbf{R}}\mathbf{Q}^T\mathbf{z} + \mathbf{b} \right) \right) \tilde{\mathbf{R}}\mathbf{Q}^T\mathbf{Q}\mathbf{R} \right) \\ &= \det \left(\mathbf{I}_M + \text{diag} \left(h' \left(\tilde{\mathbf{R}}\mathbf{Q}^T\mathbf{z} + \mathbf{b} \right) \right) \tilde{\mathbf{R}}\mathbf{R} \right) \end{aligned} \quad (298)$$

可以在 $O(M)$ 时间内计算.

给出以上流的可逆性条件

Theorem 28.1 以上流是可逆的, 若满足

$$r_{ii}\tilde{r}_{ii} > -\frac{1}{\|h'\|_\infty}$$

28.4.3 Preserving Orthogonality of Q

生成一个正交矩阵不总是可行的! 下面介绍两个显式可微地构建正交矩阵的方法, 和一个使用 permut-mat. 的方法.

Orthogonal Sylvester Flows/O-SNF 使用如下可微变换

$$\mathbf{Q}^{(k+1)} = \mathbf{Q}^{(k)} \left(\mathbf{I} + \frac{1}{2} \left(\mathbf{I} - \mathbf{Q}^{(k)\top} \mathbf{Q}^{(k)} \right) \right) \quad (299)$$

一个充分收敛性条件

$$\|\mathbf{Q}^{(0)\top} \mathbf{Q}^{(0)} - \mathbf{I}\|_2 < 1 \quad (300)$$

在本文实验中, 进行迭代直到

$$\|\mathbf{Q}^{(k)\top} \mathbf{Q}^{(k)} - \mathbf{I}\|_F \leq \epsilon \quad (301)$$

实验中大概会进行 30 次左右, 为了提高性能, 对所有流并行地计算这个正交化过程.

Householder Sylvester Flows/H-SNF

Householder reflection, with respect to $\mathbf{v} \in \mathbb{R}^D$

$$H(\mathbf{z}) = \mathbf{z} - 2 \frac{\mathbf{v}\mathbf{v}^T}{\|\mathbf{v}\|^2} \mathbf{z} \quad (302)$$

⁶ 具体使用的 Householder reflection 数量是一个超参数, 而且要求 $M = D$, 因为它使用的是方阵.

Traingular Sylvester Flows/T-SNF

考虑为一个三角阵, 其中每个正交阵都在恒等矩阵和逆转 z 顺序的 permut-mat. 之间转换 (??), 这等同于在每个流之间交换 $\mathbf{R}, \tilde{\mathbf{R}}$ 的上下三角性.

⁶In matrix sense,

$$H_{\mathbf{z}} = \mathbf{I} - \frac{2\mathbf{v}\mathbf{v}^T}{\|\mathbf{v}\|^2} \quad (303)$$

29 Contrastive Multi-view Representation Learning/MVRLG(ICML 20')

Idea 使用图上的扩散来生成增强图. 通过网络两个 view 得到不同的 repr., 让他们 contrast/作为正负样本. 来源于 CV 的对比学习 (使用 congruent/incongruent views).

可以考虑两种图增强: 初始 feature 上的增强 (如 masking/加入 Gaussian noise), 和图结构上的增强 (增减连接性, 降采样, 从最短距离或扩散矩阵生成全局 view). 前者往往会降低性能 \Rightarrow 使用全局图 + 降采样.

29.1 Augmentations

Diffusion:

$$\mathbf{S} = \sum_{k=0}^{\infty} \Theta_k \mathbf{T}^k \in \mathbb{R}^{n \times n} \quad (304)$$

, 假设 $\sum_k \theta_k = 1$, 其中 \mathbf{T} 是广义转移矩阵. 在 PRR(Personal Page Rank) 和 heat kernel 算法中,

$$\mathbf{T} = \mathbf{A}\mathbf{D}^{-1}, \quad (305)$$

$$\theta_k = \alpha(1 - \alpha)^k \text{(Geometric!)}, \text{ or } \theta_k = e^{-t} t^k / k! \text{(Poisson!)} \quad (306)$$

其中 α 可以看作概率, t 可以看作扩散时间. 可以得到闭形式的解:

$$\begin{aligned} \mathbf{S}^{\text{heat}} &= \exp(t\mathbf{A}\mathbf{D}^{-1} - t) \\ \mathbf{S}^{\text{PPR}} &= \alpha(\mathbf{I}_n - (1 - \alpha)\mathbf{D}^{-1/2}\mathbf{A}\mathbf{D}^{-1/2})^{-1} \end{aligned} \quad (307)$$

关于降采样, 从一个 view 中采样, 并从另一个 view 中取相同的 node 和 edges.

29.2 Encoders

使用 GCN 作为基准 encoder, 对于两个 view 使用分别的编码器 g_θ, g_ω , 使用邻接/扩散矩阵作为两个全等的结构视图. 其中邻接矩阵上的 GCN 使用对称正则化 ($\tilde{\mathbf{A}} = \hat{\mathbf{D}}^{-1/2}\hat{\mathbf{A}}\hat{\mathbf{D}}^{-1/2}$, $\hat{\mathbf{A}} = \mathbf{A} + \mathbf{I}_N$), 并且使用 PReLU 函数作为非线性. GCN 之后接着两层 MLP 得到图(节点)表示.

使用图池化提取全局特征 (readout): 通过 (JK-Net like)concat 所有 GCN 层的点特征的和, 并送到单层 MLP 中

$$\vec{h}_g = \sigma \left(\parallel_{l=1}^L \left[\sum_{i=1}^n \vec{h}_i^{(l)} \right] \mathbf{W} \right) \in \mathbb{R}^{h_d} \quad (308)$$

本实验中比更复杂的 readout 好 (如 DiffPool) 效果并不比这个好. 再送到共享的 proj. head, 一个双层 MLP 中.

29.3 Training

使用 Deep InfoMax, 并且最大化两个视图的 MI

$$\max_{\theta, \omega, \phi, \psi} \frac{1}{|\mathcal{G}|} \sum_{g \in \mathcal{G}} \left[\frac{1}{|g|} \sum_{i=1}^{|g|} \left[\text{MI}(\vec{h}_i^\alpha, \vec{h}_g^\beta) + \text{MI}(\vec{h}_i^\beta, \vec{h}_g^\alpha) \right] \right] \quad (309)$$

Algorithm 1 Contrastive multi-view graph representation learning algorithm.

Input: Augmentations τ_α and τ_β , sampler Γ , pooling \mathcal{P} , discriminator \mathcal{D} , loss \mathcal{L} , encoders $g_\theta, g_\omega, f_\psi, f_\phi$, and training graphs $\{\mathcal{G}|g = (\mathbf{X}, \mathbf{A}) \in \mathcal{G}\}$

for sampled batch $\{g_k\}_{k=1}^N \in \mathcal{G}$ **do**

// compute encodings:

for $k = 1$ to N **do**

$\mathbf{X}_k, \mathbf{A}_k = \Gamma(g_k)$ // sub-sample graph

$\mathbf{V}_k^\alpha = \tau_\alpha(\mathbf{A}_k)$ // first view

$\mathbf{Z}_k^\alpha = g_\theta(\mathbf{X}_k, \mathbf{V}_k^\alpha)$ // node rep.

$\mathbf{H}_k^\alpha = f_\psi(\mathbf{Z}_k^\alpha)$ // projected node rep.

$\vec{h}_k^\alpha = f_\phi(\mathcal{P}(\mathbf{Z}_k^\alpha))$ // projected graph rep.

$\mathbf{V}_k^\beta = \tau_\beta(\mathbf{A}_k)$ // second view

$\mathbf{Z}_k^\beta = g_\omega(\mathbf{X}_k, \mathbf{V}_k^\beta)$ // node rep.

$\mathbf{H}_k^\beta = f_\psi(\mathbf{Z}_k^\beta)$ // projected node rep.

$\vec{h}_k^\beta = f_\phi(\mathcal{P}(\mathbf{Z}_k^\beta))$ // projected graph rep.

end

// compute pairwise similarity:

for $i = 1$ to N and $j = 1$ to N **do**

$s_{ij}^\alpha = \mathcal{D}(\vec{h}_i^\alpha, \mathbf{H}_j^\beta), s_{ij}^\beta = \mathcal{D}(\vec{h}_i^\beta, \mathbf{H}_j^\alpha)$

end

// compute gradients:

$\nabla_{\theta, \omega, \phi, \psi} \frac{1}{N^2} \sum_{i=1}^N \sum_{j=1}^N [\mathcal{L}(s_{ij}^\alpha) + \mathcal{L}(s_{ij}^\beta)]$

end

return $[\mathbf{H}_g^\alpha + \mathbf{H}_g^\beta, \vec{h}_g^\alpha + \vec{h}_g^\beta], \forall g \in \mathcal{G}$

MI 使用一个判别器来建模

$$\mathcal{D}(\cdot, \cdot) : \mathbb{R}^{d_h} \times \mathbb{R}^{d_h} \mapsto \mathbb{R}$$

最简单的, 可以使用内积作为相似度度量.

使用联合分布来做为正样本分布, 边缘分布的乘积作为负样本 (即取不同图的视图作为负 contrast)

30 GCC: Graph Contrastive Coding for GNN Pre-Training(KDD 20')

Note 本工作使用的是结构特征, 没有用节点特征. 用于更好的预测未见过的图: 完全的 transfer across domains.

30.1 GCC Pre-Training

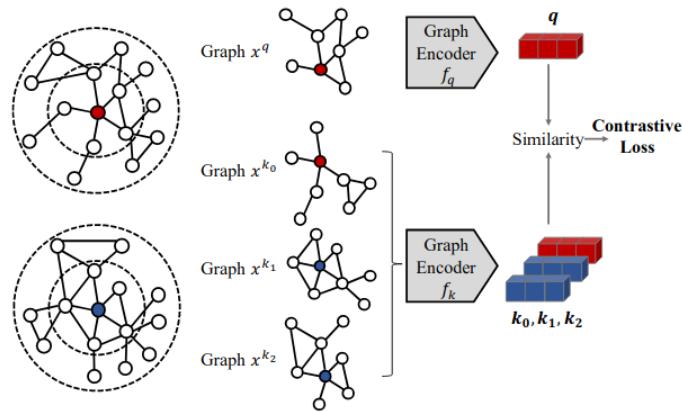


Figure 3: A running example of GCC pre-training.

Idea 使用 subgraph instance-discrimination + InfoNCE! 对于记录了一堆 encoded keys 的 dict $\{k_0, \dots, k_K\}$, 和新编码的 query q , contrast learning 找到一个匹配的 code k_+ , 可以计算 InfoNCE

$$\mathcal{L} = -\log \frac{\exp(q^\top k_+ / \tau)}{\sum_{i=0}^K \exp(q^\top k_i / \tau)} \quad (310)$$

使用子图作为对比学习中的 instance!

Definition 30.1 r -ego network G_v 是一个距离节点 v 最短距离 $\leq r$ 的节点的诱导子图.
(difference with r -hop?)

GCC 把不同 r -ego net. 作为不相似实例.

在 CV 中, 两个经过不同 aug. 的图片作为为了相似 inst. pair, GCC 中, 同样考虑同个 r -ego net. 的不同数据增强作为相似实例对, 并且使用图采样作为数据增强. 具体上讲, GCC 使 三步: 带重启的

随机游走, 子图诱导, 匿名化. 前两步的结果即 ISRW(Induced Subgraph Random Walk Sampling). 最后一步匿名化把节点重新标号. 最后, 两个经过如此图采样的方法被认为是相似实例对.

GCC 理论上使用任何 GNN 都可以 (作为 encoder), 本文使用 GIN(Graph Isomorphism Net.), 当前的 SOTA GNN. 由于大多数 GNN 使用节点 feature 作为输入, 使用 generalized positional embedding(*GPE abbr.)

Definition 30.2 *GPE*. 对于每个子图, 其 *GPE* 为正则化 *Laplacian* 的主要 *eig-vecs*. 即, 在正则化 *Laplacian* 上进行 *EVD*:

$$I - D^{-1/2} A D^{-1/2} = U \Lambda U^\top \quad (311)$$

并取 *top eig-vecs* 作为 *GPE*. (*Question*: *EVD* 不唯一)

GPE 受到 NLP 的 Transformer 的启发. 此外还增加了 one-hot 编码的节点度, 以及 binary 的是否为中心节点的特征.

训练上, 由于维护一个字典代价很高, 所以使用一些别的方法来计算 contrast, 如 E2E(end-to-end) 和 MoCo(momentum constraint).

- E2E 使用 mini-batch, 并把 batch 内的所有 instance 作为 dictionary. Drawback: 词典大小受限与 batch 大小.
- MoCo 使用基于 momentum 更新的 f_k :

$$\theta_k = \rho \theta_k + (1 - \rho) \theta_q$$

30.2 Finetuning GCC

- 对于 graph-level 下游任务, 使用图特征即可. 对于 node-level 下游任务, 使用 r-ego net. 的特征即可.
- 可以 freeze/full fine-tune.
- GCC 作为一个局部探索算法, 可以应用于大规模图和并行计算.

31 GIN: Graph Isomorphism Network(ICLR 19')

31.1 Weisfeiler-Lehman Test

图同构还没有多项式时间算法. WL 算法是一个高效的近似算法, 它在每个节点上聚合邻接节点的 label, 然后 hash 为唯一的新 label. 若在某个层面上 label 不同, 则说明 graph 不同. WL 子树 kernel 把每一次聚合的 label 作为某个子树 (的数据).

31.2 Math Intuitions

Lemma 31.1 若一个 GNN 把两个非同构图映射为不同表示，则 WL 测试也把他们区分为非同构的.

这意味着所有给予聚合的 GNN 都最多和 WL 测试一样强 (在区分非同构图上). 而且, 如果聚合操作和图 readout 函数都是单射, 则和 WL 测试同样强.

Theorem 31.2 Let $\mathcal{A} : \mathcal{G} \rightarrow \mathbb{R}^d$ be a GNN. With a sufficient number of GNN layers, \mathcal{A} maps any graphs G_1 and G_2 that the Weisfeiler-Lehman test of isomorphism decides as non-isomorphic, to different embeddings if the following conditions hold: a) \mathcal{A} aggregates and updates node features iteratively with

$$h_v^{(k)} = \phi(h_v^{(k-1)}, f(\{h_u^{(k-1)} : u \in \mathcal{N}(v)\}))$$

where the functions f , which operates on multisets, and ϕ are injective. b) \mathcal{A} 's graph-level readout, which operates on the multiset of node features $\{h_v^{(k)}\}$, is injective.

31.3 GIN

下面的引理说明了 sum-aggr 是单射, 并且可以用于表示任意函数.

Lemma 31.3 假设 \mathcal{X} 是可数的, 存在函数 $f : \mathcal{X} \mapsto \mathbb{R}^n$, 使得 $h(X) = \sum_{x \in X} f(x)$ 对于任何大小有界的 multi-set 唯一. 进一步, 任何 multiset 函数可以分解为 $g(X) = \phi(\sum_{x \in X} f(x))$.

然而, 常用的 mean-aggr. 不是单射.

Corollary 31.4 假设 \mathcal{X} 是可数的, 存在函数 $f : \mathcal{X} \mapsto \mathbb{R}^n$, 使得存在任意多 ϵ 的选择 (包括所有无理数), 使得 $h(c, X) = (1 + \epsilon)f(c) + \sum_{x \in X} f(x)$ 对于任何大小有界的 multi-set 对 (c, X) 唯一. 进一步, 任何 bi-multiset 函数可以分解为 $g(c, X) = \varphi((1 + \epsilon) \cdot f(c) + \sum_{x \in X} f(x))$.

使用 MLP 近似函数 f, φ , 而且直接表示他们的合成 $f^{(k+1)} \circ \varphi^{(k)}$.

$$h_v^{(k)} = \text{MLP}^{(k)} \left((1 + \epsilon^{(k)}) \cdot h_v^{(k-1)} + \sum_{u \in \mathcal{N}(v)} h_u^{(k-1)} \right) \quad (312)$$

31.4 Graph Readout of GIN

使用每一层级上的 repr.(JK-Net like).

$$h_G = \text{CONCAT} (\text{READOUT} (\{h_v^{(k)} \mid v \in G\}) \mid k = 0, 1, \dots, K) \quad (313)$$

根据已有定理, 如果 READOUT 函数是求和同一层中的 feature, 则确实推广了 WL.

Algorithm 1 Training procedure of Local-instance and Global-semantic Learning (GraphLoG).

Input: Training set $D = \{\mathcal{G}_j\}_{j=1}^{N_D}$, the number of training iterations N_T , hierarchical prototypes' depth L_p and exponential decay rate β .	
Output: The pre-trained GNN.	
Initialize hierarchical prototypes $\{c_i^l\}_{i=1}^{M_l}$ ($l = 1, 2, \dots, L_p$)	
for $t = 1$ to N_T do	
$B_G \leftarrow \text{RandomSample}(D)$	# Get a mini-batch of graphs
$B'_G \leftarrow \text{AttrMasking}(B_G)$	# Get the correlated graphs
$h_{\mathcal{V}_j}, h_{\mathcal{V}'_j}, h_{\mathcal{G}_j}, h_{\mathcal{G}'_j} \leftarrow \text{Eqs. (4, 5)}$ ($j = 1, 2, \dots, N$)	# Extract patch and graph embeddings
$\mathcal{L}_{\text{local}}, \mathcal{L}_{\text{global}} \leftarrow \text{Eqs. (8, 13)}$	# Compute losses
$\theta_{\text{GNN}} \leftarrow -\nabla_{\theta_{\text{GNN}}} (\mathcal{L}_{\text{local}} + \mathcal{L}_{\text{global}})$	# Update GNN's parameters
$\theta_T \leftarrow -\nabla_{\theta_T} (\mathcal{L}_{\text{local}} + \mathcal{L}_{\text{global}})$	# Update discriminator's parameters
$\{c_i^l\}_{i=1}^{M_l} \leftarrow \text{Eqs. (11, 12)}$ ($l = 1, 2, \dots, L_p$)	# Maintain hierarchical prototypes
end for	

32 GraphLoG: Self-Supervised Representation Learning with Local & Global Structure

32.1 Preliminaries

GNN and READOUT GNN-layer

$$h_v^{(l)} = \text{COMBINE}^{(l)} \left(h_v^{(l-1)}, \text{AGGREGATE}^{(l)} \left(\{(h_u^{(l-1)}, h_u^{(l-1)}, X_{uv}) : u \in \mathcal{N}(v)\} \right) \right) \quad (314)$$

Graph Readout:

$$h_{\mathcal{G}} = \text{READOUT}(\{h_v \mid v \in \mathcal{V}\}) \quad (315)$$

Mutual Info. Est. InfoNCE

$$\mathcal{L}_{\text{NCE}}(q, z_+, \{z_i\}_{i=1}^K) = -\log \frac{\exp(T(q, z_+))}{\exp(T(q, z_+)) + \sum_{i=1}^K \exp(T(q, z_i))} \quad (316)$$

其中 T 是某个参数化的判别函数.

RPCL: Rival Penalized Competitive Learning RPCL 在每个新样本上不仅推近 winning cluster(最近聚类), 也推开 rival cluster(次近聚类). 可以不预先指定聚类数地进行 clustering.

32.2 Local-Inst. Stru. Learning

使用 mask 来获得相近 (correlated) 图. 对于每个 mini-batch 获得相近图 (通过 mask). 得到他们的 node(patch)/graph repr.

$$\begin{aligned} h_{\mathcal{V}_j} &= \{h_v \mid v \in \mathcal{V}_j\} = \text{GNN}(X_{\mathcal{V}_j}, X_{\mathcal{E}_j}), \quad h_{\mathcal{V}'_j} = \{h_v \mid v \in \mathcal{V}'_j\} = \text{GNN}(X_{\mathcal{V}'_j}, X_{\mathcal{E}'_j}) \\ h_{\mathcal{G}_j} &= \text{READOUT}(h_{\mathcal{V}_j}), \quad h_{\mathcal{G}'_j} = \text{READOUT}(h_{\mathcal{V}'_j}) \end{aligned} \quad (317)$$

最小化 InfoNCE⁷

$$\begin{aligned}\mathcal{L}_{\text{patch}} &= \frac{1}{\sum_{j=1}^N |\mathcal{V}'_j|} \sum_{j=1}^N \sum_{v' \in \mathcal{V}'_j} \sum_{v \in \mathcal{V}_j} \mathbb{1}_{v \leftrightarrow v'} \cdot \mathcal{L}_{\text{NCE}}(h_{v'}, h_v, \{h_{\tilde{v}} \mid \tilde{v} \in \mathcal{V}_j, \tilde{v} \neq v\}) \\ \mathcal{L}_{\text{graph}} &= \frac{1}{N} \sum_{j=1}^N \mathcal{L}_{\text{NCE}}(h_{\mathcal{G}'_j}, h_{\mathcal{G}_j}, \{h_{\mathcal{G}_k} \mid 1 \leq k \leq N, k \neq j\}) \\ \mathcal{L}_{\text{local}} &= \mathcal{L}_{\text{patch}} + \mathcal{L}_{\text{graph}}\end{aligned}\quad (318)$$

Note 同个 batch 中的其他图 repr. 作为 negative instance. 同个图中其他 node repr. 作为 negative instance.

32.3 Global-Semantic Repr. Learning

Graph(如分子图) 常常有 hierachical structure \Rightarrow hierarchical prototypes, 建模图嵌入的分布. 他们是一些树的集合, 每个节点对应了一个 prototype, 并且指向唯一的父节点 (除非是根节点). 正式地, 这些节点是 $\{c_i^l\}_{i=1}^{M_l}$, ($l = 1, 2, \dots, L_p$), 除了树叶节点, 每个节点都有一些子节点集合 $C(c_i^l)$.

32.3.1 Init. of HP(Hierarchical Prototypes)

先用 $\mathcal{L}_{\text{local}}$ 预训练一个 epoch, 利用这个 GNN 得到所有训练集中图的表示 $\{h_{\mathcal{G}_i}\}_{i=1}^{N_D}$, 作为 HP 的叶子节点, 并且使用 RPCL 得到底层的节点

$$\left\{c_i^{L_p}\right\}_{i=1}^{M_{L_p}} = \text{RPCL}\left(\{h_{\mathcal{G}_i}\}_{i=1}^{N_D}\right) \quad (319)$$

之后 RPCL 被迭代的应用来得到整个 HP 树.

32.3.2 Maintainance of HP

训练过程中, 图嵌入在动态改变. 使用一下策略来更新 HP:

- 每有一个 batch 训练完, 得到图 embedding, 分成 M_{L_p} 个组 (属于 HP 最底层聚类), 计算每一组的平均 embedding. 并且更新节点上的 embedding, 使用 (momentum-like) 指数移动平均更新模型

$$c_i^{L_p} \leftarrow \beta c_i^{L_p} + (1 - \beta) \hat{c}_i^{L_p}, \quad 1 \leq i \leq M_{L_p} \quad (320)$$

此处 β 是衰减常数, 之后上层节点 embedding 更新为子节点的平均.

2.

为了捕捉 global-semantic structure, 让相关图属于同一聚类! 具体上讲, 对于每一个图 \mathcal{G}_j , 使⽤ cosine sim., 在每一层上寻找最接近的 embedding

$$s(\mathcal{G}_j) = \{s_1(\mathcal{G}_j), s_2(\mathcal{G}_j), \dots, s_{L_p}(\mathcal{G}_j)\}$$

这里要让相关图的 embedding 和这些 repr. 相近, 使用 InfoNCE

$$\mathcal{L}_{\text{global}} = \frac{1}{N \cdot L_p} \sum_{j=1}^N \sum_{l=1}^{L_p} \mathcal{L}_{\text{NCE}}(h_{\mathcal{G}'_j}, s_l(\mathcal{G}_j), \{c_i^l \mid 1 \leq i \leq M_l, c_i^l \neq s_l(\mathcal{G}_j)\}) \quad (321)$$

⁷如同其他 contrast learning

即, 对于每一层, 正样本是那个 (原图) 最相近的 embedding, 负样本是 (原图) 其他样本 repr.
模型在每一个 iteration 上最小化全 loss

$$\min_{\text{GNN}, T} \mathcal{L}_{\text{local}} + \mathcal{L}_{\text{global}} \quad (322)$$

32.4 Sup-GraphLoG: A Supervised Baseline

使用 label 来简单监督学习, 作为 baseline. 对于任何一个图, 找到在 HP 底层匹配的 (随机的) 一个匹配的 prototype, 把搜索路径作为正样本, 再随机选择另一条不正确的路径作为负样本. 然后 mini-batch 上的全局 loss 为

$$\mathcal{L}_{\text{global}}^{\text{sup}} = \frac{1}{N \cdot L_p} \sum_{j=1}^N \sum_{l=1}^{L_p} \mathcal{L}_{\text{NCE}}(h_{\mathcal{G}_j}, s_l(\mathcal{G}_j), s_l^n(\mathcal{G}_j)) \quad (323)$$

33 Orthogonal Weights in DNNs

33.1 Formulation & Good Properties

优化参数, 使得权重是 (伪) 正交的

$$\begin{aligned} \theta^* = \arg \min_{\theta} \mathbb{E}_{(\mathbf{x}, \mathbf{y}) \in D} [\mathcal{L}(\mathbf{y}, f(\mathbf{x}; \theta))] \\ \text{s.t. } \mathbf{W}^l \in \mathcal{O}_l^{n_l \times d_l}, l = 1, 2, \dots, L \end{aligned} \quad (324)$$

此处伪正交矩阵在一个实 Stiefel 流形 $\mathcal{O}_l^{n_l \times d_l} = \left\{ \mathbf{W}^l \in \mathbb{R}^{n_l \times d_l} : \mathbf{W}^l (\mathbf{W}^l)^T = \mathbf{I} \right\}$ 上.
 \Rightarrow OMDSM(Optim. on Multiple Dependent Stiefel Manifolds) Problem
好处

- $s = \mathbf{Wx}$, $\mathbf{W} \in \mathbb{R}^{n \times d}$, 若输入 \mathbf{x} 是白化了的, 则输出 s 也是 (0 均值且分量不相关); 如果 $n = d$, 则 $\|s\| = \|\mathbf{x}\|$; 梯度相等 $\left\| \frac{\partial \mathcal{L}}{\partial \mathbf{x}} \right\| = \left\| \frac{\partial \mathcal{L}}{\partial s} \right\|$
- 自动地, 这正则化了权重. 减少了自由度 (Stiefel Manifold 的维度减少了):
 $\dim \mathcal{O}^{n \times d} = nd - n(n+1)/2$

33.2 OWN: Orthogonal Weight Normalization

使用 Riemann(流形) 优化算法 (like in RNNs⁸) 导致了收敛的不稳定性/性能差. 显式地使用参数矩阵的重参数化

$$\phi : \mathbb{R}^{n_l \times d_l} \rightarrow \mathbb{R}^{n_l \times d_l} \quad (325)$$

并且 $\phi(\mathbf{V}) = \mathbf{W}$, s.t. $\mathbf{W}\mathbf{W}^T = \mathbf{I}$.

⁸What's that actually?

使用一个线性变换来作为函数

$$\phi(\mathbf{V}) = \mathbf{P}\mathbf{V}_C \quad (326)$$

$$\mathbf{V}_C = \mathbf{V} - \mathbf{c}\mathbf{1}_d^T \quad (327)$$

$$\mathbf{c} = \frac{1}{d}\mathbf{V}\mathbf{1}_d \quad (328)$$

$$\text{which means } \mathbf{V}_C = \mathbf{V} - \frac{1}{d}\mathbf{V}\mathbf{E}_d = \mathbf{V}(\mathbf{I} - \frac{1}{d}\mathbf{E}_d) \quad (329)$$

使用以下变换来得到权重, 使得 Jacobian 的奇异值接近于 1

$$\begin{aligned} & \min_{\mathbf{P}} \text{tr} \left((\mathbf{W} - \mathbf{V}_C)(\mathbf{W} - \mathbf{V}_C)^T \right) \\ & \text{s.t. } \mathbf{W} = \mathbf{P}\mathbf{V}_C \text{ and } \mathbf{W}\mathbf{W}^T = \mathbf{I} \end{aligned} \quad (330)$$

进而, 这个优化问题有 closed form sol.

$$\mathbf{P}^* = \mathbf{D}\Lambda^{-1/2}\mathbf{D}^T \quad (331)$$

其中 $\mathbf{V}_C\mathbf{V}_C^T = \Sigma = \mathbf{D}\Lambda\mathbf{D}^T$ 是协方差矩阵的 EVD(SVD). 最后的形式为

$$\mathbf{W} = \phi(\mathbf{V}) = \mathbf{D}\Lambda^{-1/2}\mathbf{D}^T (\mathbf{V} - \mathbf{c}\mathbf{1}_d^T) \quad (332)$$

类似但没有最小化(330)的选择 $\mathbf{P}_{var} = \Lambda^{-1/2}\mathbf{D}^T$

33.2.1 Backpropagation

计算 Jacobian

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial \Lambda} &= -\frac{1}{2}\mathbf{D}^T \frac{\partial \mathcal{L}}{\partial \mathbf{W}} \mathbf{W}^T \mathbf{D}\Lambda^{-1} \\ \frac{\partial \mathcal{L}}{\partial \mathbf{D}} &= \mathbf{D}\Lambda^{\frac{1}{2}}\mathbf{D}^T \mathbf{W} \frac{\partial \mathcal{L}}{\partial \mathbf{W}} \mathbf{D}\Lambda^{-\frac{1}{2}} + \frac{\partial \mathcal{L}}{\partial \mathbf{W}} \mathbf{W}^T \mathbf{D} \\ \frac{\partial \mathcal{L}}{\partial \mathbf{S}} &= \mathbf{D} \left(\left(\mathbf{K}^T \odot \left(\mathbf{D}^T \frac{\partial \mathcal{L}}{\partial \mathbf{D}} \right) \right) + \left(\frac{\partial \mathcal{L}}{\partial \Lambda} \right)_{\text{diag}} \right) \mathbf{D}^T \\ \frac{\partial \mathcal{L}}{\partial \mathbf{c}} &= -\mathbf{1}_d^T \frac{\partial \mathcal{L}}{\partial \mathbf{W}} \mathbf{D}\Lambda^{-\frac{1}{2}}\mathbf{D}^T - 2 \cdot \mathbf{1}_d^T (\mathbf{V} - \mathbf{c}\mathbf{1}_d^T)^T \left(\frac{\partial \mathcal{L}}{\partial \Sigma} \right)_s \\ \frac{\partial \mathcal{L}}{\partial \mathbf{V}} &= \mathbf{D}\Lambda^{-\frac{1}{2}}\mathbf{D}^T \frac{\partial \mathcal{L}}{\partial \mathbf{W}} + 2 \left(\frac{\partial \mathcal{L}}{\partial \Sigma} \right)_s (\mathbf{V} - \mathbf{c}\mathbf{1}_d^T) + \frac{1}{d} \frac{\partial \mathcal{L}^T}{\partial \mathbf{c}} \mathbf{1}_d^T \end{aligned} \quad (333)$$

此处

$$\mathbf{K}_{ij} = \frac{1}{\sigma_i - \sigma_j} [i \neq j] \quad (334)$$

是一个对角线为 0 的矩阵.

33.2.2 As Convolution

卷积层的参数 $\mathbf{W}^C \in \mathbb{R}^{n \times d \times F_h \times F_w}$, 以及上一层的输入特征 $\mathbf{X} \in \mathbb{R}^{d \times h \times w}$, 那么 activation 可以计算为

$$s_{k,\delta} = \sum_{i=1}^d \sum_{\tau \in \Omega} w_{k,i,\tau} h_{i,\delta+\tau} = \langle \mathbf{w}_k, \mathbf{h}_\delta \rangle \quad (335)$$

此处使用 $\mathbf{W} \in \mathbb{R}^{n \times p}, p = dF_hF_w$ 作为总的 filters.

33.2.3 Group Based Orthogonalization: Divided Filters

权重按行分为多组, 每组大小一致 (为一个小数 $N_G < d$, 如 $64/128$), 使得 EVD 的计算代价很小. 大大降低计算难度.

34 OrthDNNs: Orthogonal DNNs(TPAMI 19')

在数据分布 $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$ 上, ML 的目标是最优化期望风险

$$R(f) = \mathbb{E}_{z \sim P}[\mathcal{L}(f(\mathbf{x}), y)] \quad (336)$$

然而真实分布未知, 所以用样本期望 (训练集 S_m) 代替

$$R_m(f) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(f(\mathbf{x}_i), y_i) \quad (337)$$

统计学习的一大目标是计算泛化误差/gen. gap

$$\text{GE}(f_{S_m}) = |R(f_{S_m}) - R_m(f_{S_m})| \quad (338)$$

这里考虑由 DNN 建模的分类-表示模型 (Class. Repr. Learning)

$$\begin{aligned} R(f, T) &= \mathbb{E}_{z \sim P}[\mathcal{L}(f(T\mathbf{x}), y)] \\ R_m(f, T) &= \frac{1}{m} \sum_{i=1}^m \mathcal{L}(f(T\mathbf{x}_i), y_i) \end{aligned} \quad (339)$$

34.1 GE Analysis in a Robustness and Isomeric Mapping Perspectve

Definition 34.1 ($K, \epsilon(\cdot)$ -robustness) 对于 $K \in \mathbb{N}, \epsilon : \mathcal{Z}^m \mapsto \mathbb{R}$ 一个算法是 $K, \epsilon(\cdot)$ -robust 的, 如果 \mathcal{Z} 可以分为 K 个分离的集合, 记为 $\mathcal{C} = \{C_k\}_{k=1}^K$, 且

$$\begin{aligned} \forall s_i = (\mathbf{x}_i, y_i) \in C_k, \forall z = (\mathbf{x}, y) \in C_k \\ \implies |\mathcal{L}(f(\mathbf{x}_i), y_i) - \mathcal{L}(f(\mathbf{x}), y)| \leq \epsilon(S_m) \end{aligned} \quad (340)$$

对于任何健壮的算法, 有以下定理⁹

Theorem 34.2 如果一个算法是 $(K, \epsilon(\cdot)$ -robust, 并且一个损失函数 \mathcal{L} 是有界的, 则

$$\Pr \left(\text{GE}(f_{S_m}) \leq \epsilon(S_m) + M \sqrt{\frac{2K \log(2) + 2 \log(1/\nu)}{m}} \right) \geq 1 - \nu \quad (341)$$

Definition 34.3 (*Covering Number*) 给出一个度量空间 (\mathcal{M}, d) , 集合 \hat{S} 是另一个集合 S 的 γ -cover, 若 $\forall s \in S, \exists \hat{s} \in \hat{S}, s.t. d(\hat{s}, s) \leq \gamma$. 集合 S 的 γ -covering number 是

$$\mathcal{N}_\gamma(S, \rho) = \min \{|\hat{S}| : \hat{S} \text{ is a } \gamma \text{-covering of } S\} \quad (342)$$

⁹Huan Xu and Shie Mannor. Robustness and generalization.Machine Learning, 86(3):391–423, 2012.1,2,4,6,8

Definition 34.4 (δ -isometry) 映射 $T : \mathcal{P} \mapsto \mathcal{Q}$ (其中 \mathcal{P}, \mathcal{Q} 度量空间) 是 δ -保距的, 若

$$\forall \mathbf{x}, \mathbf{x}' \in \mathcal{P}, |\rho_Q(T\mathbf{x}, T\mathbf{x}') - \rho_P(\mathbf{x}, \mathbf{x}')| \leq \delta \quad (343)$$

Theorem 34.5 对与任何 CRL 问题的算法, 若 $\mathcal{L} \circ f$ (关于 $T\mathbf{x}$) 的 Lipschitz Constant 有上界 A , 且 T 是 δ -保距的, 且 \mathcal{X} 是紧的, 且有 covering number $\mathcal{N}_{\gamma/2}(\mathcal{X}, \rho)$, 则这个算法是 $(|\mathcal{Y}| \mathcal{N}_{\gamma/2}(\mathcal{X}, \rho), A(\gamma + \delta))$ -robust.

34.2 GE Analysis of DNN

10

34.3 OrthDNN by SVB(Singular Value Bound)

类似于之前的一篇, 为在 Stiefel 流形上的约束优化. 可以通过流形上的 SGD(切空间投影法) 或者 Frank-Wolfe 算法(流形投影法). 太慢! 使用 SVB 来进行优化.

- 使用估计的带偏移的梯度投影方向. 每隔一定数量迭代拉回到流形上.
- 考虑到 DNN 的优化问题含有巨大数量的局部极小/临界点. 在目标流形附近探索可能会得到更好的解(避免 local minima/critical points)
- BN 会改变权重矩阵的谱, 使得严格正交化的努力白费.

Algorithm Sketch(SVB)

1. 使用普通的 SGD 来更新参数.
2. 每隔一些 epochs, 对于每一层, 使用 SVD 来 clamp 奇异值到 $((1 + \epsilon)^{-1}, 1 + \epsilon)$:

$$\mathbf{W} = \mathbf{U}\Sigma\mathbf{V}^T \quad (344)$$

$$\mathbf{W} \leftarrow \mathbf{U}\Sigma_{\epsilon-clamped}\mathbf{V}^T \quad (345)$$

还可以使用罚函数法来作为无约束优化问题优化 (SoftRegu)

$$\min_{\Theta=\{\mathbf{W}_l, b_l\}_{l=1}^L} \mathcal{L}(\{\mathbf{x}_i, y_i\}_{i=1}^m; \Theta) + \lambda \sum_{l=1}^L \|\mathbf{W}_l^\top \mathbf{W}_l - \mathbf{I}\|_F^2 \quad (346)$$

需要假设 $n_l \geq n_{l-1}$, 为了放松这一假设, 使用自然 1-范数/谱范数 (SRIP) 而不是 Frobenius 范数

$$\min_{\Theta=\{\mathbf{W}_l, b_l\}_{l=1}^L} \mathcal{L}(\{\mathbf{x}_i, y_i\}_{i=1}^m; \Theta) + \kappa \sum_i \sigma_{\max}(\mathbf{W}_l^\top \mathbf{W}_l - \mathbf{I}) \quad (347)$$

¹⁰TO READ AND MAKE NOTES

34.3.1 BN Compatibility

BN 实际上干了

$$\text{BN}(\mathbf{h}) = \boldsymbol{\Upsilon}\boldsymbol{\Phi}(\mathbf{h} - \boldsymbol{\mu}) + \boldsymbol{\beta} \quad (348)$$

其中 $\boldsymbol{\mu} \in \mathbb{R}^n$ 是这一层 n 个神经元的 (batch 的) 均值, 对角阵 $\boldsymbol{\Phi} = \text{Diag}(1/\phi_i)$ 是神经元输出的标准差倒数矩阵 (加上小值来增加数值稳定性). 并且 $\boldsymbol{\Upsilon}, \boldsymbol{\beta}$ 是可训练的对角矩阵和 bias. 最终的均值和标准差由 running average 给出训练样本的总体估计.

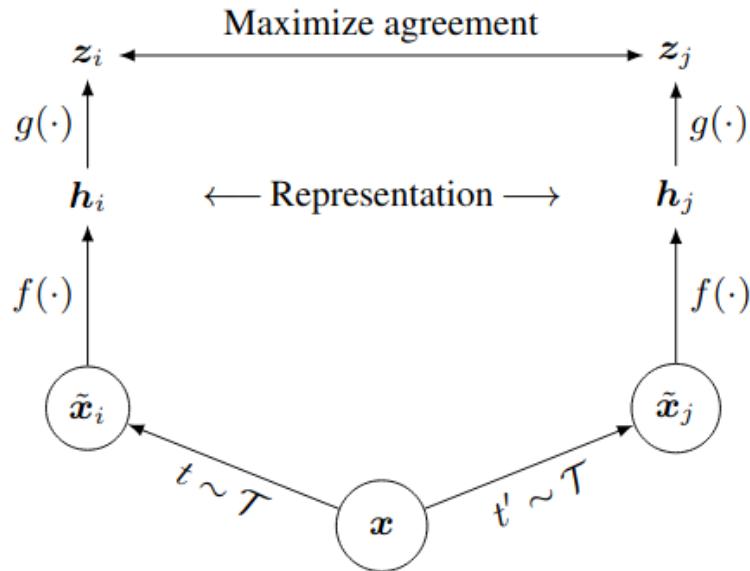
但是乘上一个对角矩阵会导致原先的权重矩阵 sig-val 不再相等! 所以让可训练的对角阵和偏差在每层都相等 \Rightarrow 对于精确形式的 OrhDNN, 使用 DBN(Degenerate Batch Norm), 其中 $\bar{\phi} = \frac{1}{n} \sum_n \phi_i$ 对于 approx. OrthDNN, 使用 BBN(Bounded BN), 控制 (通过 clamp) $\{v_i/\phi_i\}$ 在均值

$$\alpha = \frac{1}{n} \sum_i \frac{v_i}{\phi_i}, v_i/\phi_i \in [\alpha(1 + \epsilon)^{-1}, \alpha(1 + \epsilon)] \text{ 内.}$$

34.3.2 On CNN: OrthDNN as Convolution

对于 CNN 的某一层, 要求的参数形状为 $n_l \times n_{l-1} \times n_h \times n_w$, 本作使用生成 $n_l \times n_{l-1} n_h n_w$ 形状的参数矩阵来得到这些 filter.

35 SimCLR: A Simple Framework for Contrastive Learning of Visual Representations



35.1 Ideas & Basics

1. 使用随机数据增强来得到两个相关视图:

- 随机剪裁/裁剪 + 缩放至原大小

-
- 随机色彩 distortion
 - 随机高斯模糊
 - 没啥用的: cutoff, Sobel filters etc.
2. Encoder 使用经典的 ResNet 结构.
 3. 一个小的 MLP 投影头 (proj. head) 用于投影到 contrastive loss 用于计算的空间.
 4. 在最后的特征空间上最小化 InfoNCE(yet another *contrastive loss*), 最大化相关图的 MI

$$\ell_{i,j} = -\log \frac{\exp(\text{sim}(z_i, z_j) / \tau)}{\sum_{k=1}^{2N} \mathbf{1}_{[k \neq i]} \exp(\text{sim}(z_i, z_k) / \tau)} \quad (349)$$

其中相似度度量使用归一化点积 (i.e. cosine similarity) $\text{sim}(\mathbf{u}, \mathbf{v}) = \mathbf{u}^\top \mathbf{v} / \|\mathbf{u}\| \|\mathbf{v}\|$ 其中负样本来自同一 minibatch, τ 是温度系数.

本方法没有利用类似 MoCo 的 memory bank, 而是使用较大的 batchsize(256-8192).

SGD+Momentum 的优化器似乎和较大的 batchsize 上不稳定, 所以使用 LARS 优化器.

使用 Global BN: 在各个设备上本来 BN 的均值和方差是分别计算的, 这里在 BN 里使用聚合了所有设备的 mean/variance. 其他方法包括在设备间 shuffle samples, 或者使用 layer norm¹¹而不是 BN.

36 BYOL: Build Your Own Latent

36.1 Ideas & Method

使用两个网络来学习: online/target 网络, 每个网络都类似得由三个阶段组成:encoder f_θ , proj. head g_θ , predictor q_θ (discriminator). online/target 网络使用相同的结构, 但是为不同的参数, 并且 target 网络的参数 ξ 是 online 参数 θ 的 moving average(under decay const. τ)

$$\xi \leftarrow \tau \xi + (1 - \tau) \theta \quad (352)$$

¹¹在 MLP 中, 归一化每一层的权重为期望 0 标准差 1, 先计算均值和方差

$$\mu^l = \frac{1}{H} \sum_{i=1}^H w_i^l \quad \sigma^l = \sqrt{\frac{1}{H} \sum_{i=1}^H (w_i^l - \mu^l)^2} \quad (350)$$

在计算 FP 前进行权重归一化

$$\hat{\mathbf{w}}^l = \frac{\mathbf{w}^l - \mu^l}{\sqrt{(\sigma^l)^2 + \epsilon}} \quad (351)$$

LN 在权重够多的情况下和 BN 效果类似. 似乎在 CNN 上不能使用???

Remark 对于权重的更弱的限制 (相比 OrthDNN 和 OWN 中近似正交化权重), 注意随机正交矩阵必然是归一化了的, 分量还是不相关的.

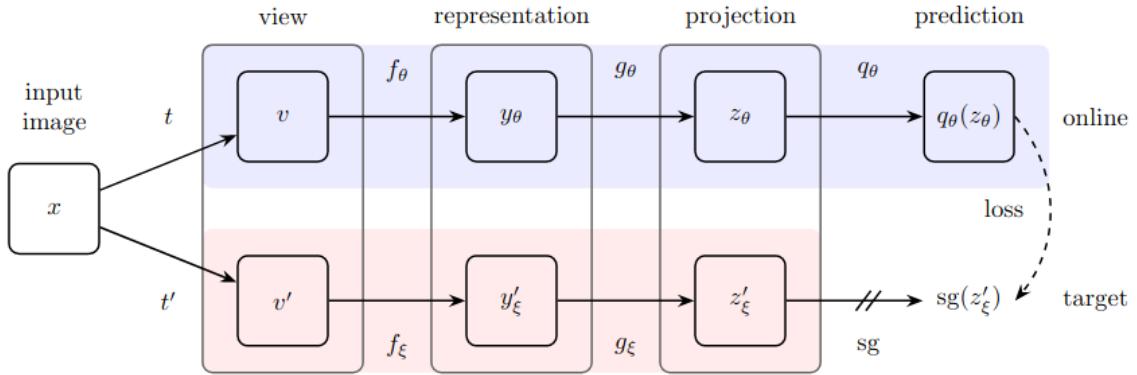


Figure 2: BYOL’s architecture. BYOL minimizes a similarity loss between $q_\theta(z_\theta)$ and $\text{sg}(z'_\xi)$, where θ are the trained weights, ξ are an exponential moving average of θ and sg means stop-gradient. At the end of training, everything but f_θ is discarded, and y_θ is used as the image representation.

假设通过两个图像变换之后的 view 是 v, v' , 然后我们使用要区分两者的 repr., 所以最小化负的 cosine-similarity¹²

$$\mathcal{L}_{\theta,\xi} \triangleq \|\bar{q}_\theta(z_\theta) - \bar{z}'_\xi\|_2^2 = 2 - 2 \cdot \frac{\langle q_\theta(z_\theta), z'_\xi \rangle}{\|q_\theta(z_\theta)\|_2 \cdot \|z'_\xi\|_2} \quad (353)$$

同样的交换两个视图送到网络的顺序, 并且得到另一个 loss, 加和来得到对称化 loss

$$\mathcal{L}_{\theta,\xi}^{\text{BYDL}} = \mathcal{L}_{\theta,\xi} + \tilde{\mathcal{L}}_{\theta,\xi} \quad (354)$$

参数的更新

$$\begin{aligned} \theta &\leftarrow \text{optimizer } (\theta, \nabla_\theta \mathcal{L}_{\theta,\xi}^{\text{BYOL}}, \eta), \\ \xi &\leftarrow \tau \xi + (1 - \tau) \theta \end{aligned} \quad (355)$$

BYOL 没有使用显式的方法防止 mode collapse(negative samples etc.). 但是根据假设, 他们证明了, 至少在 predictor 是 optimal 的时候 ($q_\theta = q^*$), 鞍点是不稳定的.

37 MoCo: Momentum Contrast for Unsupervised Visual Representation Learning

37.1 Ideas & Method

InfoNCE on memory bank

$$\mathcal{L}_q = -\log \frac{\exp(q \cdot k + / \tau)}{\sum_{i=0}^K \exp(q \cdot k_i / \tau)} \quad (356)$$

Momentum encoder: 按照 moving average 更新 k-encoder

$$\theta_k \leftarrow m\theta_k + (1 - m)\theta_q \quad (357)$$

Tricks: Shuffling BNs

¹²equivalent to MSE of normalized feature

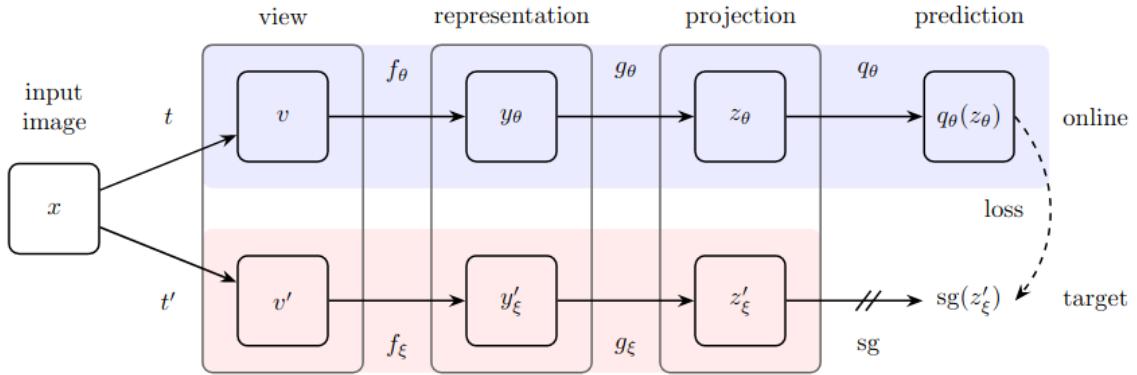


Figure 2: BYOL’s architecture. BYOL minimizes a similarity loss between $q_\theta(z_\theta)$ and $\text{sg}(z'_\xi)$, where θ are the trained weights, ξ are an exponential moving average of θ and sg means stop-gradient. At the end of training, everything but f_θ is discarded, and y_θ is used as the image representation.

38 SimSiam: Exploring Simple Siamese Representation Learning

38.1 Intuitions

BYOL 既没有使用 negative samples, 但是使用了 momentum encoder+ 不对称的 proj. head+stop-grad 策略来避免 collapse, 并且不需要巨大的 batchsize! 所以本方法可以考虑为 BYOL-momentum encoder, 使用两个共享权重的网络 (like SwAV without online clustering/SimCLR without negative pairs), 即使如此, 不会导致 mode collapse!

Remark stop-grad 是关键的, 没有 stop-grad 则会直接收敛到 trivial 解 (mode collapse)

38.2 Method

使用共享权重的 backbone(like ResNet)+proj. head(MLP), 最小化不同视图最后特征的负 cosine-sim.

$$\mathcal{D}(p_1, z_2) = -\frac{p_1}{\|p_1\|_2} \cdot \frac{z_2}{\|z_2\|_2} \quad (358)$$

并且使用对称化 loss

$$\mathcal{L} = \frac{1}{2}\mathcal{D}(p_1, z_2) + \frac{1}{2}\mathcal{D}(p_2, z_1) \quad (359)$$

并且不对一边的网络进行参数更新.

Settings

- Optimizer: SGD for pre-training. $lr = lr_{base} \times \text{BatchSize}/256$, base lr 为 0.05, 使用 cosine decay(annealing) schedule, L2 正则化 1e-4, 动量 0.9.
- Batch Size 512, Synchronized BN across devices.

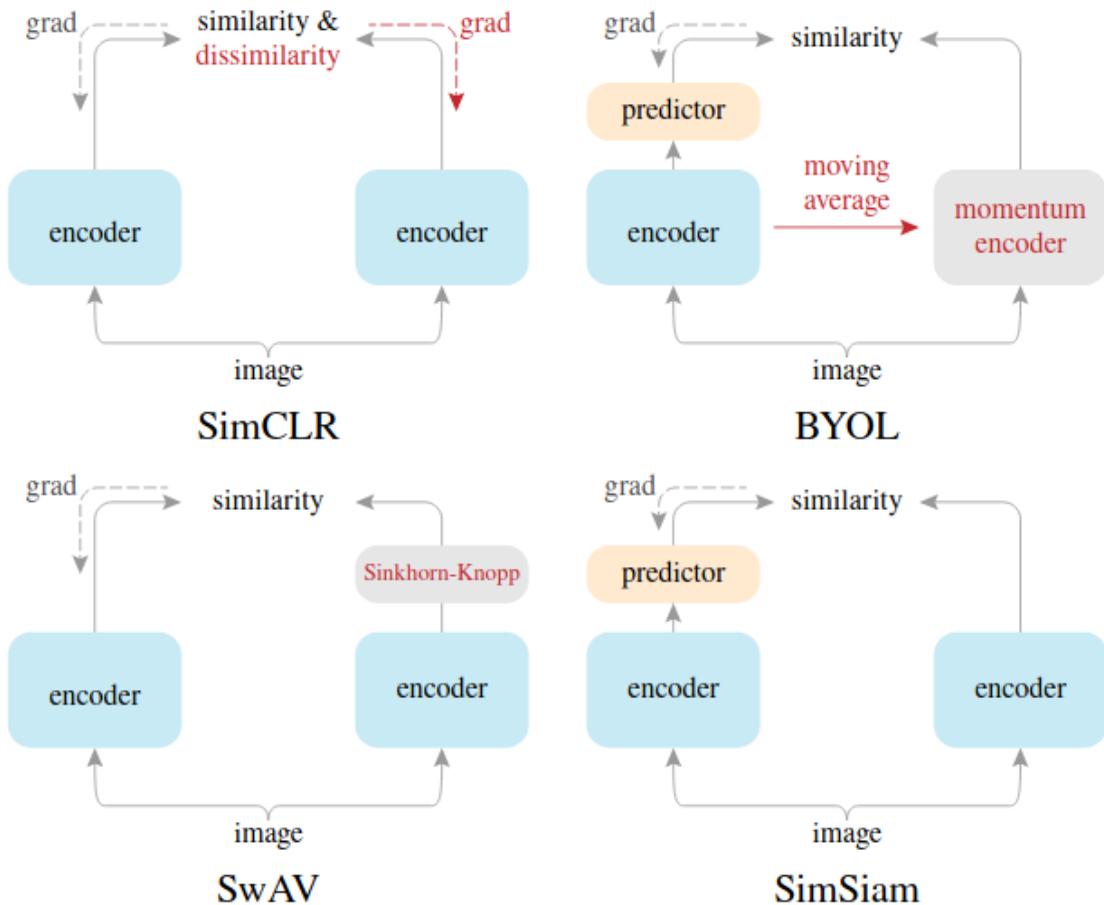


Figure 3. Comparison on Siamese architectures. The encoder includes all layers that can be shared between both branches. The dash lines indicate the gradient propagation flow. In BYOL, SwAV, and SimSiam, the lack of a dash line implies stop-gradient, and their symmetrization is not illustrated for simplicity. The components in red are those missing in SimSiam.

-
- Proj. Head: 3FCN of 2048-d.
 - Pred. Head: 2FCN of 2048-512-2048-d.

39 Graph Layouts by t-SNE

39.1 Backgrounds

Dimension Reduction

投影高维数据到低维, 可分为 distance-preserving 方法, 最小化 (aggregated normalized) stress

$$\sigma = \sum_{i,j} \left(\frac{d(\mathbf{x}_i, \mathbf{x}_j) - \|\mathbf{y}_i - \mathbf{y}_j\|}{d(\mathbf{x}_i, \mathbf{x}_j)} \right)^2 \quad (360)$$

以及 neighborhood-preserving 方法, 最大化 kNN 的重叠. distance-preserving 方法在特别高维时失效很严重 (由于高维空间的各种奇性).

39.2 Method

39.3 tsNET

40 细粒度图像数据分类 by Xiangteng He

困难: 标注成本巨大, 依赖人工先验, 忽略辨识速度 (like RPN, 2-stages), 忽略语义关联

40.1 细粒度图像数据分类

基于显著性图

40.2 RL-based 图像部件/对象识别

40.3 多层注意力区域辨识

基于 IoU 的奖励函数强化学习

40.4 多模态

- 引入文本信息, 音频, 视频
- 数据集构建 PKU FG-Xmedia
- ResNet-based 统一处理各种模态, 视频抽帧, 音频使用 mel 谱图, 文本使用 word embedding

Future: 大规模细粒度图像分类, 细粒度视觉推理, 图像到跨媒体迁移

41 Dirac Operator for Extrinsic Shape Analysis

为什么需要微分算子? 他们提供了一组流形上的基函数 (Hilbert 空间的基), 以及对应特征值. 流形上的函数可以在这些基上展开, 并且提供了傅立叶变换和卷积. 在紧流形上, 一个算子有离散的特征基, 若他是自伴和椭圆算子 (矩阵为对称的和正定的). 然而很多微分算子并不满足这些条件:

- Laplace-Beltrami Op., i.e. Δ , 一种曲面光滑性的度量.
- Hessian 矩阵, 和 modified Dirichlet 能量. $\sum_i E_D(N^i \phi)$
- 各项异性 Laplacian $\Delta_A = \operatorname{div} A \nabla$

对于曲面 M , L 算子可以看做 Dirichlet 能量的 Hessian. 我们提出 Dirac Op.

41.1 Math

使用四元数来表示曲面 (的一个嵌入 $f : M \hookrightarrow \operatorname{Im} \mathbb{H}$). 假设这个嵌入是共形 (conformal) 的. 使用 κ_1, κ_2 代表两个主曲率, 有高斯曲率 $K = \kappa_1 \kappa_2$ 和平均曲率 $H := \frac{1}{2}(\kappa_1 + \kappa_2)$. $|df|^2$ 代表嵌入 f 诱导的体积元.

对于实值函数, L 算子是最低阶非平凡自伴椭圆算子. 对于复变函数, 则存在 1-order 微分算子: Dirac 算子. 经典例子是 Cauchy-Riemann(Poincare) 算子 $\bar{\partial}$, 以及微分形式上的 Hodge-Dirac 算子 $\star d + d\star$.

我们使用的是 extrinsic D-Op.

$$D\psi := -\frac{df \wedge d\psi}{|df|^2} \quad (361)$$

其中 $|df|^2$ 的除法为应用 2-形式的 Hodge star 算子.

平方, 有

$$D^2\psi = \Delta\psi + \frac{dN \wedge d\psi}{|df|^2} \quad (362)$$

相对微分算子

$$D_{f_1, f_2}\psi := -\frac{df_2 \wedge d\psi}{|df_1|^2} \quad (363)$$

单参数内插

$$L(\tau) := (1 - \tau)\Delta + \tau D_N \quad (364)$$

算子 D_N 只是半正定的. 但是增加一些 L 算子的部分就能成为强正定的. 这些算子和曲率也有一定关系, $\Delta = \Delta_{\mathbb{R}^2} e^{-2u}$, u 是对数共形尺度因子, 且满足 Yamabe 方程 $\Delta u = K$. 由于要最小化 $\langle\langle \Delta\varphi, \varphi \rangle\rangle$, $\Delta\varphi$ 在有很大高斯曲率的地方会很小.

MDE(modified Dirichlet energy) 也可写成

$$E_{MDE}(\varphi) = \langle\langle \Delta\varphi, \varphi \rangle\rangle + \int_M \varphi^2 (\kappa_1^2 + \kappa_2^2) dA \quad (365)$$

这意味着它的 Hessian 是 $\Delta + U$, 后者是 Willmore 势能 $U = \kappa_1^2 + \kappa_2^2$, 这鼓励特征函数在曲率很大地方有小的方差.

41.2 离散化

对以一个 2-流形的三角 mesh $K = (V, E, F)$, 不同于连续的情况, 使用法向量 \mathbf{N} 代替 \mathbf{f} . 对于任何 mesh 上的三角 $ijk \in F, \psi : V \mapsto \mathbb{H}$ 有

$$(D_N \psi)_{ijk} := -\frac{1}{2\mathcal{A}_{ijk}} \sum_{pqr \in \mathcal{C}(ijk)} (N_r - N_q) \psi_p \quad (366)$$

其中 \mathcal{A}_{ijk} 是三角面积, $\mathcal{C}(ijk)$ 是轮换. 这个算子可以写成张量 $D \in \mathbb{H}^{|F| \times |V|}$

$$D_{ijk,p} = -(N_r - N_q) / 2\mathcal{A}_{ijk} \quad (367)$$

41.3 实值表示

用 4×4 矩阵代替四元数运算. 为了找到典型表示, 需要找到乘积常数 q , 使得某个特征向量范数接近 1. 近似的有

$$q := \sum_{i=1}^{|V|} \mathcal{A}_i \phi_i^{-1} \quad (368)$$

其中 A_i 是对偶区域的面积.

41.4 有界区域: 边界条件

边界条件可以是 Dirichlet/Neumann 条件. 这里提出无限势阱边界. 增加一个势能项 U , 并且在有界区域之外迅速趋向无穷, 使得优化 $\langle \{(\Delta + U)\psi, \psi\} \rangle$ 让波函数迅速在区域外收敛到 0.

$$U(p) := \frac{c}{1 + (e^{-(d(p,q)-\beta)})^\gamma} \quad (369)$$

q 是质心, c 是大常数.

为了使用这个罚函数, 在算子中加入对角矩阵 $U_{ii} = A_i U(p_i)$

42 Mesh-Based Simulation with GNNs

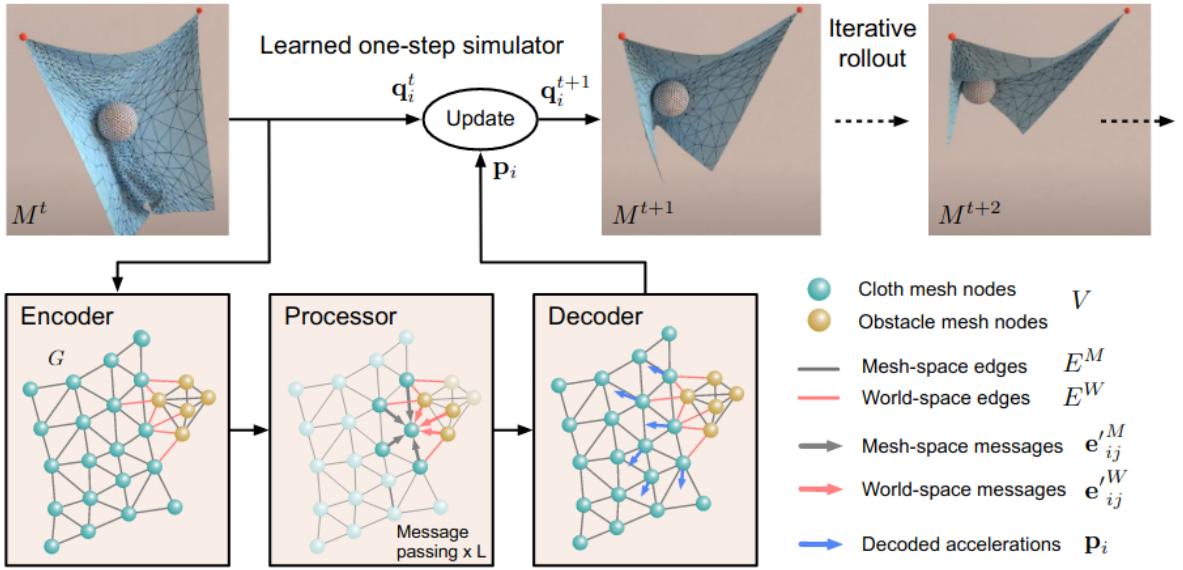
使用 World/Mesh 两种空间的输入. 预测物理体系的时间序列. 使用 $M^t = (V, E^M)$ 描述 t 时间的系统状态. 每个结点都有 mesh-space 坐标 \mathbf{u}_i 和动量/速度坐标 \mathbf{q}_i . 对于欧拉系统, 建模的是固定的 mesh 上的速度场, 对于 Lagrange 系统, mesh 是动态移动的, 并且包含了绝对世界坐标 \mathbf{x}_i .

42.1 结构

Encoder: 把 M^t 编码成多重图 $G = (V, E^M, E^W)$, 并且增加了世界坐标下的边 E^W . 这里使用简单的 rNN 采样, 选取半径为最短的 mesh 边长 r_W . 并且 mesh 图边权为相对坐标 $\mathbf{u}_{ij} = \mathbf{u}_i - \mathbf{u}_j$, 同理世界坐标图中也为相对坐标.

Processor: GNN. 其中

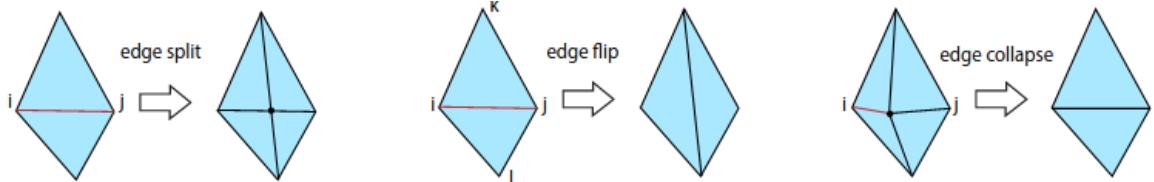
$$\mathbf{e}'_i = f^X(\mathbf{e}_i^X, \mathbf{v}_i, \mathbf{v}_j) \quad (370)$$



Decoder: 使用在结点 feature 上的 MLP. 得到输出向量 p_i

Updater: 对于一阶系统, $\mathbf{q}_i^{t+1} = \mathbf{p}_i + \mathbf{q}_i^t$, 二阶系统: $\mathbf{q}_i^{t+1} = \mathbf{p}_i + 2\mathbf{q}_i^t - \mathbf{q}_i^{t-1}$. 输出向量的额外维度也用于预测附加属性(压力/张力 etc.)

4.2.2 Adaptive Remeshing



在需要的地方增加 mesh 数量.(有限元那味儿) 使用 sizing filed 方法. 定义 sizing field 张量

$\mathbf{S}(\mathbf{u}) \in \mathbb{R}^{2 \times 2}$, 同时一条边是合理的 iff. $\mathbf{u}_{ij}^T \mathbf{S}_i \mathbf{u}_{ij} \leq 1$, 否则需要拆分这条边.

使用动态的 remeshing. 使用如上文的 GNN 的结构来学习 sizing 张量场. 每一步, 我们都计算动量信息和 resizing 张量场, 再用 remesher 计算下一步的 mesh.

给定了 sizing 张量场 \mathbf{S}_i , 对于每条边,

- 需要分割, 若 $\mathbf{u}_{ij}^T \mathbf{S}_{ij} \mathbf{u}_{ij} > 1$, $\mathbf{S}_{ij} = \frac{1}{2} (\mathbf{S}_i + \mathbf{S}_j)$
- 需要 collapse, 若 collapse 不会产生新的无效边
- 需要 flip, 若满足各项异性 Delaunay 条件

$$(\mathbf{u}_{jk} \times \mathbf{u}_{ik}) \mathbf{u}_{il}^T \mathbf{S}_A \mathbf{u}_{jl} < \mathbf{u}_{jk}^T \mathbf{S}_A \mathbf{u}_{ik} (\mathbf{u}_{il} \times \mathbf{u}_{jl}), \quad \mathbf{S}_A = \frac{1}{4} (\mathbf{S}_i + \mathbf{S}_j + \mathbf{S}_k + \mathbf{S}_l) \quad (371)$$

顺序上, 先尽可能 split 所有边, 再尽可能 collapse 所有边, 再尽可能 flip 所有边
若没有能估计 sizing field 的数据, 我们从 mesh 序列中估计之. 我们要找到能诱导出下一时刻 mesh 的 sizing field. 假设 remesher 几乎是最优的, 所有结果边都有效, 那么最大化度规 \mathbf{S} 下的边长

$$\mathbf{S}_i = \operatorname{argmax} \sum_{j \in \mathcal{N}_i} \mathbf{u}_{ij}^T \mathbf{S}_i \mathbf{u}_{ij}, \quad \text{s.t. } \forall j \in \mathcal{N}_i : \mathbf{u}_{ij}^T \mathbf{S}_i \mathbf{u}_{ij} \leq 1 \quad (372)$$

这是简单的凸优化 (找到最小面积, 0-中心的包含 \mathbf{u}_{ij} 的椭圆).

43 SENet

Squeeze-and-Excitation Module.




Fig. 1. A Squeeze-and-Excitation block.

计算 squeeze 变换 (单纯的计算每个 channel 的均值)

$$z_c = \mathbf{F}_{sq}(\mathbf{u}_c) = \frac{1}{H \times W} \sum_{i=1}^H \sum_{j=1}^W u_c(i, j) \quad (373)$$

计算 excitation, σ 是 sigmoid, δ 是 ReLU

$$\mathbf{s} = \mathbf{F}_{ex}(\mathbf{z}, \mathbf{W}) = \sigma(g(\mathbf{z}, \mathbf{W})) = \sigma(\mathbf{W}_2 \delta(\mathbf{W}_1 \mathbf{z})) \quad (374)$$

其中 $\mathbf{W}_1 \in \mathbb{R}^{\frac{C}{r} \times C}$, $\mathbf{W}_2 \in \mathbb{R}^{C \times \frac{C}{r}}$ 维度形成某种 bottleneck, 降低模型复杂度.

将 excitation 乘到 channel 上

$$\tilde{\mathbf{x}}_c = \mathbf{F}_{scale}(\mathbf{u}_c, s_c) = s_c \mathbf{u}_c \quad (375)$$

13

¹³in matrix-sense,

$$\mathbf{z} = \frac{1}{HW} \mathbf{u}_{ij} \mathbf{1}^{ij} \in \mathbb{R}^{C \times 1} \quad (376)$$

$$\mathbf{s} = \mathbf{F}_{ex}(\mathbf{z}) \in \mathbb{R}^{C \times 1} \quad (377)$$

$$\tilde{\mathbf{X}} = \mathbf{s}^T \odot \mathbf{U} \in \mathbb{R}^{N(-H \times W) \times C} \quad (378)$$

44 Continuous-Time Spiking Neural Network

模拟真实神经元. alternative to Hebb 神经元. 传统的 SNN 模拟手段涉及时间驱动 (time-driven) 和事件驱动 (event-driven), 前者涉及时间格点上的离散化, 可能带来误差. 采用 event-driven 来实现连续时间.

44.1 Neurons

$$\begin{aligned} S &= S_p + P_r P_w - T_l, \text{ for } S < S_{th} \\ S &= S_p + P_r P_w + T_r, \text{ for } S \geq S_{th} \end{aligned} \quad (379)$$

- S_p , 之前的状态
- P_r : presynaptic weight, 受到的刺激
- P_w : postsynaptic weight, 神经元连接强度
- T_l : 电位 leakage $Ld\Delta t$
- T_r : $T_r = \frac{(S_p-1)^2 \Delta t}{1-(S_p-1)\Delta t}$
- t_f : firing time, obeys firing equation $t_f = \frac{1}{S-1}$

可以有 excitatory/inhibitory 的输入区别.

44.2 Network Topology

一般来说, en 比 in 多, 后者占网络的 15-25%, 这样可以增进网络的稳定性. 没有 in-in 突触, 这回导致不稳定性, 因为这会让 in 减少, 进而导致不受控制的激发.

44.3 突触塑性规则

”后突触规则”: 是相对于同一突触的突触总体模式/时间点控制了后突触效用.

- 对数衰减. 所有后突触权重 postsynaptic weight 对数衰减.

$$P_w = P_{w,\min} + (P_w - P_{w,\min}) e^{-\frac{\Delta t}{\tau}}$$

- 同突触增强. 当一个 spiking event 在一个突触上发生, PW 增加, 和同神经元 + 同突触的上一个刺激成函数 (在一个时间窗口内).
- 异突触增强. 当一个 spiking event 在一个突触上发生, PW 增加, 和同神经元 + 其他突触的上一个刺激成函数 (在一个时间窗口内).

同/异突触增强使用以下方程控制.

$$\Delta P_w = \eta (P_{w,\max} - P_w) \quad (380)$$

子后, 还可以采取其他策略, 如 STDP, Synaptic Scaling etc.

45 Towards Deep Learning Models Resistant to Adversarial Attacks

从优化的角度来看, 这是个 robust-optimization 的问题. 给出 empirical loss $\mathbb{E}_{(x,y) \sim \mathcal{D}}[J(x, y, \theta)]$, 我们可以适用 robust optimization 模型

$$\begin{aligned} & \min_{\theta} \rho(\theta) \\ \text{where } \rho(\theta) = & \mathbb{E}_{(x,y) \sim \mathcal{D}} [\max_{\delta \in \mathcal{S}} J(\theta, x + \delta, y)] \end{aligned} \tag{381}$$

其中 perturbation set 可以是 l_∞ -球. 当然也可以基于别的视觉相似度.

45.1 Inner Maximize Prob. : 如何提出好的对抗样本

关于对抗样本和攻击, 存在两个问题:

1. 如何获得强的对抗样本
2. 如何训练一个没有/难以获得对抗样本的网络.

Fast Gradient Sign Method(FGSM) 是一个如下获得 l_∞ -有界对抗样本的

$$x + \varepsilon \operatorname{sgn}(\nabla_x J(\theta, x, y)) \tag{382}$$

一个更强大的方法可以进行多步:FGSM^k, 基本上就是多步 Projected GD

$$x^{t+1} = \Pi_{x+\mathcal{S}} (x^t + \alpha \operatorname{sgn}(\nabla_x J(\theta, x, y))) \tag{383}$$

实验证明, l_∞ -球中包含了很多有着相似 loss 值的 local maxima. 具体上讲,

1. PGD 导致的 loss 上升具有同种趋势 (曲线形状)
2. PGD 导致的最终 loss 非常密集
3. 计算了这些最终点的距离, 都和平均距离相差不大, 且趋向于成 90°.

这说明了 PGD 是泛用的一阶对抗方法.

猜想: PGD 对抗样本不会比其他一阶样本的 loss 小. 那么, 就是用 PGD 来最大化内部最大化问题.

45.2 最小化问题

SGD, 使用 perturbed points 来进行最小化的输入. Danskin 定理说对于连续可微函数¹⁴, 此时的梯度确实是 saddle/minmax optimization 的一个下降梯度. 虽然 ReLU/max-pooling 并不是连续可微的, 但是大概用也行.

¹⁴待补充

45.3 网络能力 & 对抗健壮性

对抗健壮性和网络能力相关. high-level 的讲, 更强的分类器趋向于有好的健壮性, 因为对抗样本的存在让分类边界更复杂.

一些观察

1. 网络容量提高了对抗性能. 训练超长时间提高了对抗性能.
2. one-step attack 也能提高健壮性
3. 弱模型可能再 saddle opt. 中不会收敛
4. 更高的容量和强对抗降低 transferred attacks 的有效性.

46 CAS: Channel-wise Activation Suppressing Module for Adversarial Robustness

考虑 channel-wise 的激活强度. 他们发现对抗训练可以显著降低对抗样本的 channel-wise act. magnitude, 以至于和自然样本相当近似. 考虑 channel-wise 激活频率, 可以发现对抗样本的激活频率非常均匀 (按照自然样本输入的激活频率排序), 对抗训练可以对齐这两个趋势, CAS 可以直接大大减少 activation 频率.

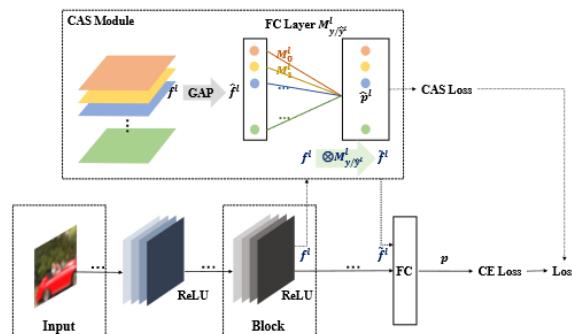


Figure 3: Framework of our proposed Channel-wise Activation Suppressing (CAS).

假设第 1 层的网络输出是 f^l , 考虑 in-channel average

$$\hat{f}_k^l = \frac{1}{H \times W} \sum_{i=1}^H \sum_{j=1}^W f_k^l(i, j) \quad (384)$$

这被送进一个辅助 MLP(FC) classifier

$$\hat{p}^l = \text{softmax}(\hat{f}^l M^l) \in \mathbb{R}^C \quad (385)$$

其中每个类别的参数可以写成 $M^l = [M_1^l, M_2^l, \dots, M_C^l] \in \mathbb{R}^{K \times C}$ 用 GT 的 label 分量去按通道乘 (在 test set 上直接用 classifier 的结果来计算)act. map

$$\tilde{\mathbf{f}}^l = \begin{cases} \mathbf{f}^l \otimes M_y^l, & (\text{training phase}) \\ \mathbf{f}^l \otimes M_{\hat{y}^l}^l, & (\text{test phase}) \end{cases} \quad (386)$$

模型训练基于两个 loss, 其中每个 CAS loss 为

$$\mathcal{L}_{\text{CAS}}(\hat{\mathbf{p}}^l(\mathbf{x}', \theta, M), y) = - \sum_{c=1}^C \mathbb{1}\{c = y\} \cdot \log \hat{p}_c^l(\mathbf{x}') \quad (387)$$

注意这就是在辅助分类器上的 CE loss. 同时还有主 Loss

$$\mathcal{L}(\mathbf{x}', y; \theta, M) = \mathcal{L}_{\text{CE}}(\mathbf{p}(\mathbf{x}', \theta), y) + \frac{\beta}{S} \cdot \sum_{s=1}^S \mathcal{L}_{\text{CAS}}^s(\hat{\mathbf{p}}^s(\mathbf{x}', \theta, M), y) \quad (388)$$

使用标准的 AT 技术. 也可使用 TRADES/MART.

47 Resisting Adversarial Attacks by *k*-Winners-Takes-All

不使用梯度信息. 试图使用 C^0 不连续函数, 即 k -Winners-Takes-All(kWTA) 激活函数 (代替 ReLU). kWTA 激活函数的优势是, 具有稠密的不连续点, 使得难以利用小的 perturbation 可控地改变结果, 因为这会经过不连续的部分, 导致梯度信息失效. 可以证明, 即使是无穷小位移也会导致神经元激活模式彻底改变!

* k-WTA 可以看作是 ReLU 和 max-pool 的结合.

47.1 Related Work: Obfuscated Gradients

一种类型是利用随机性: 在输入前进行随机变换/增加随机化的计算层. 然而梯度仍然可以使用 MC 估计来得到, 仍然不是很 robust. 另一种是 shattered gradient, 它们让梯度对于攻击者不存在或者不正确, 通过故意离散化输入/人工地提高梯度数值不稳定性. 这些方法仍然 vulnerable, 可以通过 backward pass differentiable approx.(BPDA) 来估计梯度. 假设 $f = f_1 \circ f_2 \circ \dots \circ f_n$, 其中 $f_i(\mathbf{x})$ 是不连续的函数组分, 可以使用 smooth delegate function g 来近似这些不连续函数! 对于 k-WTA, 他们没找到任何平滑的近似函数.

47.2 *k*-Winners-Takes-All

k-Winners-Takes-AllModule:

$$\phi_k(\mathbf{y})_j = \begin{cases} y_j, & y_j \in \{k \text{ largest elements of } \mathbf{y}\} \\ 0, & \text{Otherwise} \end{cases} \quad (389)$$

接着显然要决定 k 的大小. 使用自适应的大小: 定义 $\gamma \in (0, 1)$ 作为比例系数, 称为稀疏度比例. 并且对所有层都固定. k-WTA 可以在每个 channel 上进行, 但是计算最有效/最简单的是直接在整个 tensor 上进行.

渐进时间复杂度 $O(N)$ ¹⁵

47.3 Training

小的 γ 是 preferable 的, 但是越小训练越难以收敛. 所以使用一种迭代 finetune 的方法. 使用线性退火降低稀疏系数, 每次降低后训练两个 epoch.

47.4 Theory Understand of the Discontinuity

定义 activation pattern

$$\mathcal{A}(\mathbf{x}) := \{i \in [l] \mid x_i \text{ is one of the } k \text{ largest values in } \mathbf{x}\} \subseteq [l] \quad (390)$$

可以证明, 极小的输入变化也会导致 $\mathcal{A}(\mathbf{x})$ 的变化. 整个网络是一个分段 linear 的映射.

注意 $\phi_k(W\mathbf{x} + \mathbf{b})$ 的 act. pat. 对于正定线性变换不变, 应该衡量垂直于输入 \mathbf{x} 的分量的距离 $d(\mathbf{x}, \mathbf{x}')$, 且 $\mathbf{x}' = c \cdot (\mathbf{x} + d(\mathbf{x}, \mathbf{x}') \mathbf{x}_\perp)$ ¹⁶, 那么若权值矩阵 $W_{ij} \sim \mathcal{N}(0, 1/l)$, $\mathbf{b} = 0$ 初始化, 有以下定理

Theorem 47.1 (*Dense discontinuities*). Given any input $\mathbf{x} \in \mathbb{R}^m$ and some β , and $\forall \mathbf{x}' \in \mathbb{R}^m$ such that $\frac{d^2(\mathbf{x}, \mathbf{x}')}{\|\mathbf{x}\|_2^2} \geq \beta$, if the following condition

$$l \geq \Omega\left(\left(\frac{m}{\gamma} \cdot \frac{1}{\beta}\right) \cdot \log\left(\frac{m}{\gamma} \cdot \frac{1}{\beta}\right)\right)$$

is satisfied, then with a probability at least $1 - 2^{-m}$, we have $\mathcal{A}(W\mathbf{x} + \mathbf{b}) \neq \mathcal{A}(W\mathbf{x}' + \mathbf{b})$.

此处 l 是层宽.

为什么 k-WTA 可以被训练? 即使那么不连续.

Theorem 47.2 Consider N data points $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N \in \mathbb{R}^m$. Suppose $\forall i \neq j$, $\frac{\mathbf{x}_i}{\|\mathbf{x}_i\|_2} \neq \frac{\mathbf{x}_j}{\|\mathbf{x}_j\|_2}$. If N is sufficiently large, then with a high probability, we have $\forall i \neq j$, $\mathcal{A}(W\mathbf{x}_i + \mathbf{b}) \cap \mathcal{A}(W\mathbf{x}_j + \mathbf{b}) = \emptyset$.

所以网络对于不同的输入可以独立的训练, 实际上这些 act. pat. 往往是弱相关的.

实际上, 这些性质在 ReLU 中也有体现

47.5 Related Works

White-box attack model: PDG, Deep-fool, C&W Attack, Momentum Iterative Method(MIM).

Black-box(BB) attacks (Papernot et al., 2017). 基于 transfer.

Adversarial Training: AT, TRADES, free adversarial training(FAT).

¹⁵ Problem: How to impl. in CUDA?

¹⁶i.e., $d(\mathbf{x}, \mathbf{x}') = (\mathbf{x} - \mathbf{x}')/\mathbf{x}_\perp$

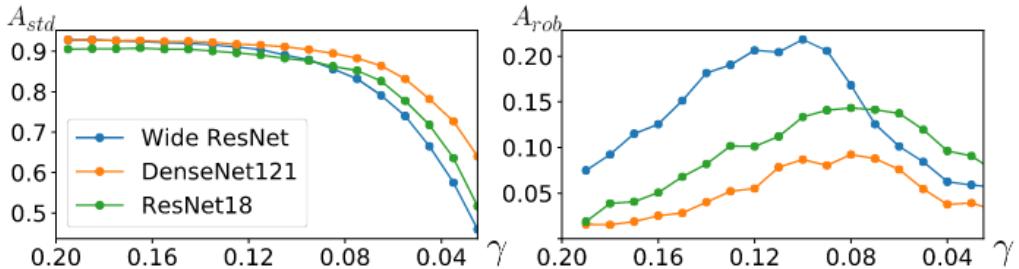


Figure 4: Robustness changing w.r.t. γ on CIFAR. When γ decreases, the standard test accuracy (left) starts to drop after a certain point. The robust accuracy (right) first increases then decreases.

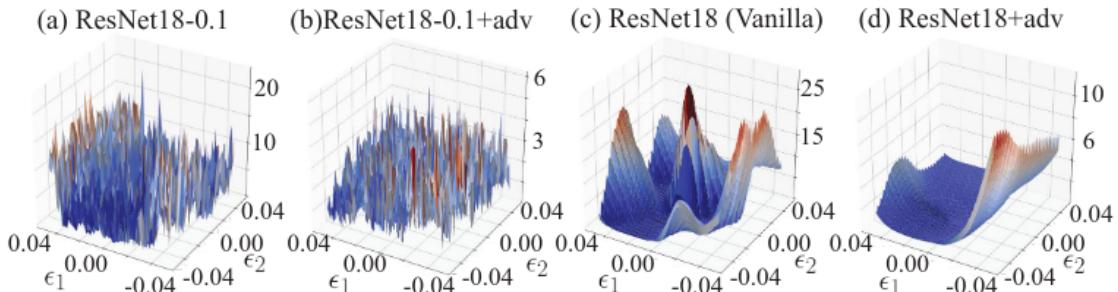


Figure 5: Gradient-based attack's loss landscapes in k -WTA (a, b) and conventional ReLU models (c, d). (a,b) k -WTA Models have much more non-convex and non-smooth landscapes. Also, the model optimized by adversarial training (b) has a lower absolute value of loss.